



University of  
Pittsburgh

# Introduction to Operating Systems CS 1550



Spring 2023  
Sherif Khattab  
ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013**)

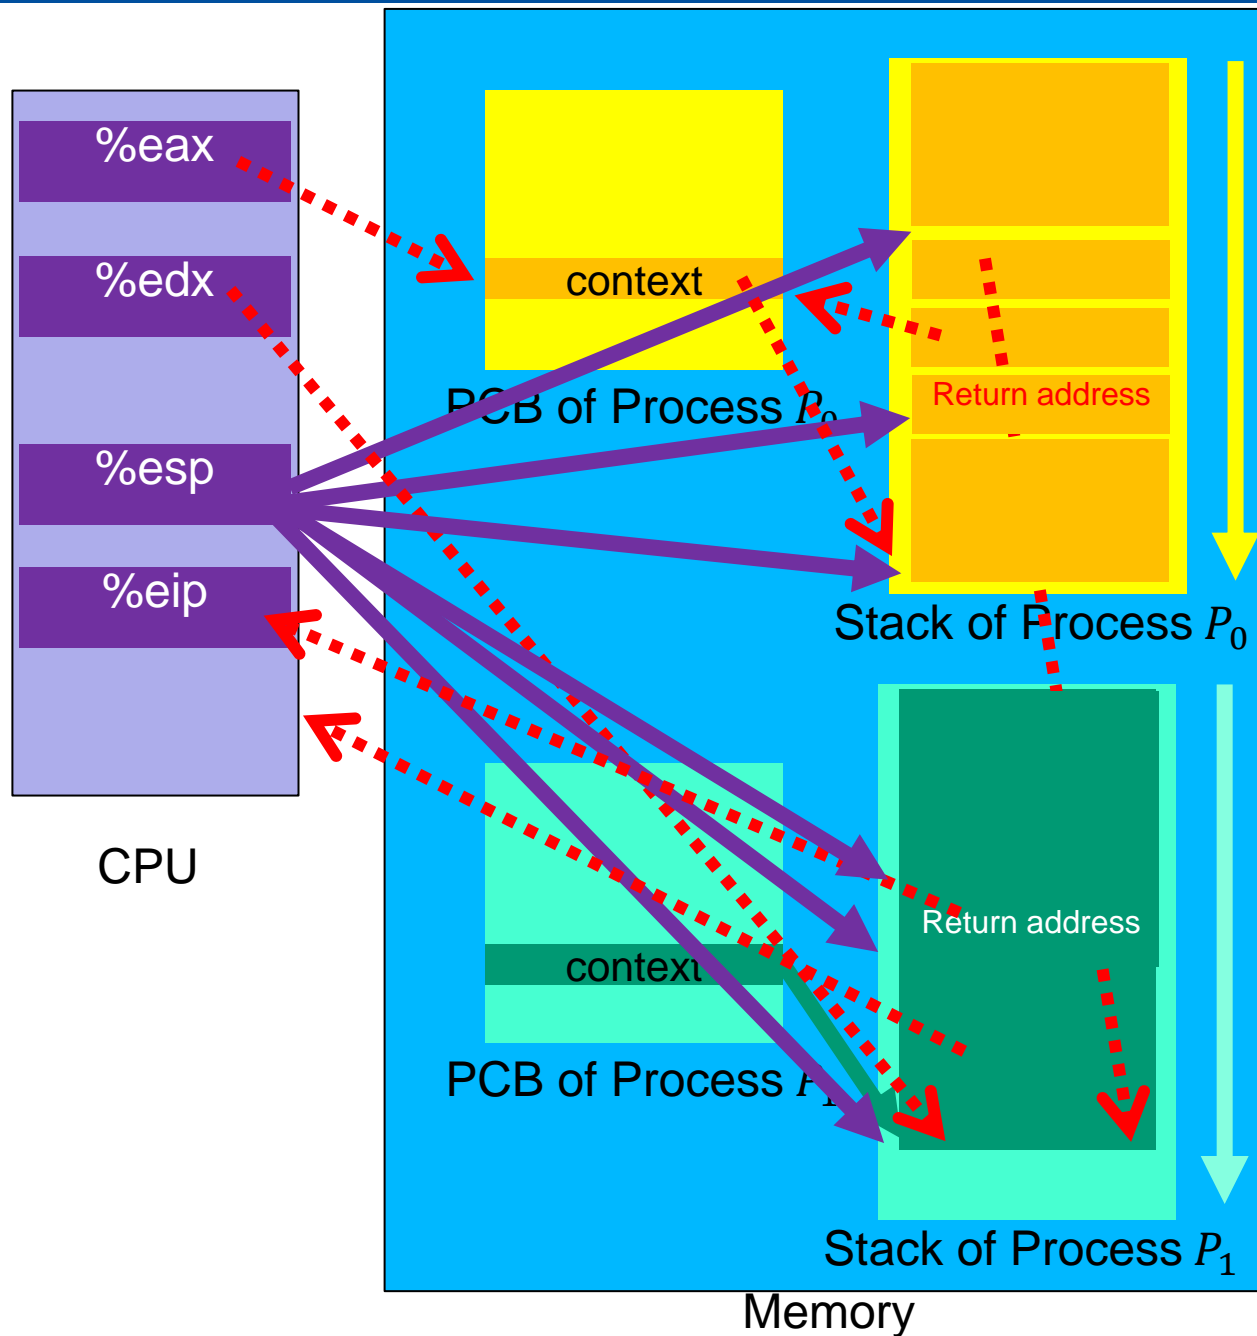
# Announcements

- Homework 1 is due tomorrow at 11:59 pm
- Recitations start this week
- Project 1 will be posted on Canvas this Friday
- Docker images for labs and projects are available on Canvas
  - As an alternative to running the labs and projects on the Thoth server

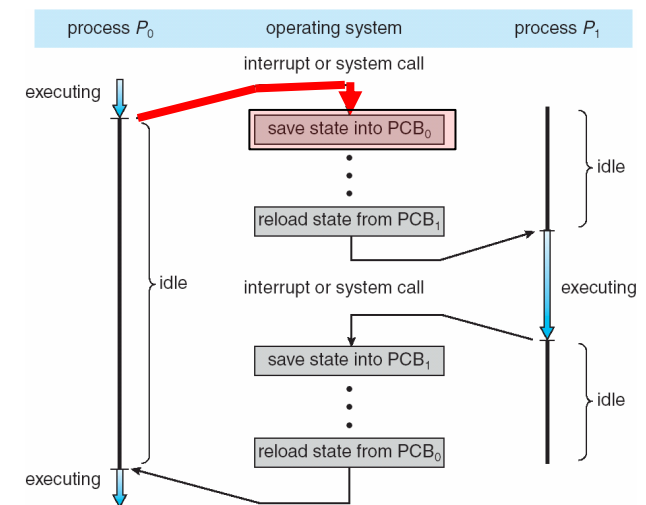
# Agenda

- Context Switching
- Critical Region as a solution to the Race Condition problem
- Spinlocks to implement Critical Region
- Busy Waiting Problem
- Why does it happen?
- What are its implications?

# Context Switching in Xv6




```
1 .globl switch
2 switch:
3     movl 4(%esp), %eax
4     movl 8(%esp), %edx
5
6     # Save old callee-saved registers
7     pushl %ebp
8     pushl %ebx
9     pushl %esi
10    pushl %edi
11
12    # Switch stacks
13    movl %esp, (%eax)
14    movl %edx, %esp
15
16    # Load new callee-saved registers
17    popl %edi
18    popl %esi
19    popl %ebx
20    popl %ebp
21    ret
```



# Critical regions

- Back to the race conditions problem
- Use critical regions to provide *mutual exclusion* and help fix race conditions
- Let's put the statement `x++` in a critical region



# How to implement critical regions?

- Turn-based solutions
- Spinlocks
- Semaphores
- Monitors

# Using Spinlocks

```
Spinlock lock;
```

**Code for process  $P_i$**

```
While(1) {  
    Lock(lock)  
    // critical section  
    Unlock(lock);  
    // remainder of code  
}  
}
```

# Spinlock implementation (1/2)

- Solution: use hardware
- Several hardware methods
  - Test & set: test a variable and set it in one instruction
  - Atomic swap: switch register & memory in one instruction
  - Turn off interrupts: process won't be switched out unless it asks to be suspended
- The first two methods can be implemented in user land
  - Why can't we implement the third method in user land?



# Busy Waiting

- A process that is trying to acquire a locked spinlock is running!
  - It continuously checks:
    - can I get the lock? No, lock is held by another process
    - can I get the lock? No, lock is held by another process
    - ...
  - This continuous check is called **spinning** or **busy waiting**
  - But what is wrong with that?
    - Busy waiting wastes CPU cycles
    - on a single-core system it delays the process that is holding the lock from releasing it

# Today's problem: Busy Waiting

While P1 is in the critical region, P2 is busy waiting

## Shared Data

Spinlock lk;

int x;

## Process P1

lock(lk);

//critical region (e.g., x++)

unlock(lk);

## Process P2

lock(lk);

//critical region (e.g., x++)

unlock(lk);

# But why?

Why does busy waiting happen with spinlocks?

# Atomic TestAndSet

- TestAndSet is an atomic instruction
- Works for single-core and multi-core Symmetric Multi-Processing (SMP)

```
int TestAndSet(int &x){  
    lock memory access to x  
  
    int temp = *x;  
  
    *x = 1;  
  
    unlock memory access to x  
  
    return temp;  
}
```

# Spinlock implementation using TestAndSet

- Single **shared** variable: lock
- Works for any number of processes

```
int lock = 0;  
  
Lock () {  
    while (TestAndSet (&lock) )  
        ;  
}  
  
Unlock () {  
    lock = 0;  
}
```

# Atomic Swap

- Swap is an atomic instruction
- Works for single-core and multi-core Symmetric Multi-Processing (SMP)

```
int Swap(int &x, int y){  
    lock memory access to x  
  
    int temp = *x;  
  
    *x = y;  
  
    unlock memory access to x  
  
    return temp;  
}
```

# Spinlock implementation using Swap

- Single **shared** variable: lock
- Works for any number of processes

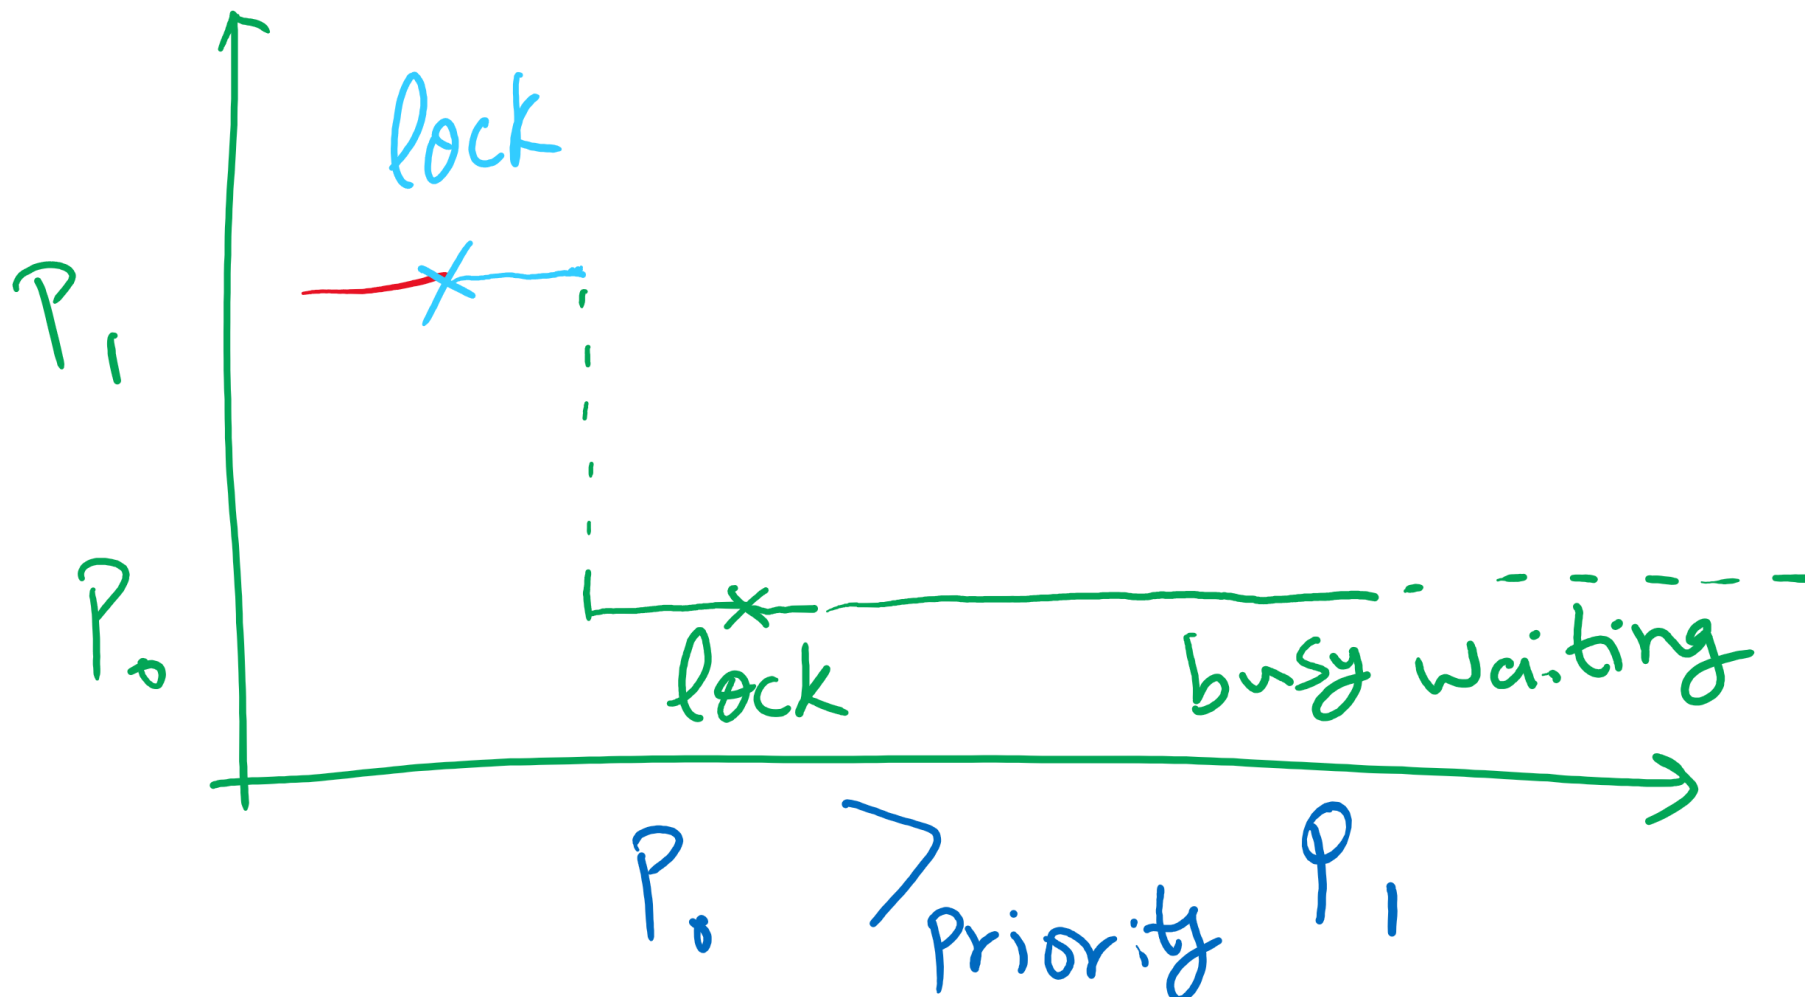
```
int lock = 0;

Lock() {
    while (Swap(&lock, 1))
        ;
}

Unlock() {
    lock = 0;
}
```

# Implication of Busy Waiting

Subproblem: *priority inversion* (higher priority process busy waits for lower priority process)





# Implementation Detail

compiler and/or hardware may reorder instructions

