



University of
Pittsburgh

Introduction to Operating Systems CS 1550



Spring 2023
Sherif Khattab
ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013**)

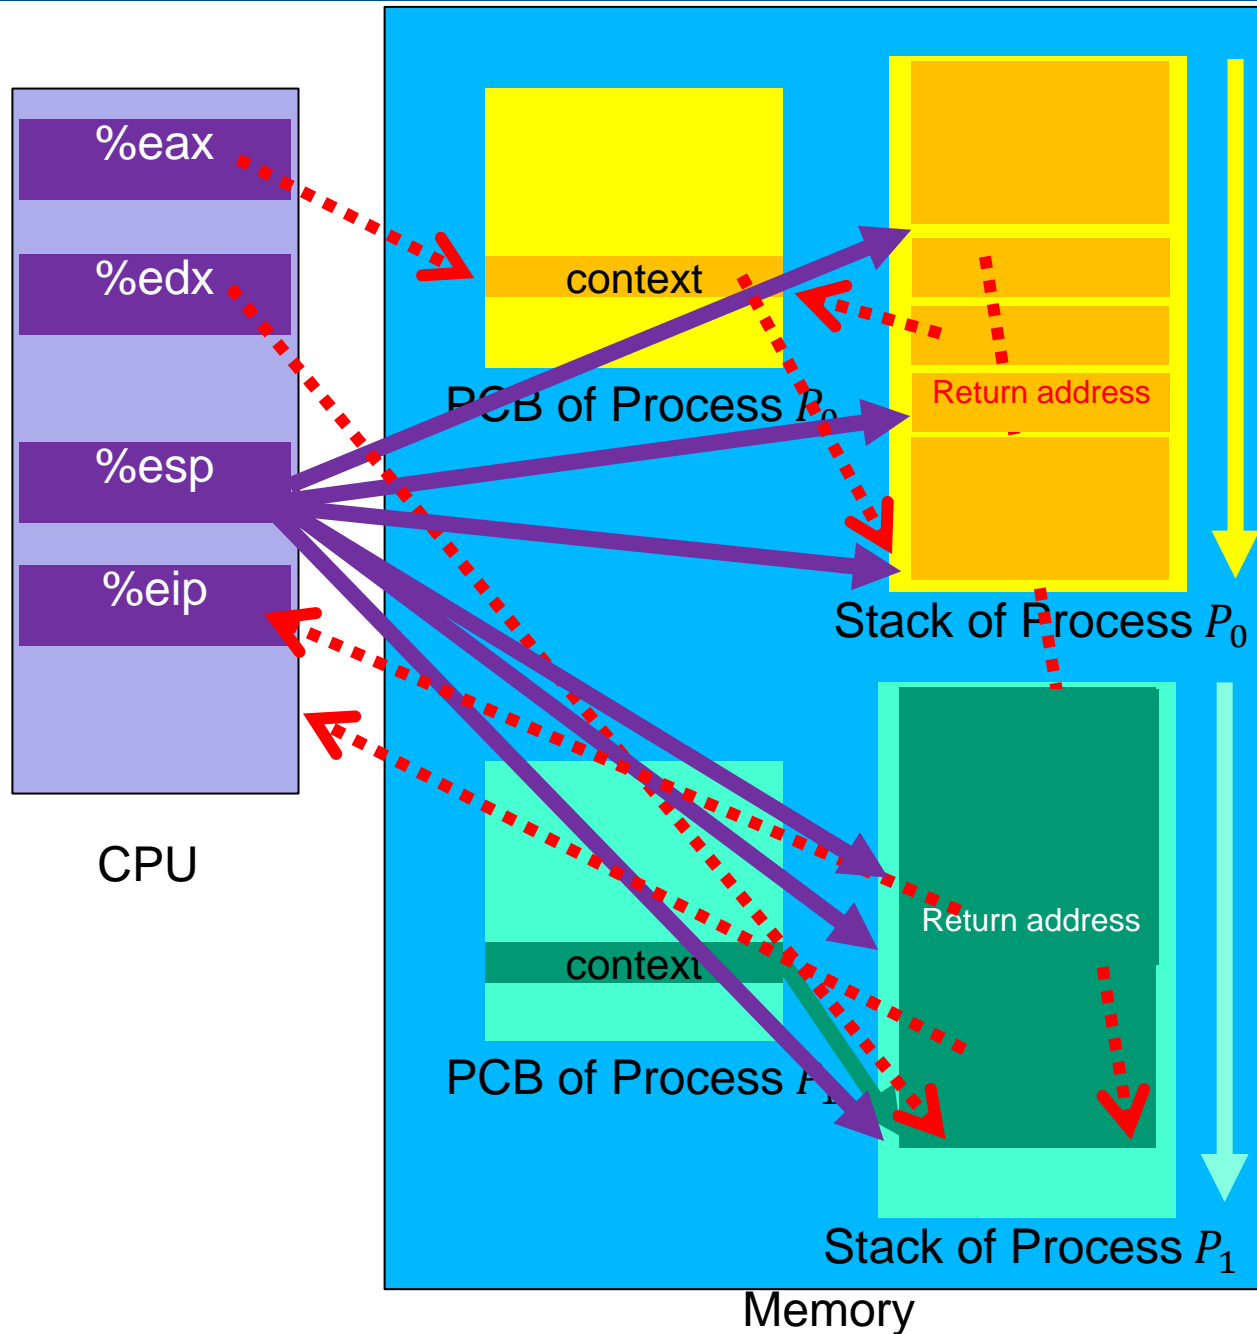
Announcements

- Homework 1 is due tomorrow at 11:59 pm
- Recitations start this week
- Project 1 will be posted on Canvas this Friday
- Docker images for labs and projects are available on Canvas
 - As an alternative to running the labs and projects on the Thoth server

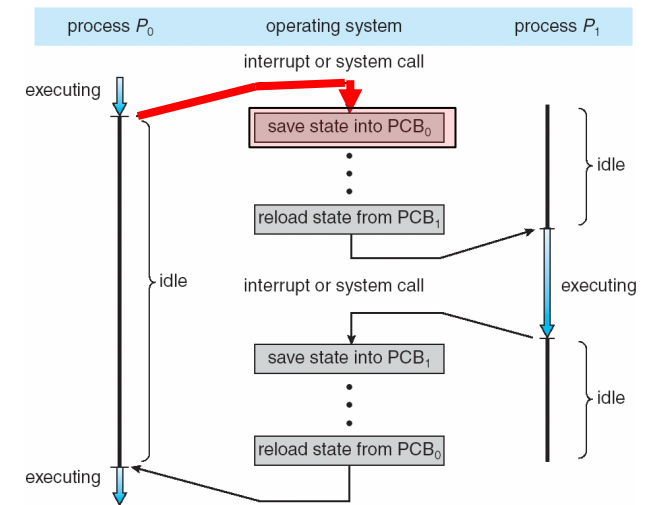
Agenda

- Context Switching
- Critical Region as a solution to the Race Condition problem
- Spinlocks to implement Critical Region
- Busy Waiting Problem
- Why does it happen?
- What are its implications?
- How to solve it?
 - Sempahores

Context Switching in Xv6




```
1 .globl switch
2 switch:
3     movl 4(%esp), %eax
4     movl 8(%esp), %edx
5
6     # Save old callee-saved registers
7     pushl %ebp
8     pushl %ebx
9     pushl %esi
10    pushl %edi
11
12    # Switch stacks
13    movl %esp, (%eax)
14    movl %edx, %esp
15
16    # Load new callee-saved registers
17    popl %edi
18    popl %esi
19    popl %ebx
20    popl %ebp
21    ret
```



Critical regions

- Back to the race conditions problem
- Use critical regions to provide *mutual exclusion* and help fix race conditions
- Let's put the statement `x++` in a critical region



How to implement critical regions?

- Turn-based solutions
- Spinlocks
- Semaphores
- Monitors

Using Spinlocks

```
Spinlock lock;
```

Code for process P_i

```
While(1) {  
    Lock(lock)  
    // critical section  
    Unlock(lock);  
    // remainder of code  
}  
}
```

Spinlock implementation (1/2)

- Solution: use hardware
- Several hardware methods
 - Test & set: test a variable and set it in one instruction
 - Atomic swap: switch register & memory in one instruction
 - Turn off interrupts: process won't be switched out unless it asks to be suspended
- The first two methods can be implemented in user land
 - Why can't we implement the third method in user land?

Busy Waiting

- A process that is trying to acquire a locked spinlock is running!
 - It continuously checks:
 - can I get the lock? No, lock is held by another process
 - can I get the lock? No, lock is held by another process
 - ...
 - This continuous check is called **spinning** or **busy waiting**
 - But what is wrong with that?
 - Busy waiting wastes CPU cycles
 - on a single-core system it delays the process that is holding the lock from releasing it

Today's problem: Busy Waiting

While P1 is in the critical region, P2 is busy waiting

Shared Data

Spinlock lk;

int x;

Process P1

lock(lk);

//critical region (e.g., x++)

unlock(lk);

Process P2

lock(lk);

//critical region (e.g., x++)

unlock(lk);

But why?

Why does busy waiting happen with spinlocks?

Atomic TestAndSet

- TestAndSet is an atomic instruction
- Works for single-core and multi-core Symmetric Multi-Processing (SMP)

```
int TestAndSet(int &x){  
    lock memory access to x  
  
    int temp = *x;  
  
    *x = 1;  
  
    unlock memory access to x  
  
    return temp;  
}
```

Spinlock implementation using TestAndSet

- Single **shared** variable: lock
- Works for any number of processes

```
int lock = 0;

Lock () {
    while (TestAndSet (&lock) )
        ;
}

Unlock () {
    lock = 0;
}
```

Atomic Swap

- Swap is an atomic instruction
- Works for single-core and multi-core Symmetric Multi-Processing (SMP)

```
int Swap(int &x, int y){  
    lock memory access to x  
  
    int temp = *x;  
  
    *x = y;  
  
    unlock memory access to x  
  
    return temp;  
}
```

Spinlock implementation using Swap

- Single **shared** variable: lock
- Works for any number of processes

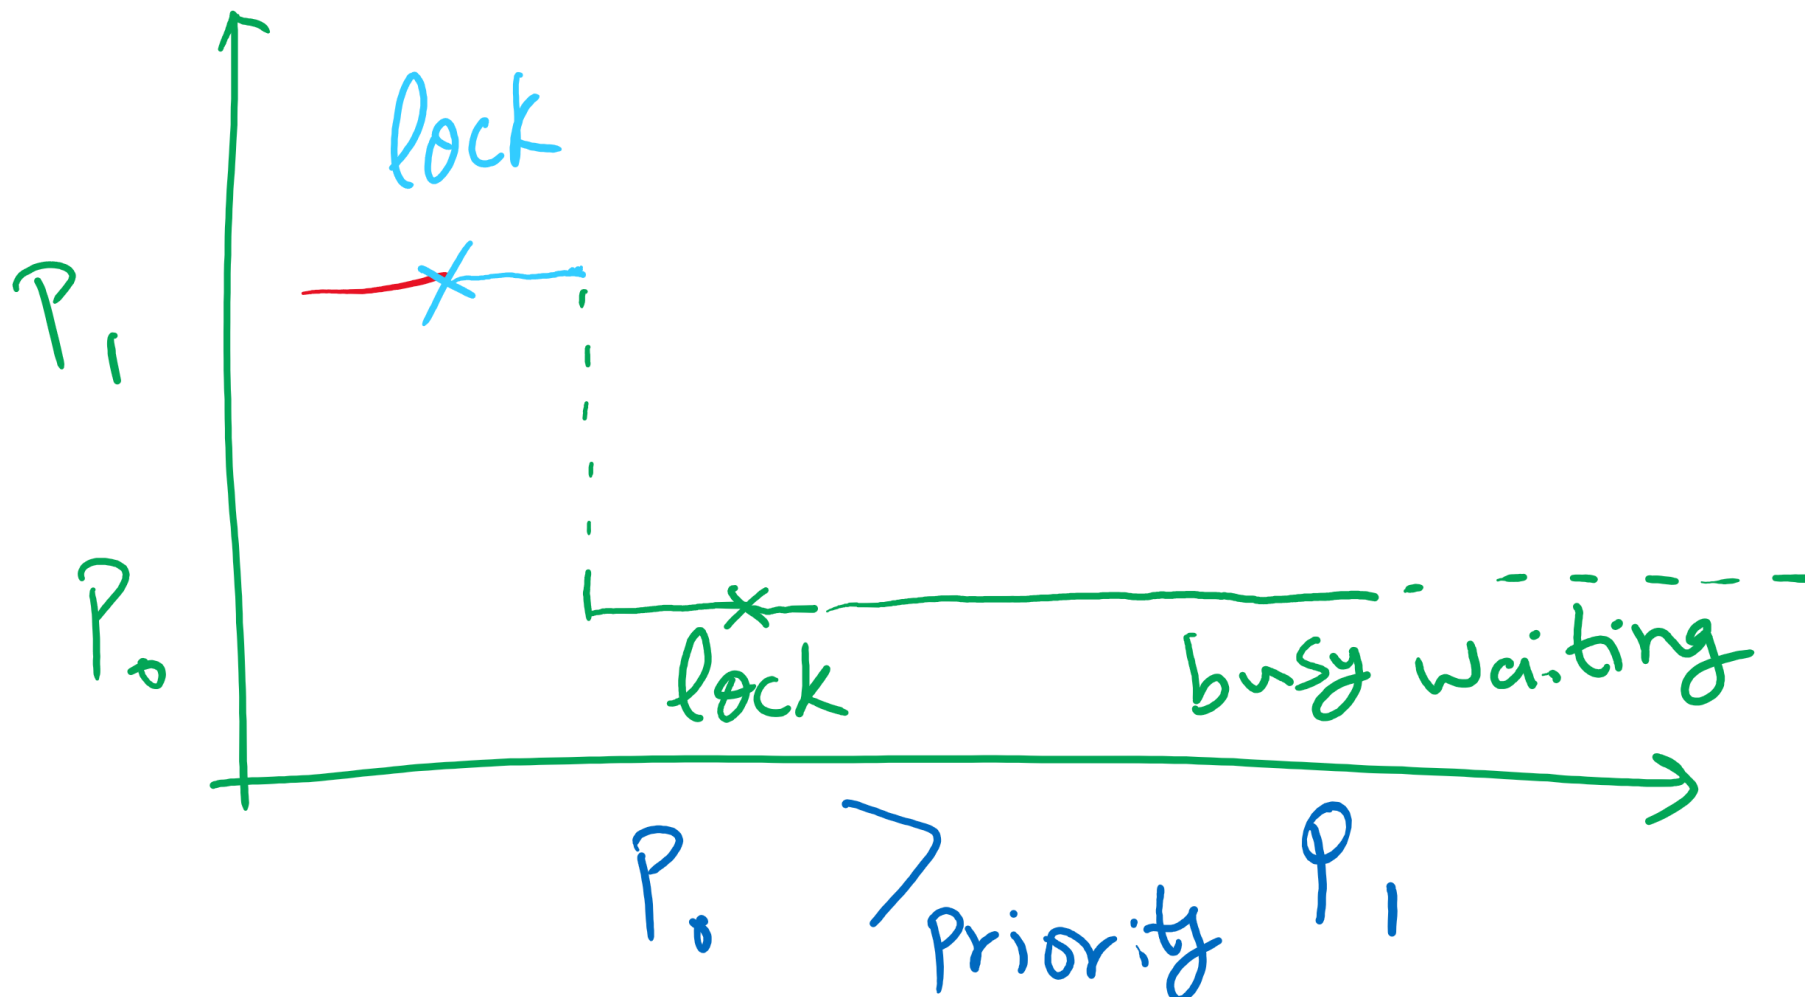
```
int lock = 0;

Lock() {
    while (Swap(&lock, 1))
        ;
}

Unlock() {
    lock = 0;
}
```

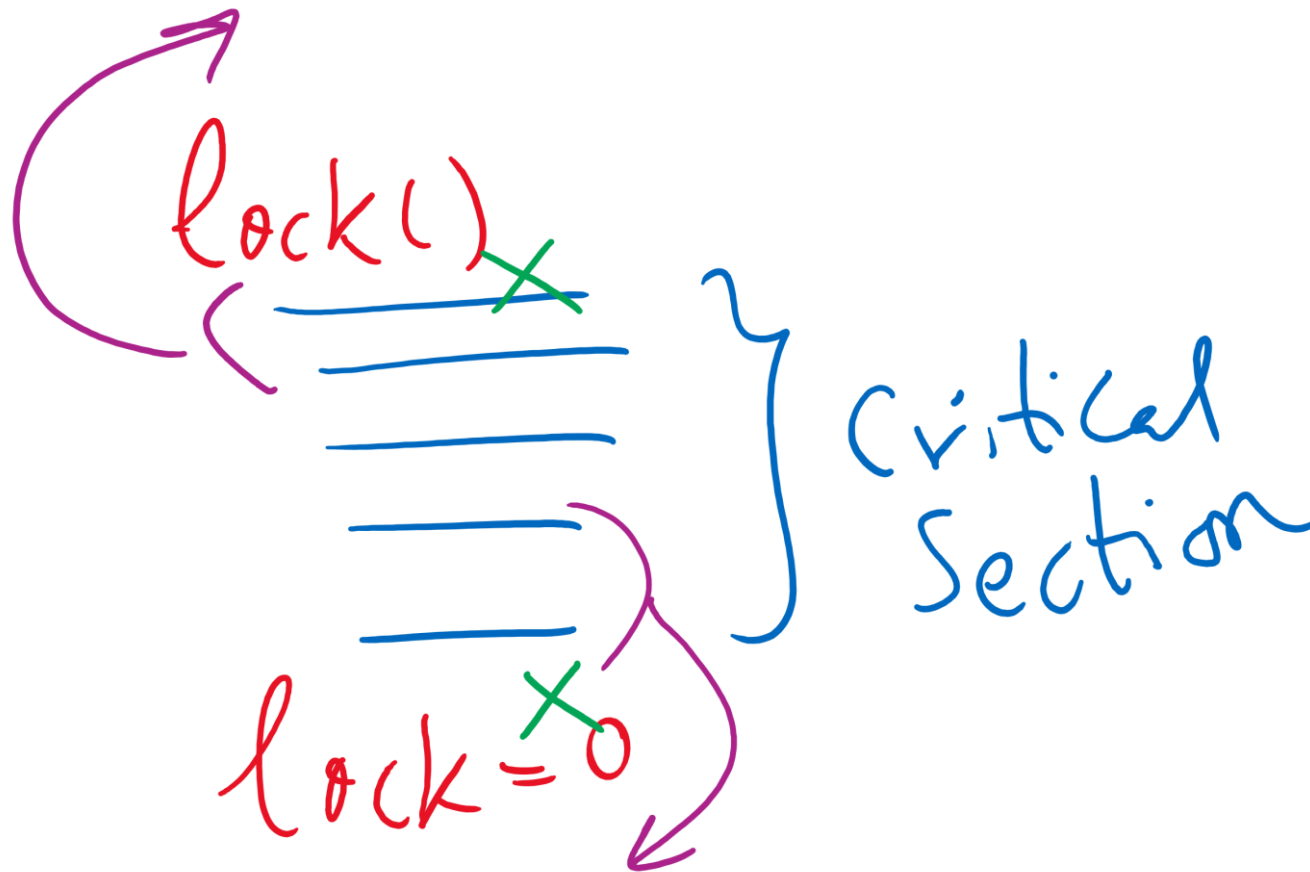
Implication of Busy Waiting

Subproblem: *priority inversion* (higher priority process busy waits for lower priority process)



Implementation Detail

compiler and/or hardware may reorder instructions



Xv6 Walkthrough

- Spinlocks
 - `__sync_synchronize()` is a memory barrier instruction

```
// Mutual exclusion lock.
struct spinlock {
    uint locked;      // Is the lock held?

};
```

```
void
initlock(struct spinlock *lk, char *name)
{

    lk->locked = 0;

}
```

```
void
acquire(struct spinlock *lk)
{

    // The xchg is atomic.
    while(xchg(&lk->locked, 1) != 0)
        ;

    // Tell the C compiler and the processor to not move loads or stores
    // past this point, to ensure that the critical section's memory
    // references happen after the lock is acquired.
    __sync_synchronize();

}
```

```
void
release(struct spinlock *lk)
{

    __sync_synchronize();


    // Release the lock, equivalent to lk->locked = 0.
    // This code can't use a C assignment, since it might
    // not be atomic. A real OS would use C atomics here.
    asm volatile("movl $0, %0" : "+m" (lk->locked) : );

}
```

Semaphores

- Solution: use semaphores
 - Synchronization mechanism that doesn't require busy waiting
- Implementation
 - Semaphore S accessed by two atomic operations
 - Down(S): decrement the semaphore if > 0 ; block otherwise
 - Up(S): increment the semaphore and wakeup one blocked process if any
 - Down() is another name for P()
 - Up() is another name for V()

Busy waiting vs. Blocking



Blocking involves 2 context switches

Critical sections using semaphores

Shared variables

```
Semaphore sem(1);
```

Code for process P_i

```
while (1) {  
    // non-critical section  
    down(sem);  
    // critical section  
    up(sem);  
    // non-critical section  
}
```

Types of semaphores

- Two different types of semaphores
 - Counting semaphores
 - Binary semaphores
- Counting semaphore
 - Value can range over an unrestricted range
- Binary semaphore
 - Only two values possible
 - 1 means the semaphore is available
 - 0 means a process has acquired the semaphore
 - May be simpler to implement
- Possible to implement one type using the other

Semaphore Implementation

But how do semaphores avoid busy waiting?

Implementing semaphores with blocking

- Assume two operations:
 - Sleep(): suspends current process
 - Wakeup(P): allows process P to resume execution
- Semaphore data structure
 - Tracks value of semaphore
 - Keeps a list of processes waiting for the semaphore

```
struct Semaphore {  
    int value;  
    ProcessList pl;  
};
```

```
down ()  
{  
    value -= 1;  
    if (value < 0) {  
        // add this process to pl  
        Sleep ();  
    }  
}  
  
up () {  
    Process P;  
    value += 1;  
    if (value <= 0) {  
        // remove a process P  
        // from pl  
        Wakeup (P);  
    }  
}
```

How to protect these shared variables??

Spinlocks in Semaphores

```
struct Semaphore {  
    int value;  
    ProcessList pl;  
  
};
```

```
down ()  
{  
  
    value -= 1;  
    if (value < 0) {  
        // add this process to pl  
        Sleep ();  
    }  
  
}  
  
up () {  
    Process P;  
    value += 1;  
    if (value <= 0) {  
        // remove a process P  
        // from pl  
        Wakeup (P);  
    }  
}
```

Spinlocks are sometimes better than Semaphores

- Very (very) short waiting time to enter the critical section $<$ the 2 context switches needed for blocking
 - Multi-core
 - so that the spinlock can be unlocked while the process is busy waiting
 - Few contending processes for the critical section
 - Short critical section code