



University of
Pittsburgh

Introduction to Operating Systems CS 1550



Spring 2023
Sherif Khattab
ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013**)

Announcements

- Upcoming deadlines
 - Homework 2 is due this Friday at 11:59 pm
 - Lab 1 is due on Tuesday 2/7 at 11:59 pm
 - Project 1 is due on Friday 2/17 at 11:59 pm
- Student Support Hours available on the syllabus page

Previous Lecture ...

- Critical Region
- Using Spinlock to implement Critical Regions
 - Busy Waiting problem

Xv6 Walkthrough of Spinlock Implementation

`__sync_synchronize()` is a memory barrier instruction

```
void
release(struct spinlock *lk)
{
    __sync_synchronize();

    // Release the lock, equivalent to lk->locked = 0.
    // This code can't use a C assignment, since it might
    // not be atomic. A real OS would use C atomics here.
    asm volatile("movl $0, %0" : "+m" (lk->locked) : );
}
```

Semaphores

- Solution to the Busy Waiting problem: use semaphores
 - Semaphores are Synchronization mechanism that doesn't require busy waiting
- Implementation
 - Semaphore S accessed by two atomic operations
 - Down(S): decrement the semaphore if > 0 ; *block* otherwise
 - Up(S): increment the semaphore and *wakeup* one blocked process if any
 - Down() is another name for P()
 - also called wait
 - Up() is another name for V()
 - also called signal

Busy waiting vs. Blocking



Blocking involves 2 context switches

Critical sections using semaphores

Shared variables

```
Semaphore sem(1);
```

Code for process P_i

```
while (1) {  
    // non-critical section  
    down(sem);  
    // critical section  
    up(sem);  
    // non-critical section  
}
```

Semaphore Implementation

But how do semaphores avoid busy waiting?

Implementing semaphores with blocking

- Assume two operations:
 - Sleep(): suspends current process
 - Wakeup(P): allows process P to resume execution
- Semaphore data structure
 - Tracks value of semaphore
 - Keeps a list of processes waiting for the semaphore

```
struct Semaphore {  
    int value;  
    ProcessList pl;  
};
```

```
down ()  
{  
    value -= 1;  
    if (value < 0) {  
        // add this process to pl  
        Sleep ();  
    }  
}  
  
up () {  
    Process P;  
    value += 1;  
    if (value <= 0) {  
        // remove a process P  
        // from pl  
        Wakeup (P);  
    }  
}
```

How to protect these shared variables??

Spinlocks in Semaphores

```
struct Semaphore {  
    int value;  
    ProcessList pl;  
  
};
```

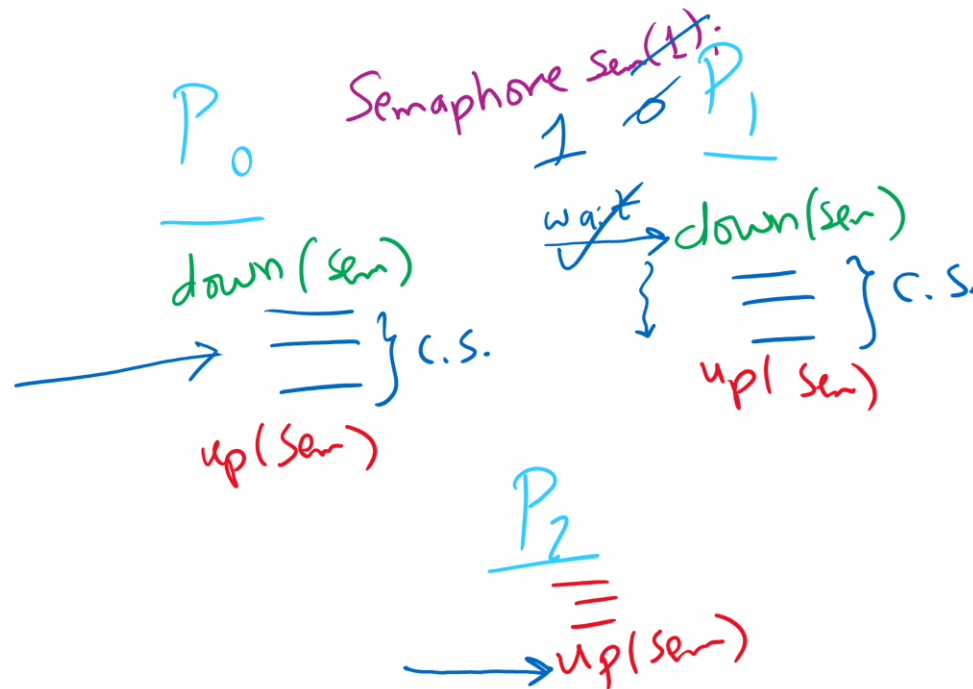
```
down ()  
{  
  
    value -= 1;  
    if (value < 0) {  
        // add this process to pl  
        Sleep ();  
    }  
  
}  
  
up () {  
    Process P;  
    value += 1;  
    if (value <= 0) {  
        // remove a process P  
        // from pl  
        Wakeup (P);  
    }  
}
```

Spinlocks are sometimes better than Semaphores

- Very (very) short waiting time to enter the critical section $<$ the 2 context switches needed for blocking
 - Multi-core
 - so that the spinlock can be unlocked while the process is busy waiting
 - Few contending processes for the critical section
 - Short critical section code

Semaphore Usage Problem: Compromising Mutual Exclusion

- Any process can up() the semaphore



- Solution: A **Mutex** can be up()'d only by the same process that down()'d it

Semaphore Usage Problem: Deadlock and Starvation

- Deadlock: two or more processes are waiting indefinitely for an event that can only be caused by a waiting process
 - P0 gets A, needs B
 - P1 gets B, needs A
 - Each process waiting for the other to signal
- Starvation: indefinite blocking
 - Process is never removed from the semaphore queue in which its suspended
 - May be caused by ordering in queues (priority)

Shared variables

```
Semaphore A (1) , B (1) ;
```

Process P₀


```
A.down () ;  
B.down () ;  
.  
.  
.  
B.up () ;  
A.up () ;
```

Process P₁

```
B.down () ;  
A.down () ;  
.  
.  
.  
A.up () ;  
B.up () ;
```

Semaphore Usage Problem: Priority Inversion

- Priority inversion is still possible using semaphores
 - Slightly less likely
 - Needs at least three processes



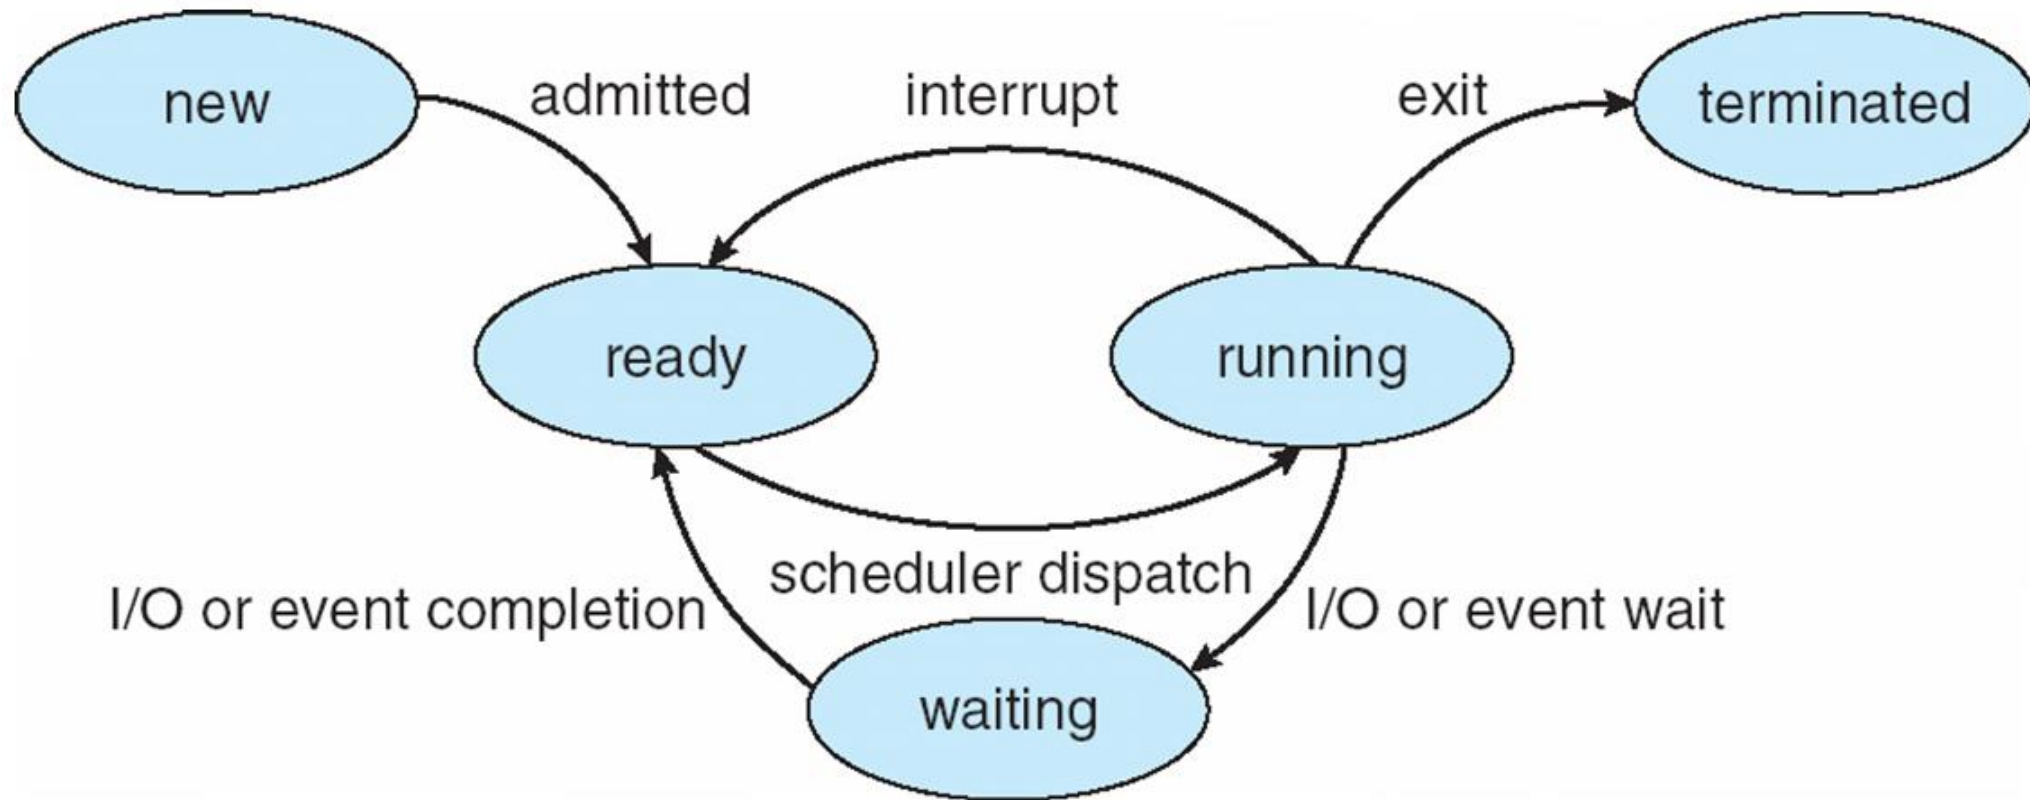
Types of semaphores

- Two different types of semaphores
 - Counting semaphores
 - Binary semaphores
- Counting semaphore
 - Value can range over an unrestricted range
- Binary semaphore
 - Only two values possible
 - 1 means the semaphore is available
 - 0 means a process has acquired the semaphore
 - May be simpler to implement
- Possible to implement one type using the other

Question

How are processes created, maintained, and terminated?

Process Lifecycle (AKA Process States)



Process Creation

- Via `fork()` syscall
- Parent process: the process that calls `fork()`
- Child process: the process that gets created
- Child process has a new context
 - new PCB

Process Creation

Memory of parent process copied to child process

- Too much copying
- Even not necessary sometimes
 - e.g., `fork()` followed by `exec()` to run a different program
- Optimization trick:
 - **copy-on-write**
 - copy when any of the two processes writes into its memory
 - copy the affected memory “part” only
 - How would the OS know when a process writes to its memory?


fork() tracing

≡
≡
≡
int ret = fork();

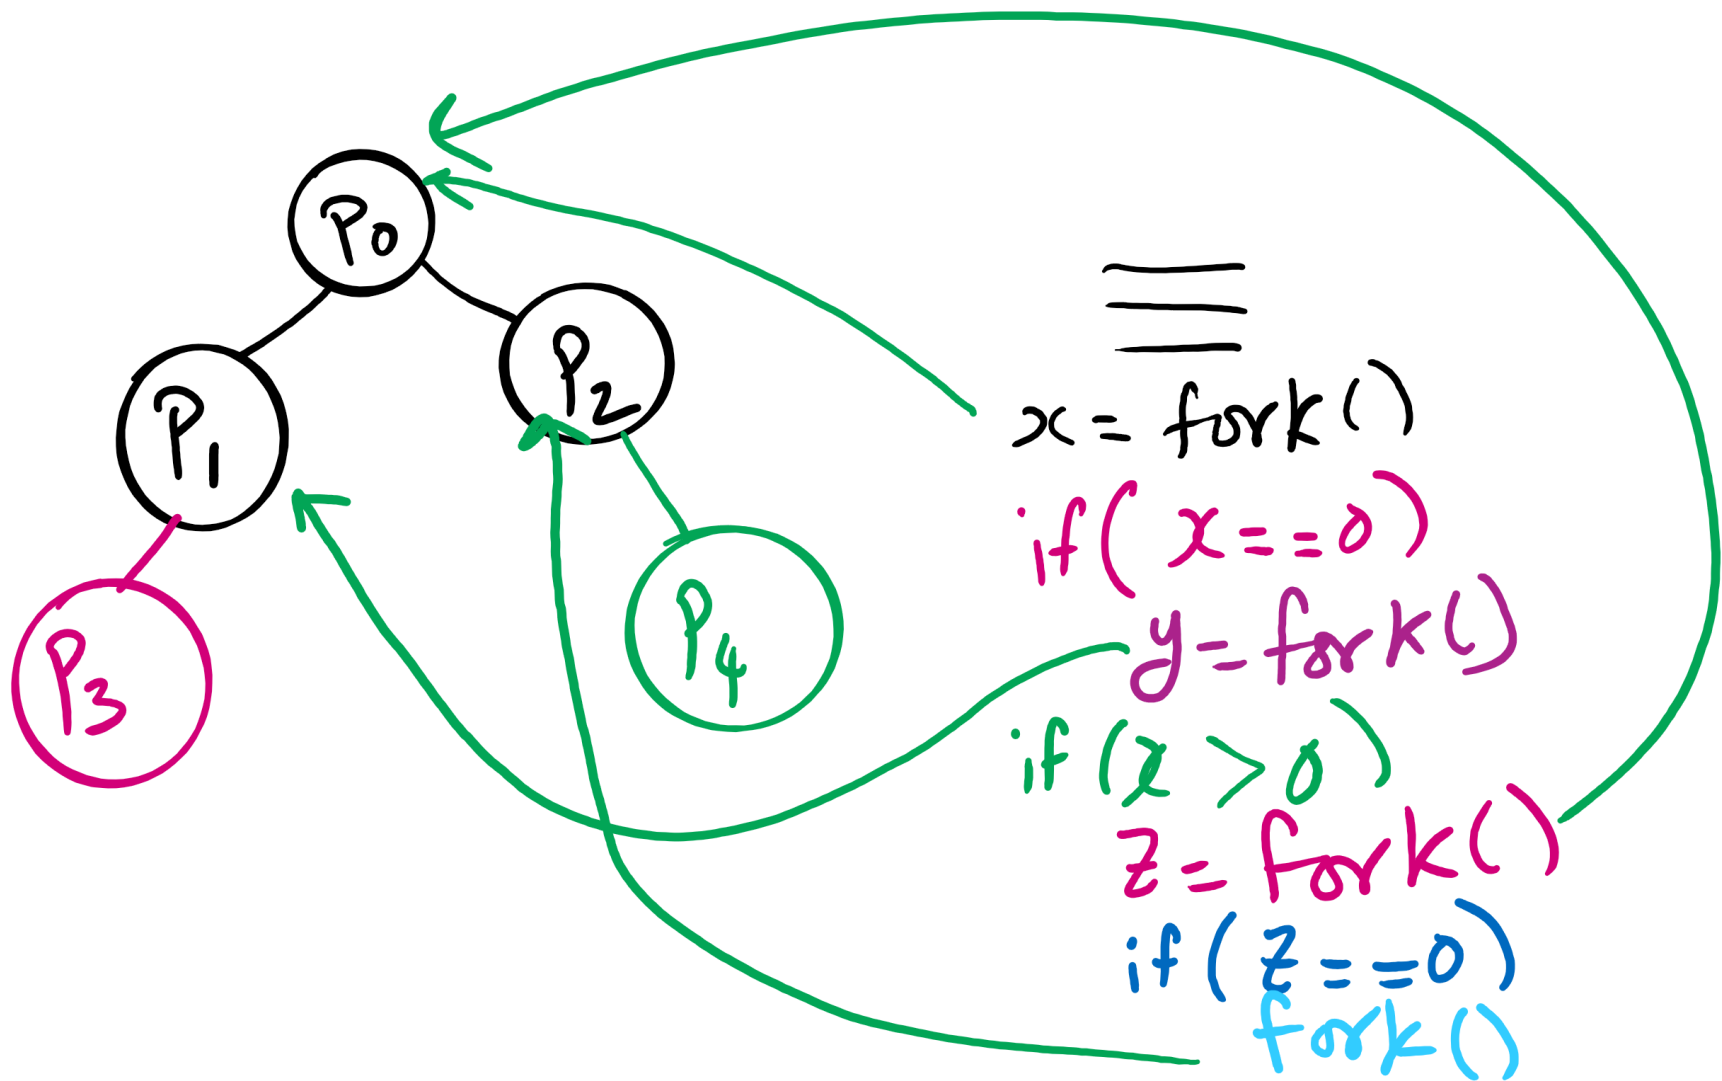
≡
≡
≡
if (ret == 0)

≡
≡
≡
if (ret > 0)

≡
≡
≡
≡



fork()'s of fork()'s



Process vs. Thread

