



University of
Pittsburgh

Introduction to Operating Systems CS 1550



Spring 2023
Sherif Khattab
ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013**)

Announcements

- Upcoming deadlines
 - Homework 6 is due **this Friday**
 - Quiz 1 and Lab 2 due on Tuesday 2/28 at 11:59 pm
 - Project 2 is due Friday 3/17 at 11:59 pm
- Midterm exam on Thursday 3/2
 - In-person, on paper, closed book
 - Study guide, old exam, and practice Midterm on Canvas
- Lost points because autograder or simple mistake?
 - please reach out to Grader TA over Piazza
- Navigating the Panopto Videos
 - Video contents
 - Search in captions

Previous lecture ...

- Sleepy Barbers solution using Condition Variables

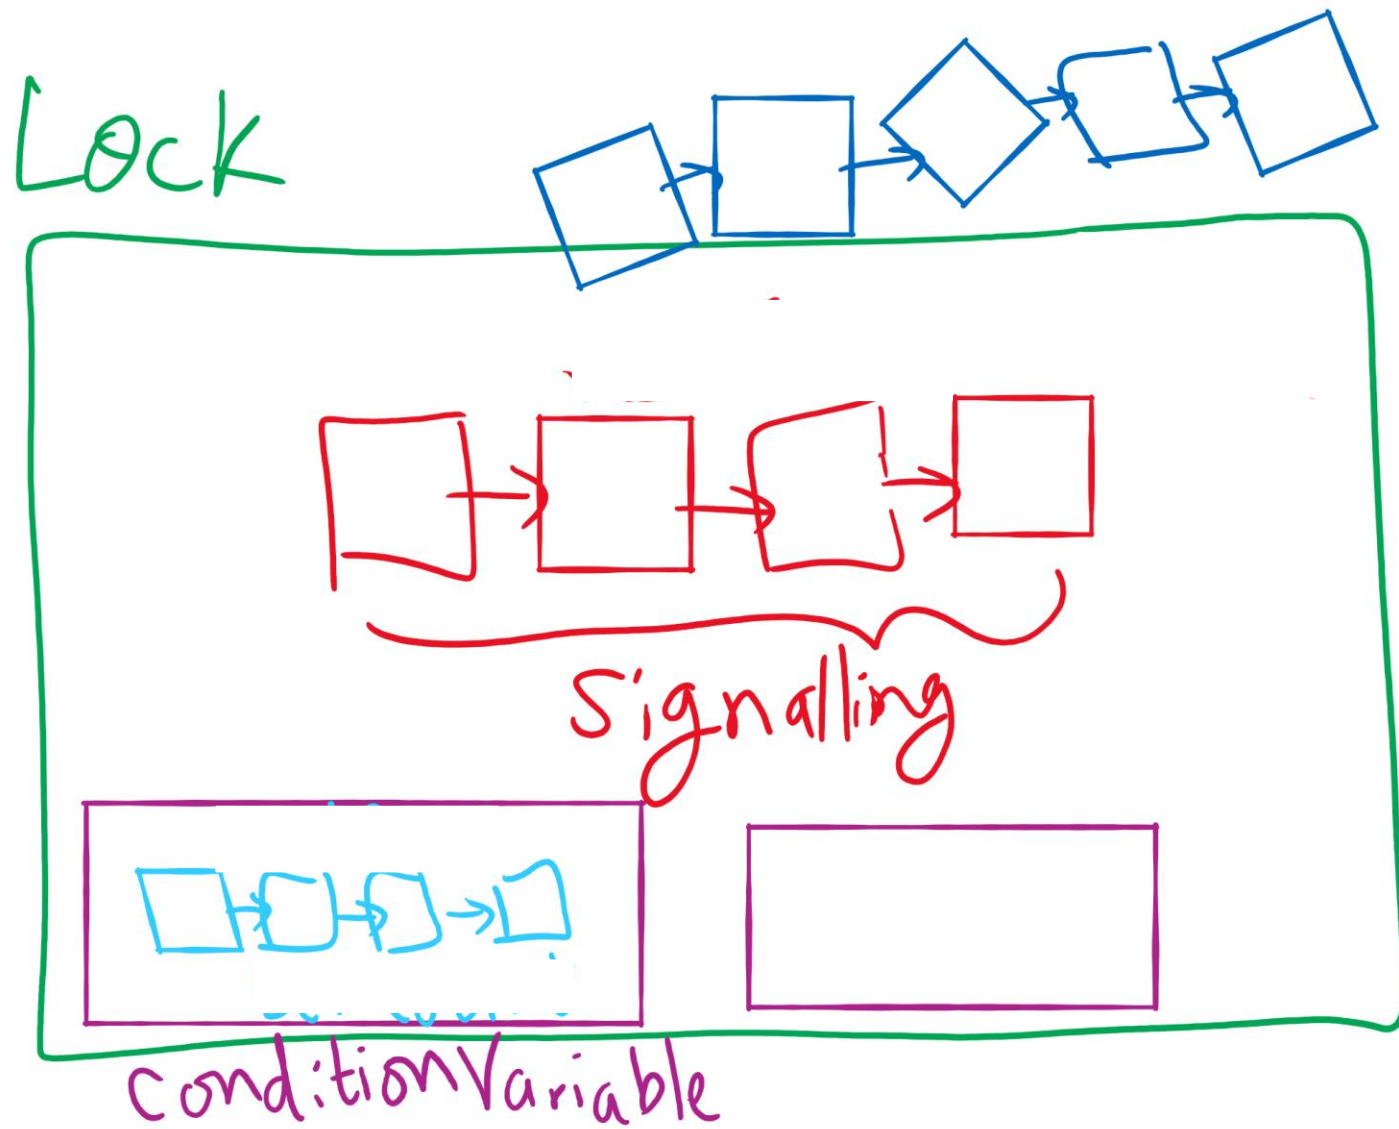
Today ...

- How to implement condition variables
- Reflections on using semaphores and condition variables
- CPU Scheduling

User-level implementation of Condition Variables

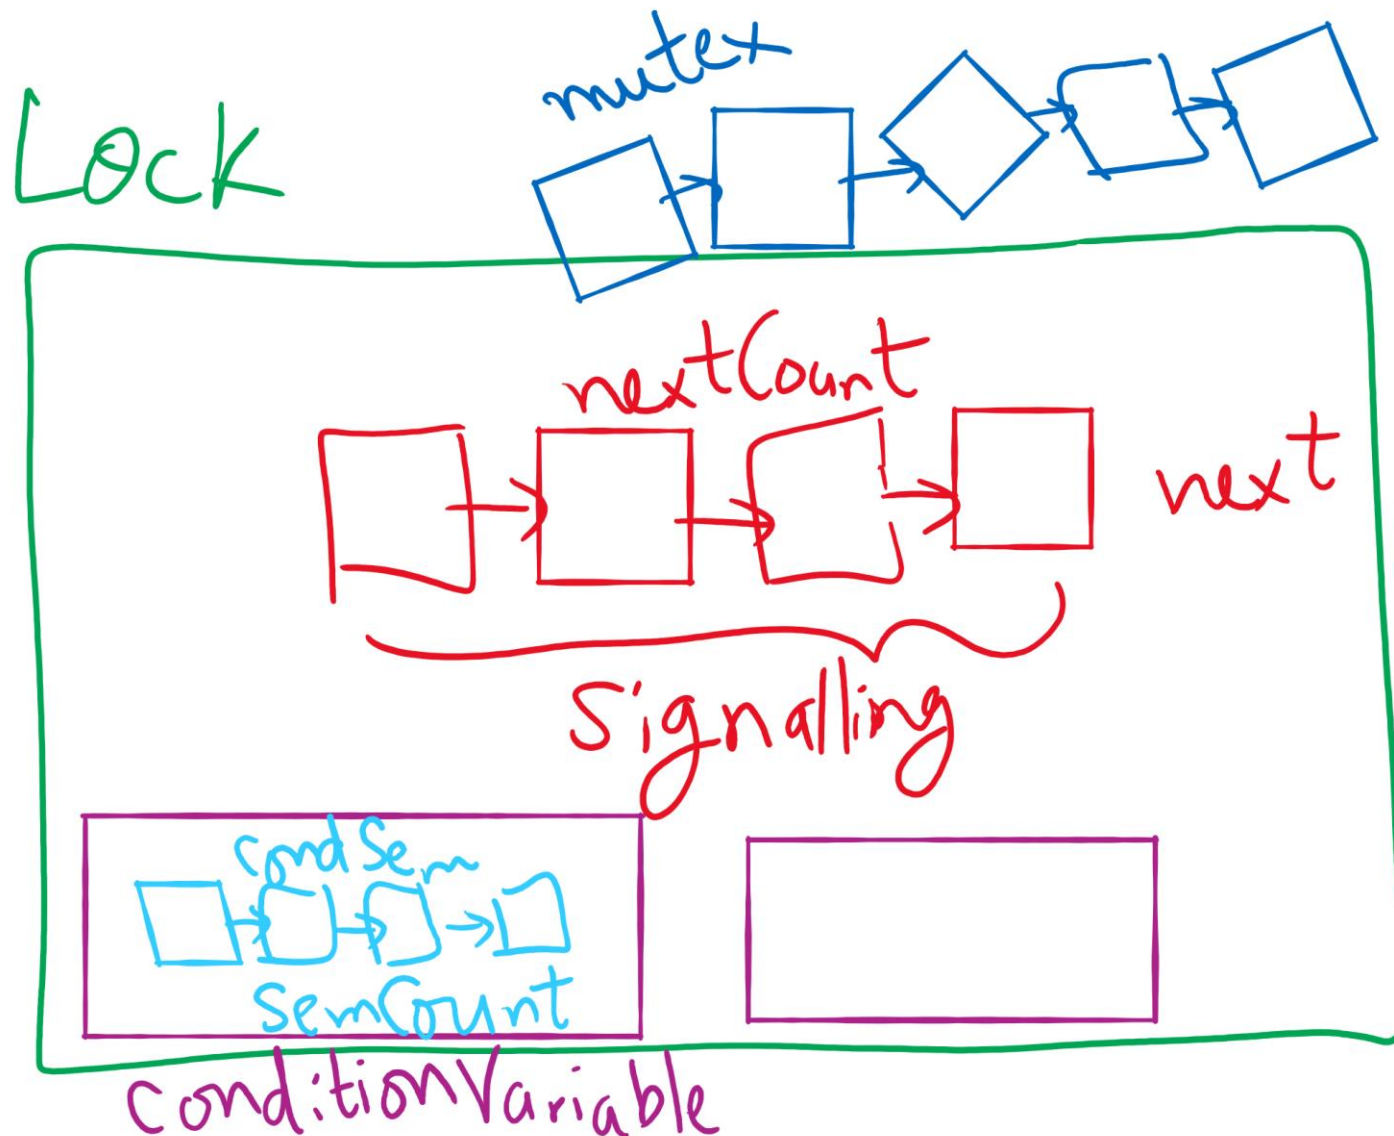
- Why?
 - Some operating systems don't have condition variables
 - Another exercise on solving synchronization problems with semaphores
- What are the waiting situations?
 - waiting to acquire the lock
 - waiting on the condition variable
 - waiting after signaling on a condition variable
 - This is Hoare semantics
 - signaling process waits
 - Compare to Mesa semantics
 - signaled process waits

Lock and Condition Variable Implementation



Lock and Condition Variable Implementation

- Let's have a semaphore for each waiting situation



User-level implementation of Condition Variables

A Lock with two waiting queues

```
struct Lock {  
    Semaphore mutex(1);  
    Semaphore next(0);  
    int nextCount = 0;  
}
```

```
Acquire(){  
    mutex.down();  
}
```

```
Release(){  
    if(nextCount > 0){  
        next.up();  
        nextCount--;  
    } else mutex.up();  
}
```


Condition Variable

```
struct ConditionVariable {  
    Semaphore condSem(0);  
    int semCount = 0;  
    Lock *lk;  
}
```

```
Wait(){
```

```
    if(lk->nextCount > 0)
```

```
        lk->next.up();
```

```
        lk->nextCount--;
```

```
    else lk->mutex.up();
```

```
    semCount++;
```

```
    condSem.down();
```

```
    semCount--;
```

```
}
```

```
Signal(){
```

```
    if(semCount > 0){
```

```
        condSem.up()
```

```
        lk->nextCount++
```

```
        lk->next.down();
```

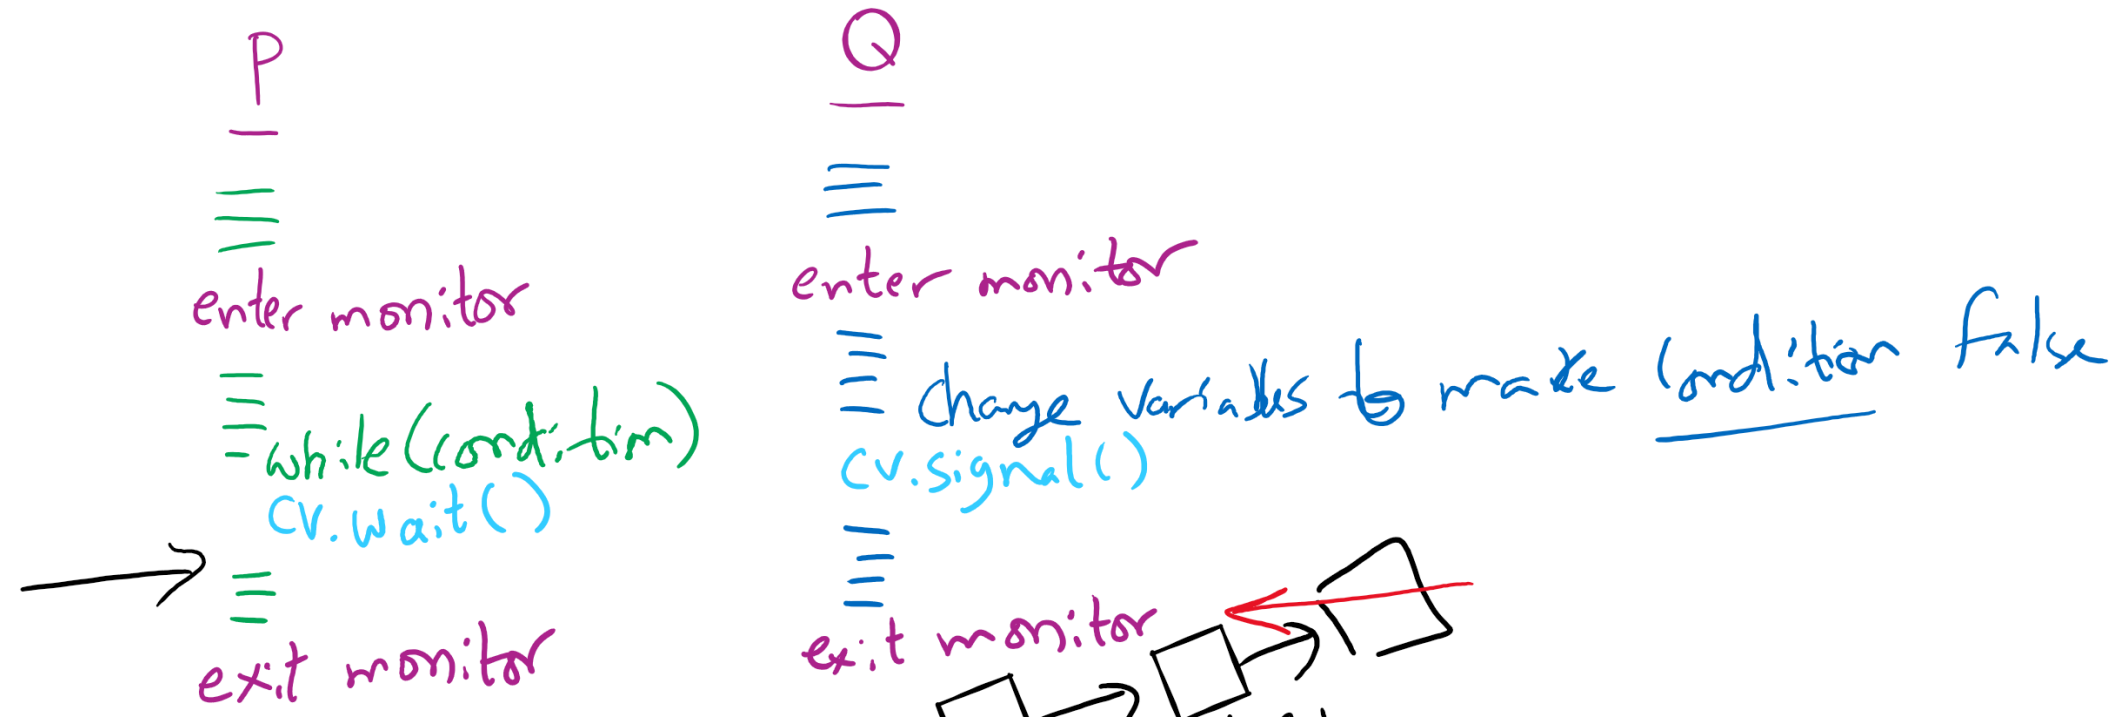
```
        lk->nextCount—
```

```
    }
```

```
}
```

Let's trace our solution!

- Note: Monitor is another name for Lock



Reflections on semaphore usage

- Semaphores can be used as
 - Resource counters
 - Waiting spaces
 - For mutual exclusion

Reflections on Condition Variables

- Define a class and put all shared variables inside the class
- Include a mutex and a condition variable in the class
- For each public method of the class
 - Start by locking the mutex lock
 - If need to wait, use a while loop and wait on the condition variable
 - Before **broadcasting** on the condition variable, make sure to change the waiting condition

Final Remarks on Process Synchronization

- Many other synchronization mechanisms
 - Message passing
 - Barriers
 - Futex
 - Re-entrant locks
 - AtomicInteger, AtomicX

Problem of the Day: CPU Scheduling

How does the ***short-term scheduler*** select the next process to run?

CPU Scheduling

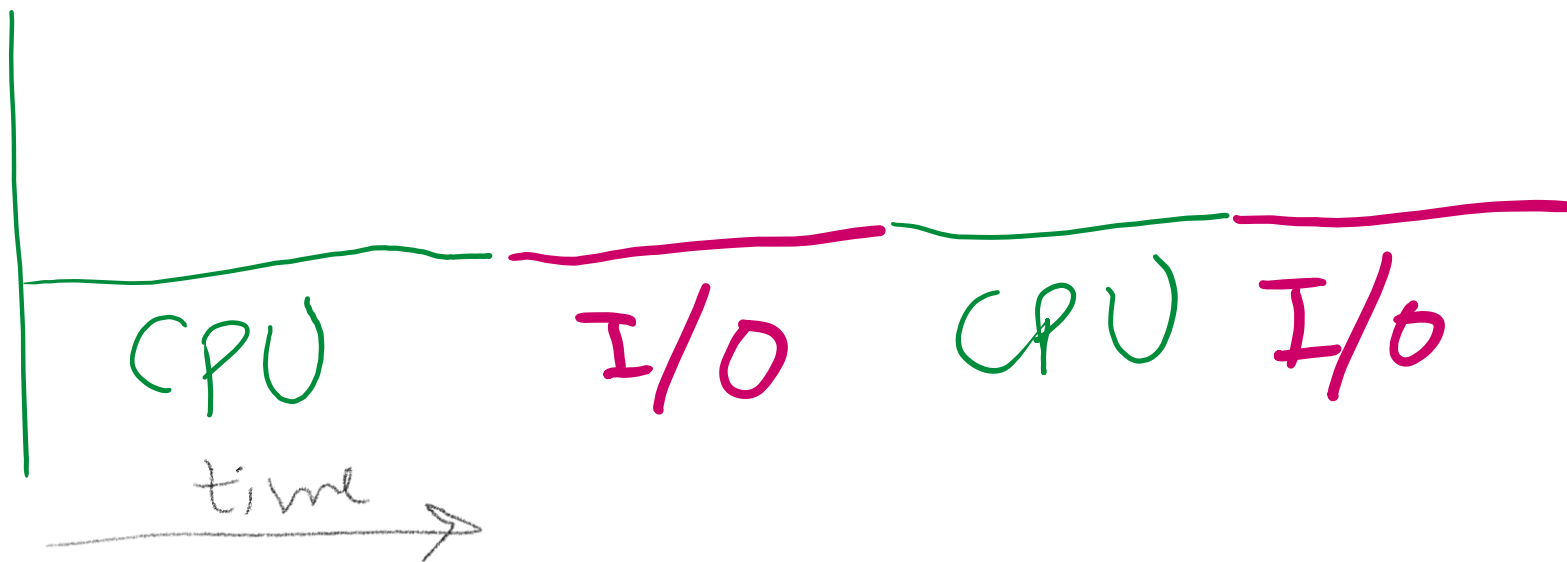
- Scheduling the processor among all ready processes
- User-oriented criteria
 - **Response Time**: Elapsed time between the submission of a request and the receipt of a response
 - **Turnaround Time**: Elapsed time between the submission of a process to its completion
- System-oriented criteria
 - Processor utilization
 - Throughput: number of process completed per unit time
 - Fairness

Short-Term Scheduler Dispatcher

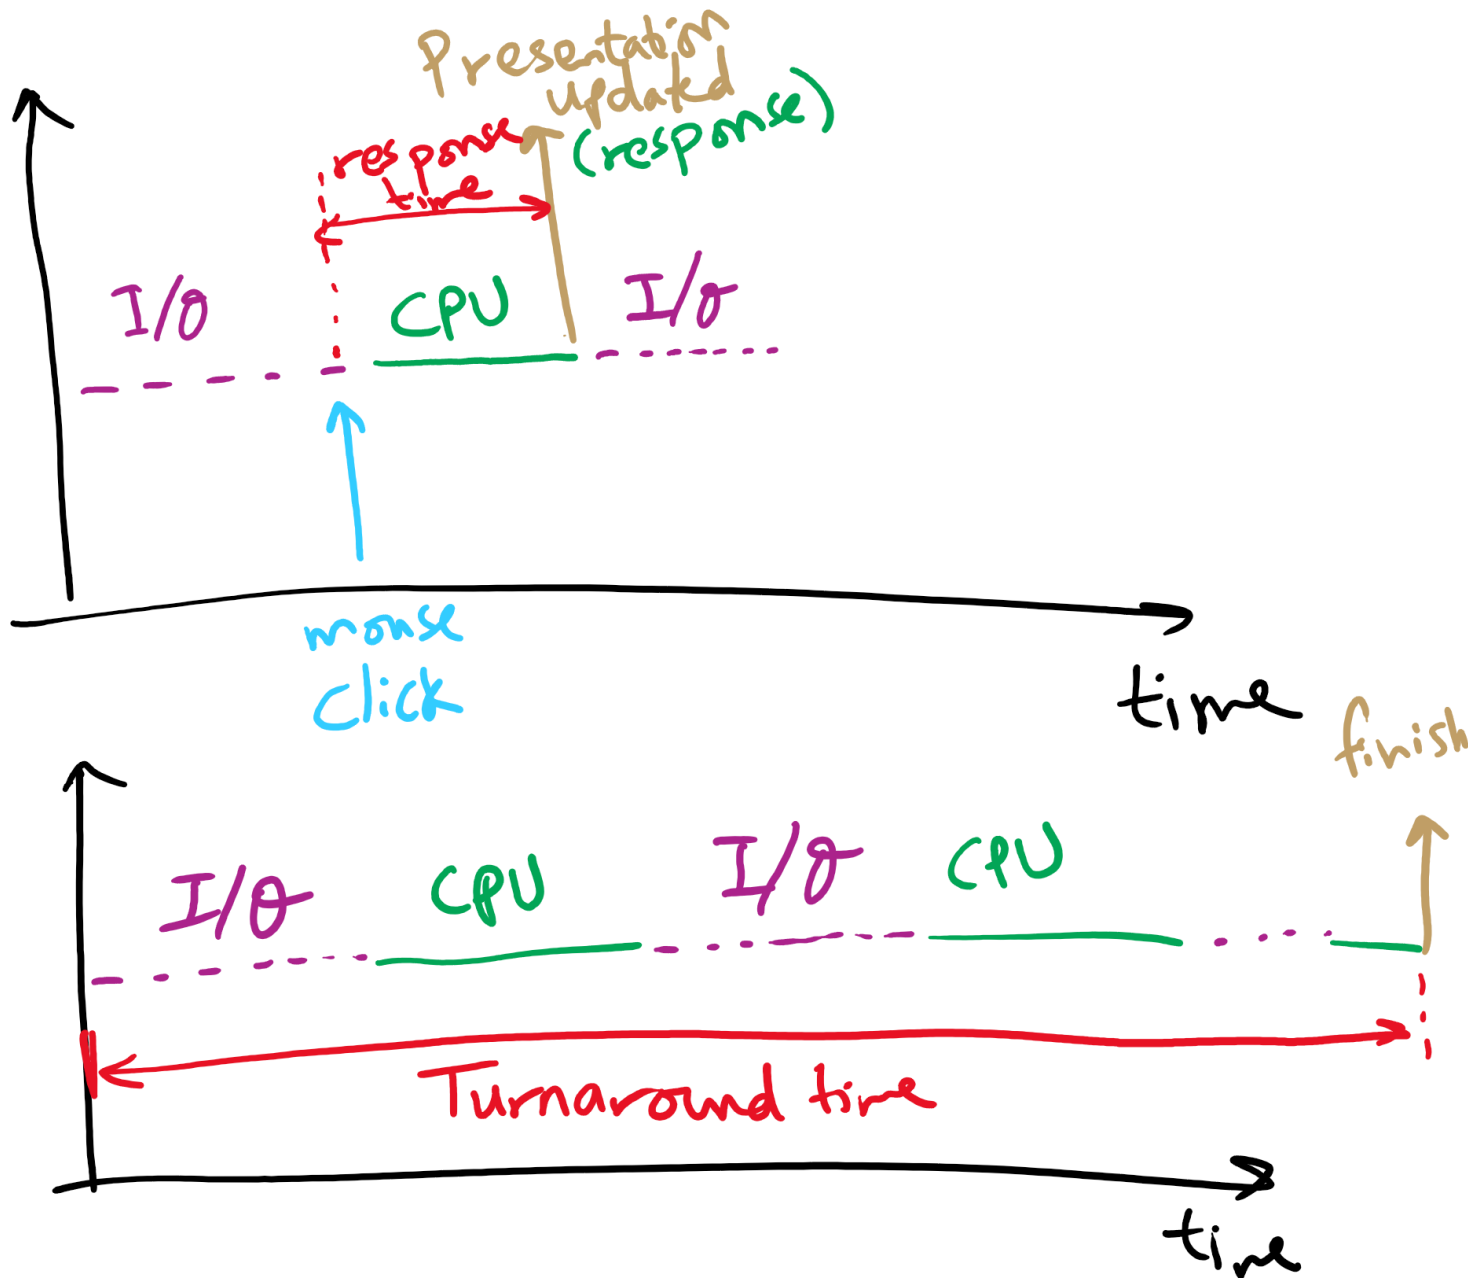
- The dispatcher is the module that gives control of the CPU to the process selected by the short-term scheduler
- The functions of the dispatcher include:
 - Switching context
 - Switching to user mode
 - Jumping to the location in the user program to restart execution
- The dispatch latency must be minimal

The CPU-I/O Cycle

- Processes require alternate use of processor and I/O in a repetitive fashion
- Each cycle consist of a CPU burst followed by an I/O burst
 - A process terminates on a CPU burst
- CPU-bound processes have longer CPU bursts than I/O-bound processes



Response time vs. Turnaround time



Scheduling Algorithms

- First-Come, First-Served Scheduling
- Shortest-Job-First Scheduling
 - Also referred to as Shortest Process Next
- Priority Scheduling
- Round-Robin Scheduling
- Multilevel Queue Scheduling
- Multilevel Feedback Queue Scheduling

Characterization of Scheduling Policies

- The **selection function** determines which ready process is selected next for execution
- The **decision mode** specifies the instants in time the selection function is exercised
 - Nonpreemptive
 - Once a process is in the running state, it will continue until it terminates or blocks for an I/O
 - Preemptive
 - Currently running process may be interrupted and moved to the Ready state by the OS
 - Prevents one process from monopolizing the processor

Process Mix Example

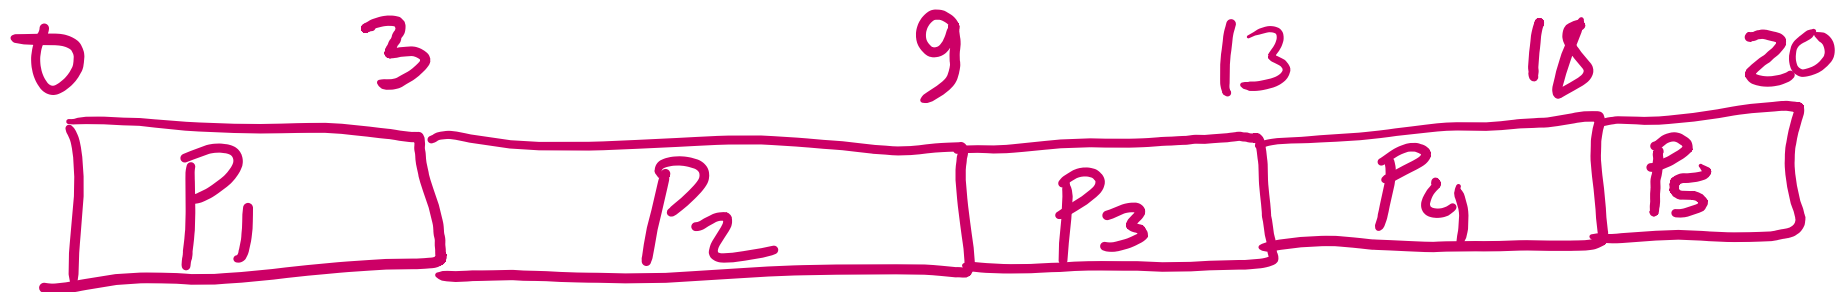
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

Service time = total processor time needed in one (CPU-I/O) cycle
Jobs with long service time are CPU-bound jobs and are referred to as “long jobs”

First Come First Served (FCFS)

- Selection function: the process that has been waiting the longest in the ready queue (hence, FCFS)
- Decision mode: non-preemptive
 - a process runs until it blocks for an I/O

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Average Response Time

$$\text{Average Response Time} = \frac{P_1 + P_2 + P_3 + P_4 + P_5}{5}$$

Handwritten calculation showing the average response time for five processes (P₁ to P₅). The numerator is the sum of the response times: (3-0) + (9-2) + (13-4) + (18-6) + (20-8). The denominator is 5.

FCFS drawbacks

- Favours CPU-bound processes
 - CPU-bound processes monopolize the processor
 - I/O-bound processes have to wait until completion of CPU-bound process
 - I/O-bound processes may have to wait even after their I/Os are completed (poor device utilization)
 - Convoy effect
- Better I/O device utilization could be achieved if I/O bound processes had higher priority

Convoy Effect

