



University of  
Pittsburgh

# Introduction to Operating Systems CS 1550



Spring 2023  
Sherif Khattab  
ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013**)

# Announcements

- Upcoming deadlines
  - Homework 8 is due **this Friday**
  - Quiz 1 and Lab 2 due on Tuesday 2/28 at 11:59 pm
  - Project 2 is due Friday 3/17 at 11:59 pm
- Talk by candidate faculty
  - This Wednesday 3/15 @ 10 am at 5317 Sennott Square
  - Donuts will be served!

# Last Lecture ...

- How to implement Condition Variables and Locks using semaphores
- CPU scheduling
  - SJF
  - Priority
  - RR

# Today ...

- CPU scheduling
  - Multi-Level Feedback Queues
  - CPU burst estimation

# Multilevel Feedback Scheduling

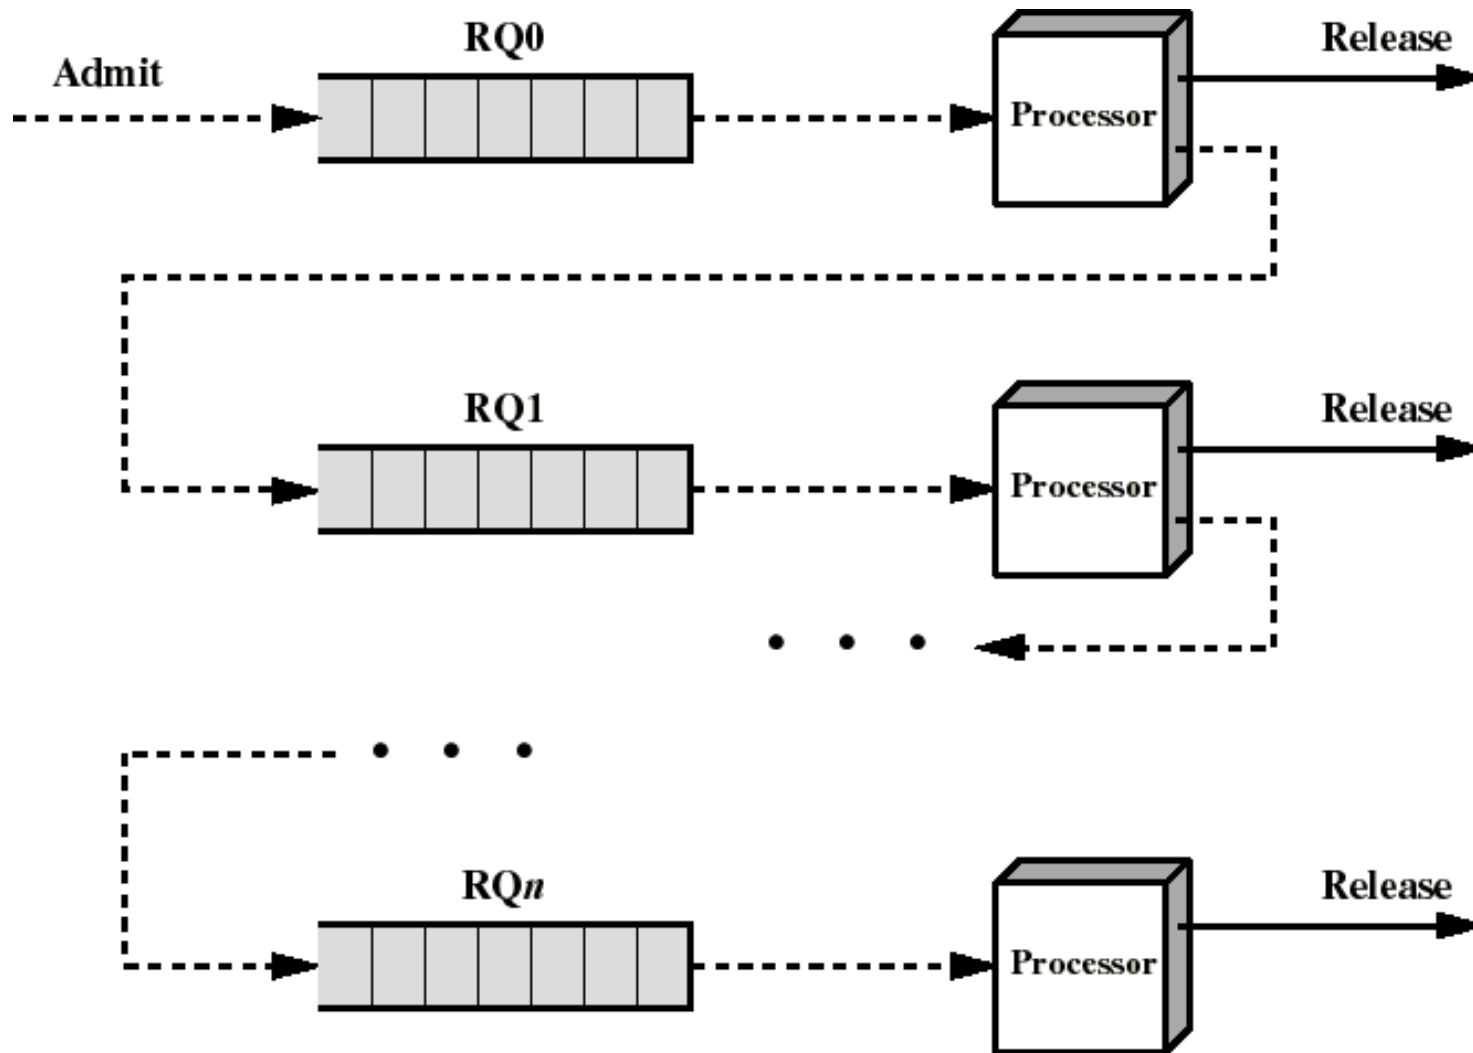
- Preemptive scheduling with dynamic priorities
- N ready to execute queues with decreasing priorities:
  - $P(RQ_0) > P(RQ_1) > \dots > P(RQ_{N-1})$
- Dispatcher selects a process for execution from  $RQ_i$  only if  $RQ_{i-1}$  to  $RQ_0$  are empty

# Multilevel Feedback Scheduling

- New process are placed in  $RQ_0$
- After the first quantum, they are moved to  $RQ_1$ , and to  $RQ_2$  after the second quantum, ... and to  $RQ_{N-1}$  after the Nth quantum
- I/O-bound processes remain in higher priority queues.
  - CPU-bound jobs drift downward.
  - Hence, long jobs may starve

# Multiple Feedback Queues

Different RQs may have different quantum values



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

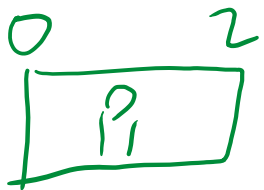
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

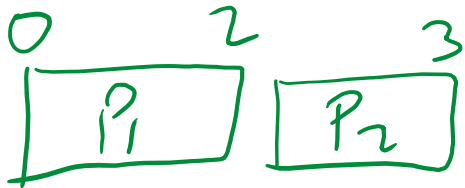
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

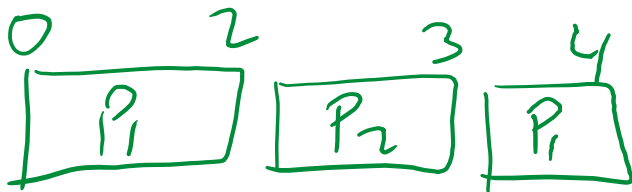
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

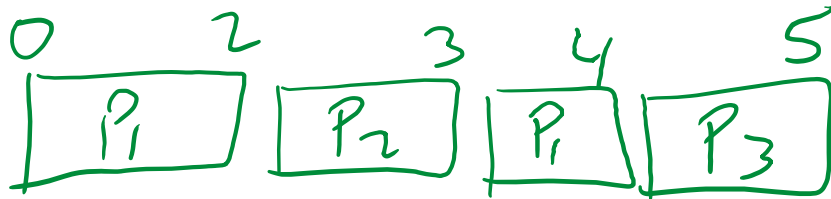
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

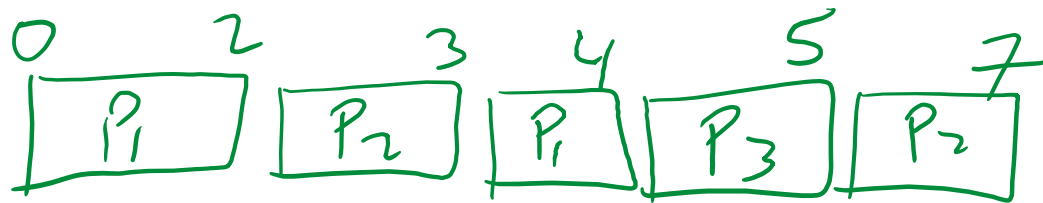
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

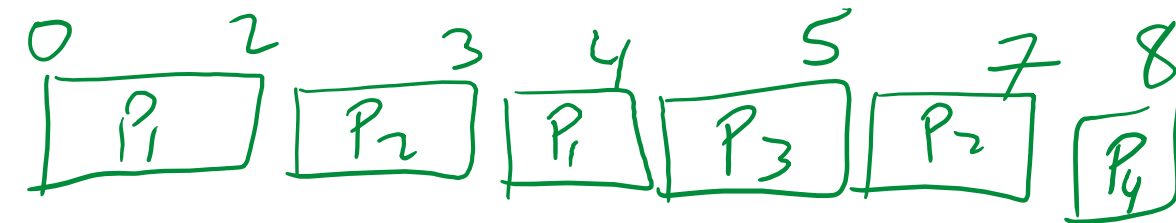
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

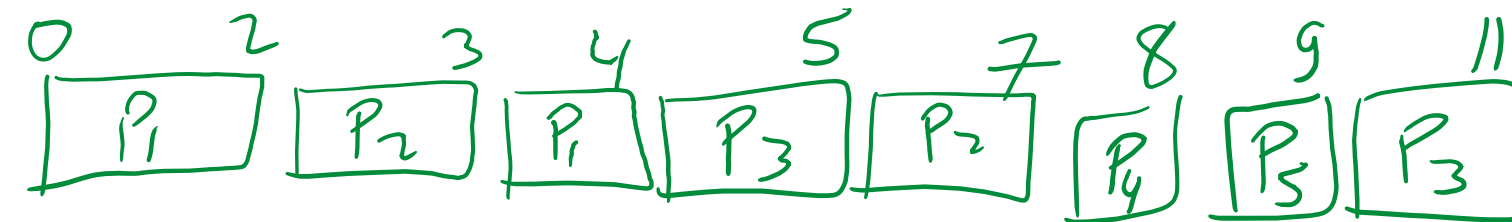
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

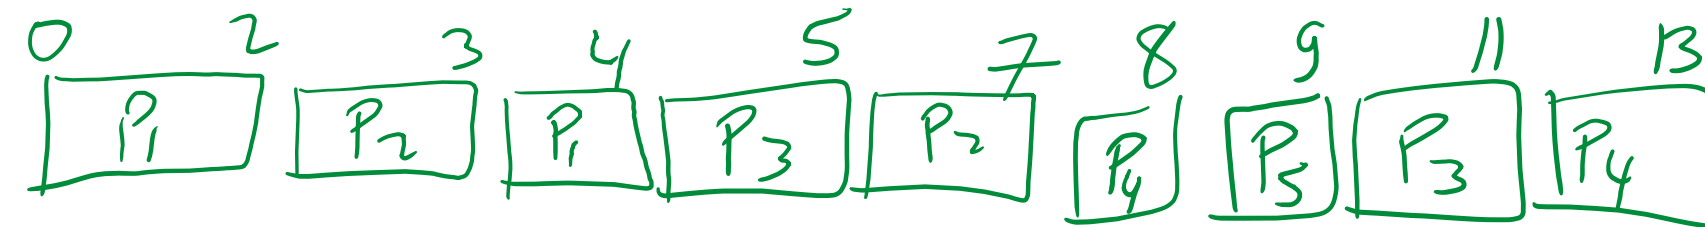




# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

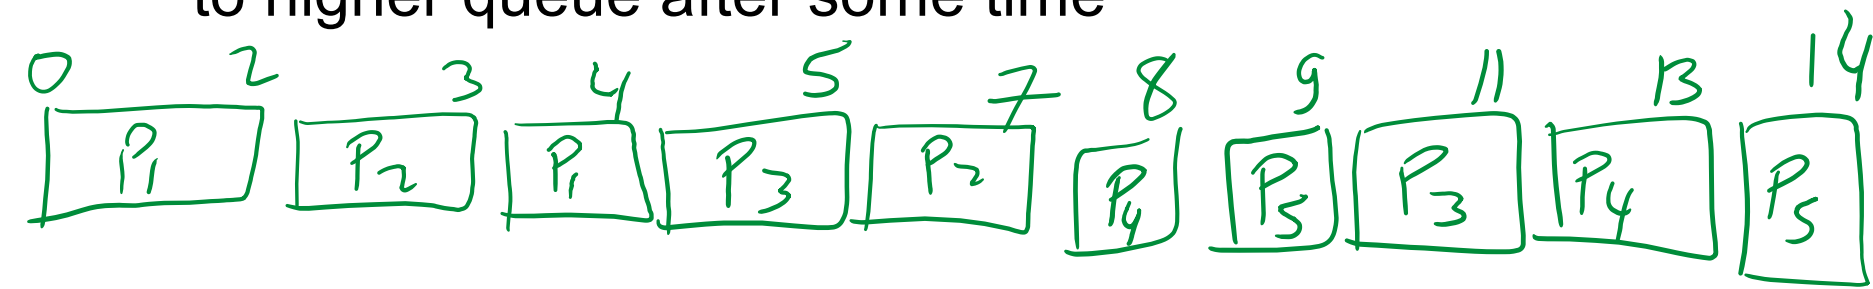
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

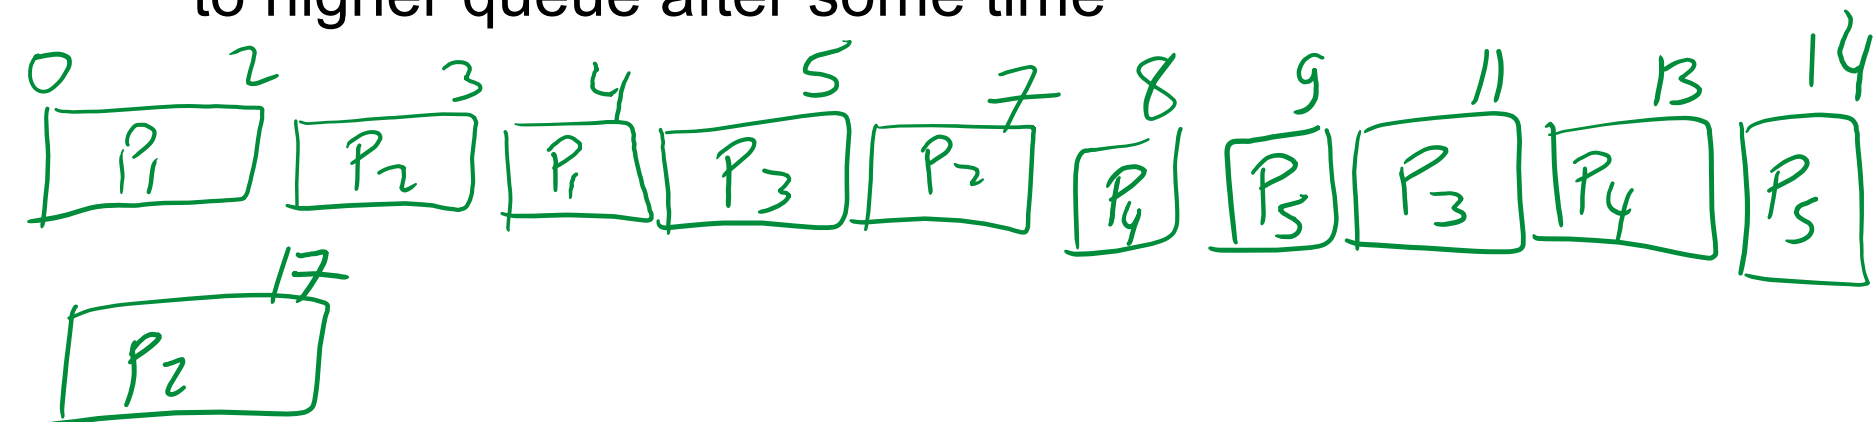
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

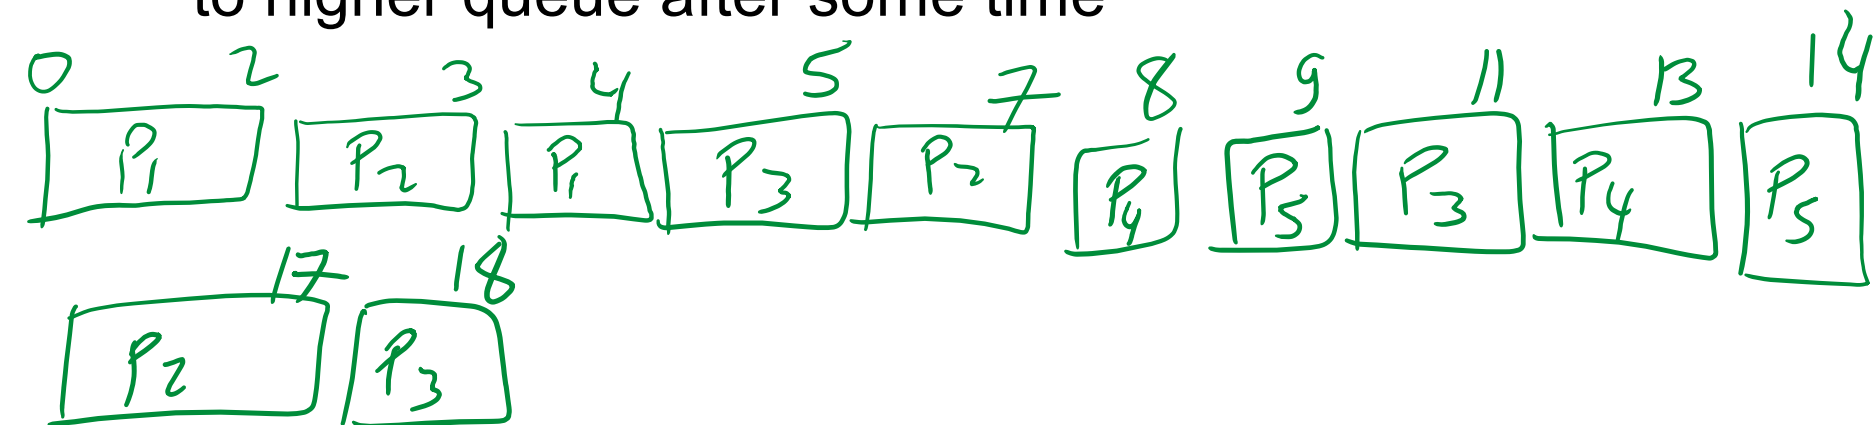
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

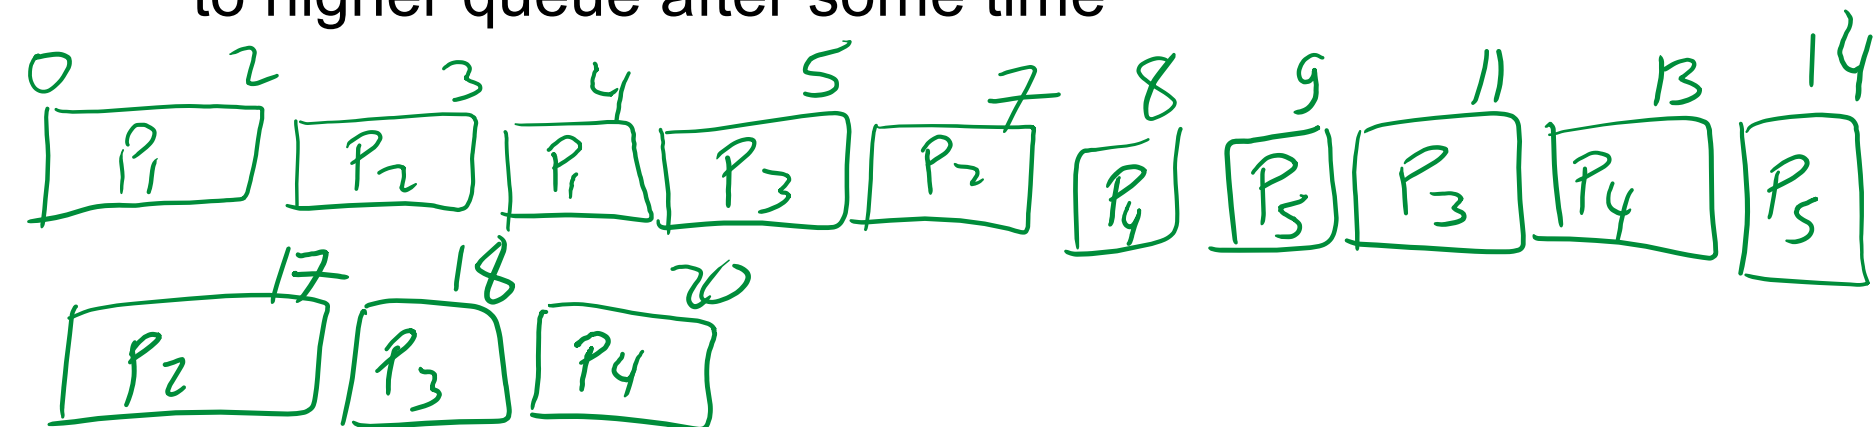
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



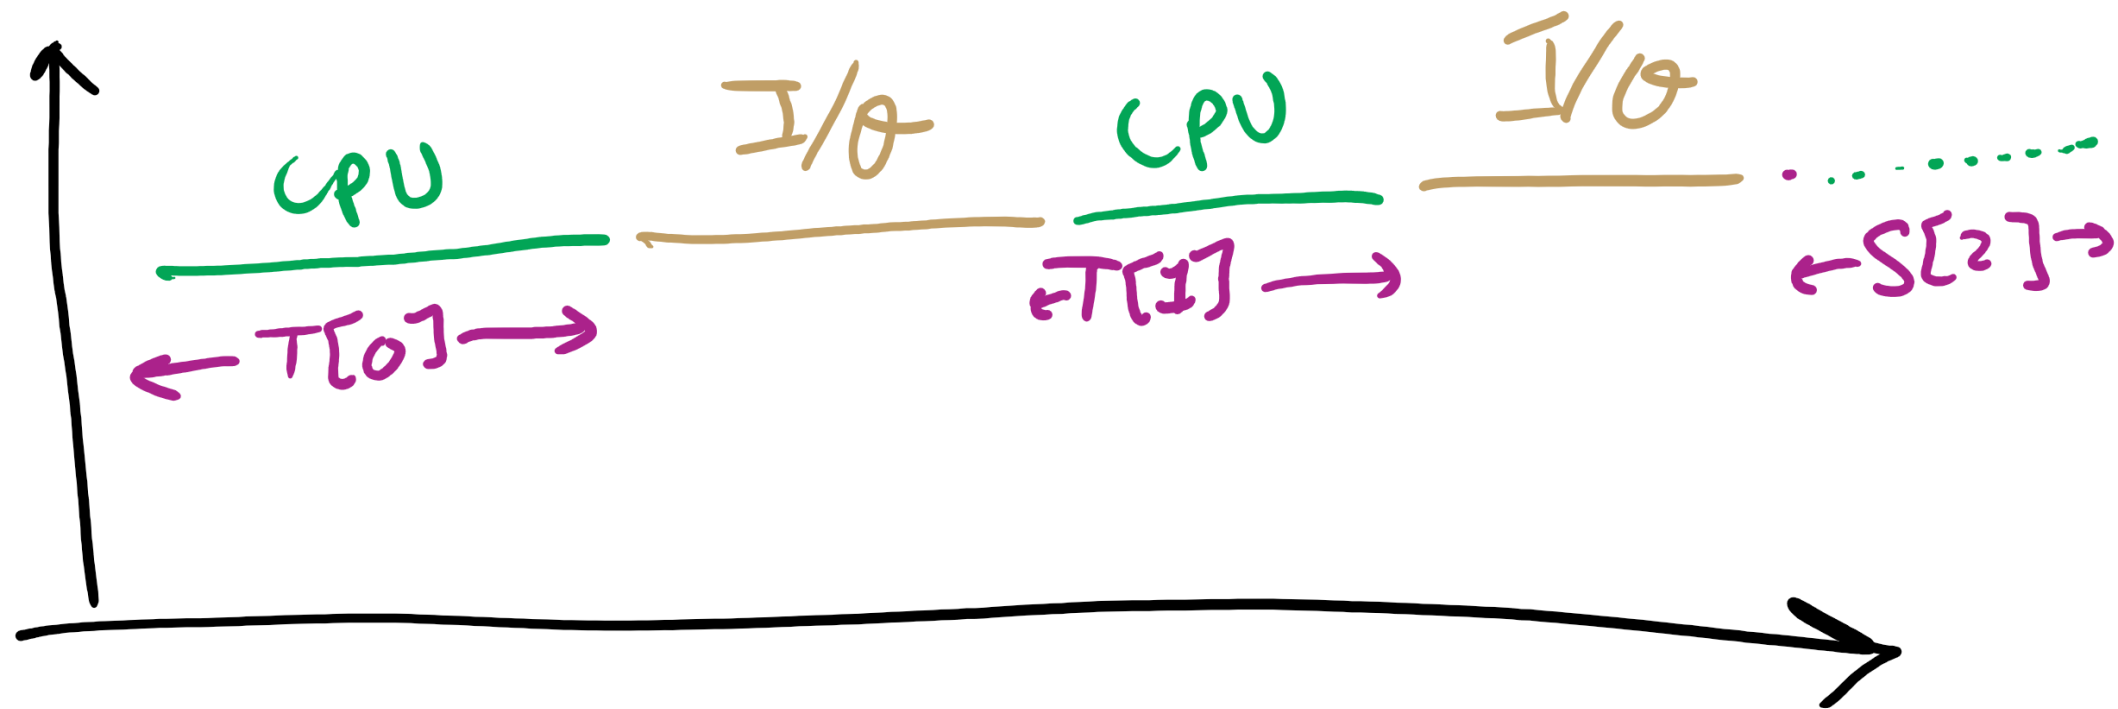
# Algorithm Comparison

- Which one is the best?
- The answer depends on many factors:
  - the system workload (extremely variable)
  - hardware support for the dispatcher
  - relative importance of performance criteria (response time, CPU utilization, throughput...)
  - The evaluation method used (each has its limitations...)

# Back to SJF: CPU Burst Estimation

- Let  $T[i]$  be the execution time for the  $i$ th instance of this process: the actual duration of the  $i$ th CPU burst of this process
- Let  $S[i]$  be the predicted value for the  $i$ th CPU burst of this process. The simplest choice is:
  - $S[n+1] = (1/n)(T[1] + \dots + T[n]) = (1/n) \sum_{\{i=1 \text{ to } n\}} T[i]$
- This can be more efficiently calculated as:
  - $S[n+1] = (1/n) T[n] + ((n-1)/n) S[n]$
- This estimate, however, results in equal weight for each instance

# CPU Burst Estimation





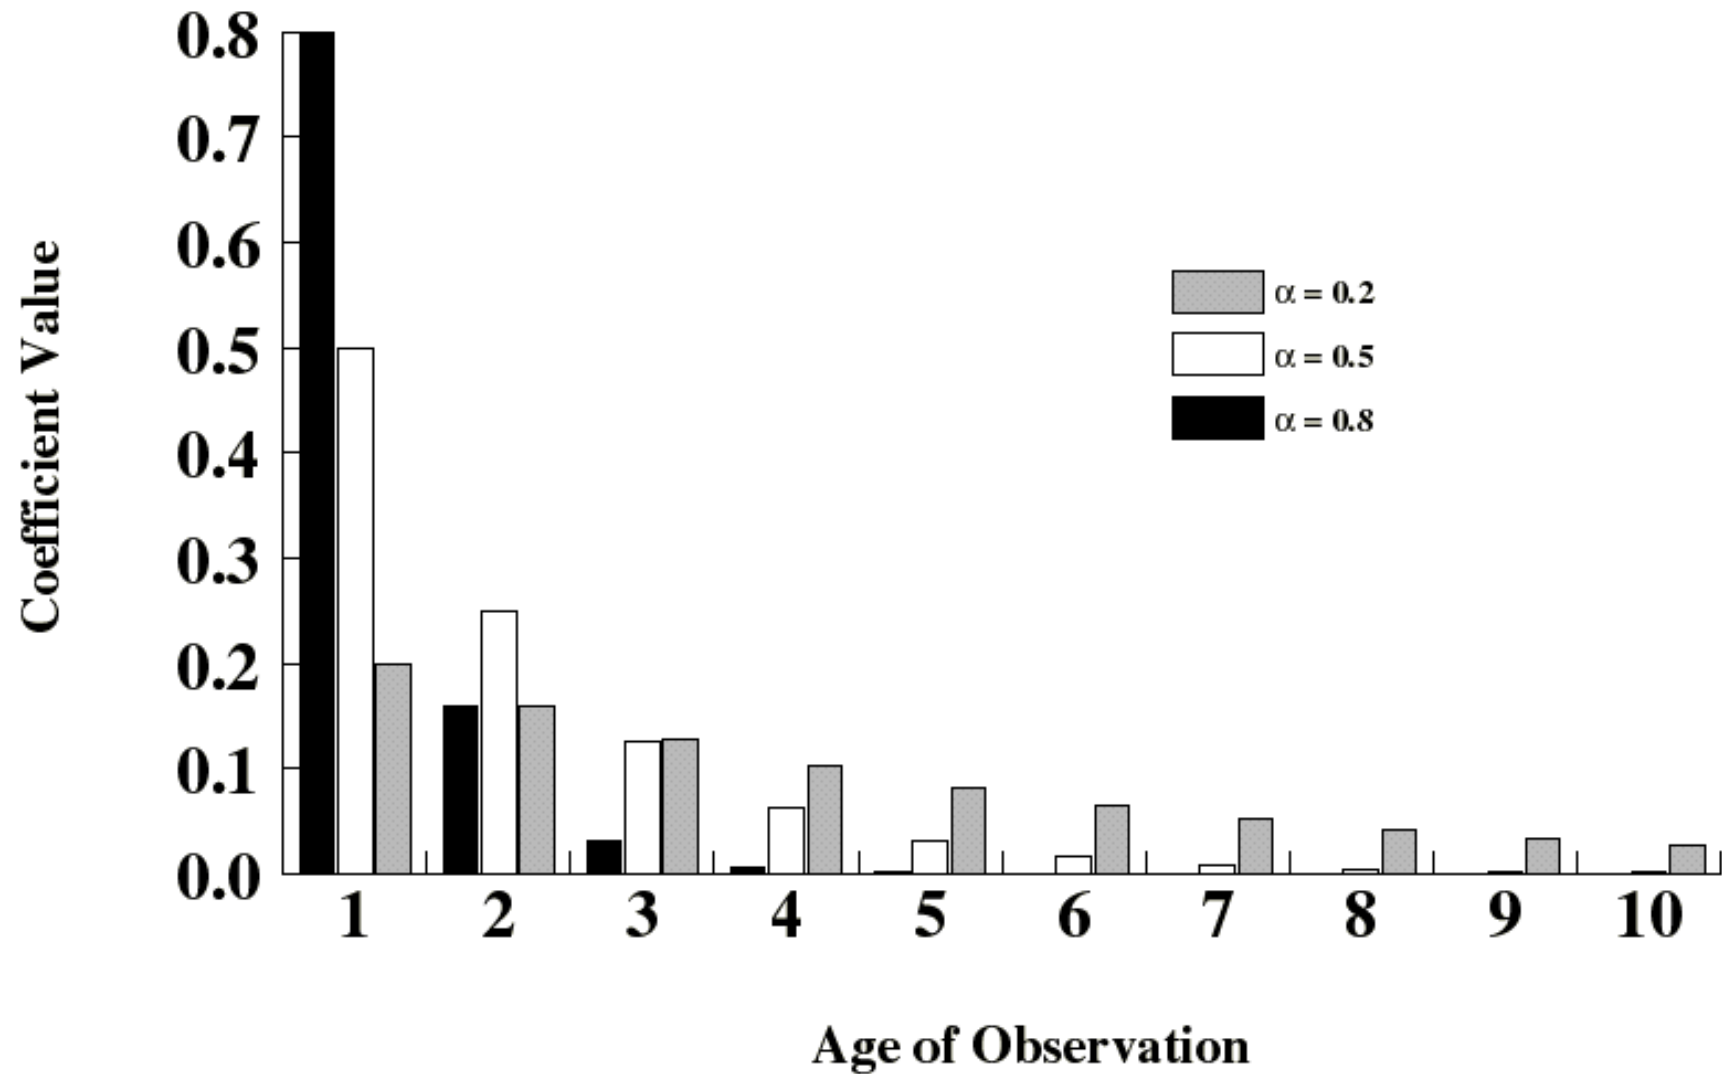
# Estimating the required CPU burst

- Recent instances are more likely to better reflect future behavior
- A common technique to factor the above observation into the estimate is to use **exponential averaging** :
  - $S[n+1] = \alpha T[n] + (1 - \alpha) S[n] ; \quad 0 < \alpha < 1$

# CPU burst Estimate (Exponential Average)

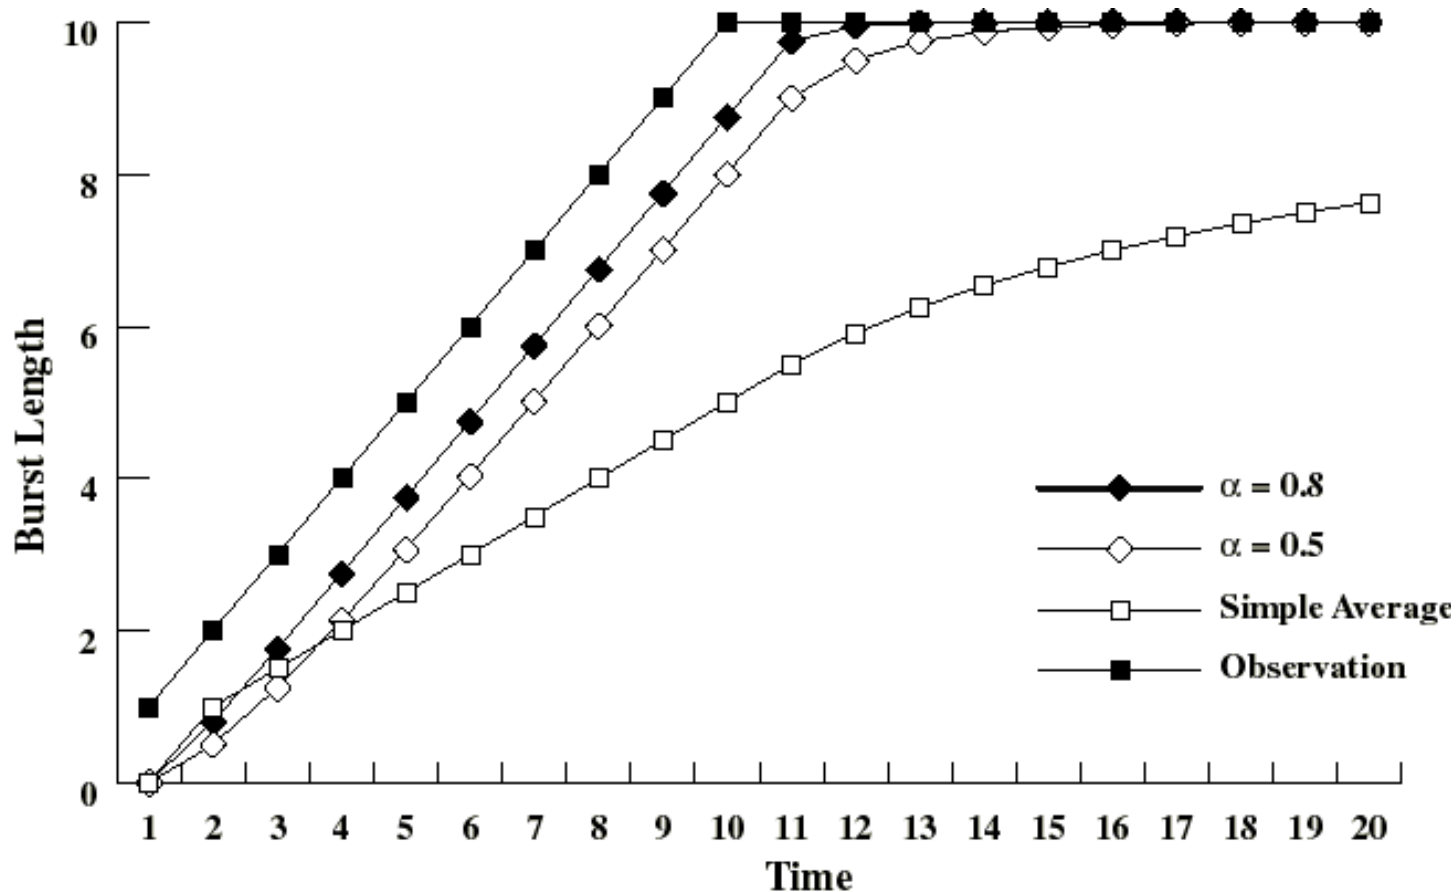
- Recent instances have higher weights, whenever  $\alpha > 1/n$
- Expanding the estimated value shows that the weights of past instances decrease exponentially
  - $$S[n+1] = \alpha T[n] + (1-\alpha) \alpha T[n-1] + \dots (1-\alpha)^{i} \alpha T[n-i] + \dots + (1-\alpha)^{n} S[1]$$
  - The predicted value of 1st instance,  $S[1]$ , is usually set to 0 to give priority to new processes

# Exponentially Decreasing Coefficients

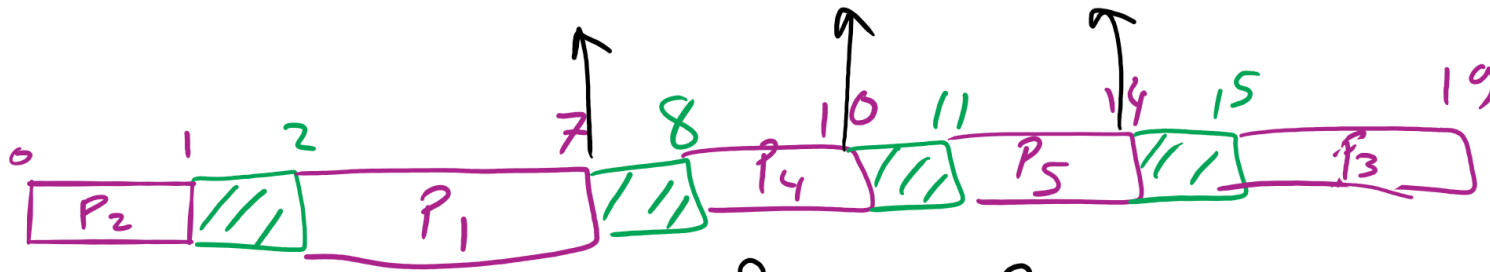


# Exponentially Decreasing Coefficients

- $S[1] = 0$  to give high priority to new processes
- Exponential averaging tracks changes in process behavior much faster than simple averaging



# FCFS Problem in HW7



$$\begin{aligned}
 \text{A.R.T.} &= \frac{\overset{P_1}{(7-2)} + \overset{P_2}{(1-0)} + \overset{P_3}{(19-8)} + \overset{P_4}{(10-4)} + \overset{P_5}{(4-7)}}{5} \\
 &= \frac{5 + 1 + 11 + 6 + 7}{5} = 6\text{ms}
 \end{aligned}$$

$$\text{CPU Util} = \frac{19-4}{19} = \frac{15}{19}$$

# CPU Burst Estimation

$$S[n+1] = \alpha T[n] + (1-\alpha) S[n]$$

$$0 < \alpha < 1$$

