



University of  
Pittsburgh

# Introduction to Operating Systems CS 1550



Spring 2023  
Sherif Khattab  
ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013**)

# Announcements

- Upcoming deadlines
  - Homework 7 is due **this Friday**
  - Quiz 1 and Lab 2 due on Tuesday 2/28 at 11:59 pm
  - Project 2 is due Friday 3/17 at 11:59 pm
- Talk by candidate faculty
  - This Wednesday 3/15 @ 10 am at 5317 Sennott Square
  - Donuts will be served!

# Last Lecture ...

- How to implement Condition Variables and Locks using semaphores
- CPU scheduling
  - SJF
  - Priority
  - RR

# Today ...

- CPU scheduling
  - Multi-Level Feedback Queues
  - CPU burst estimation
- Memory

# Multilevel Feedback Scheduling


- Preemptive scheduling with dynamic priorities
- N ready to execute queues with decreasing priorities:
  - $P(RQ_0) > P(RQ_1) > \dots > P(RQ_{N-1})$
- Dispatcher selects a process for execution from  $RQ_i$  only if  $RQ_{i-1}$  to  $RQ_0$  are empty

# Multilevel Feedback Scheduling

- New process are placed in  $RQ_0$
- After the first quantum, they are moved to  $RQ_1$ , and to  $RQ_2$  after the second quantum, ... and to  $RQ_{N-1}$  after the Nth quantum
- I/O-bound processes remain in higher priority queues.
  - CPU-bound jobs drift downward.
  - Hence, long jobs may starve

# Multiple Feedback Queues

Different RQs may have different quantum values



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2






# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



# Time Quantum for feedback Scheduling

- With a fixed quantum time, the turnaround time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
  - Time quantum of  $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
  - Possible fix is to promote a process to higher queue after some time

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



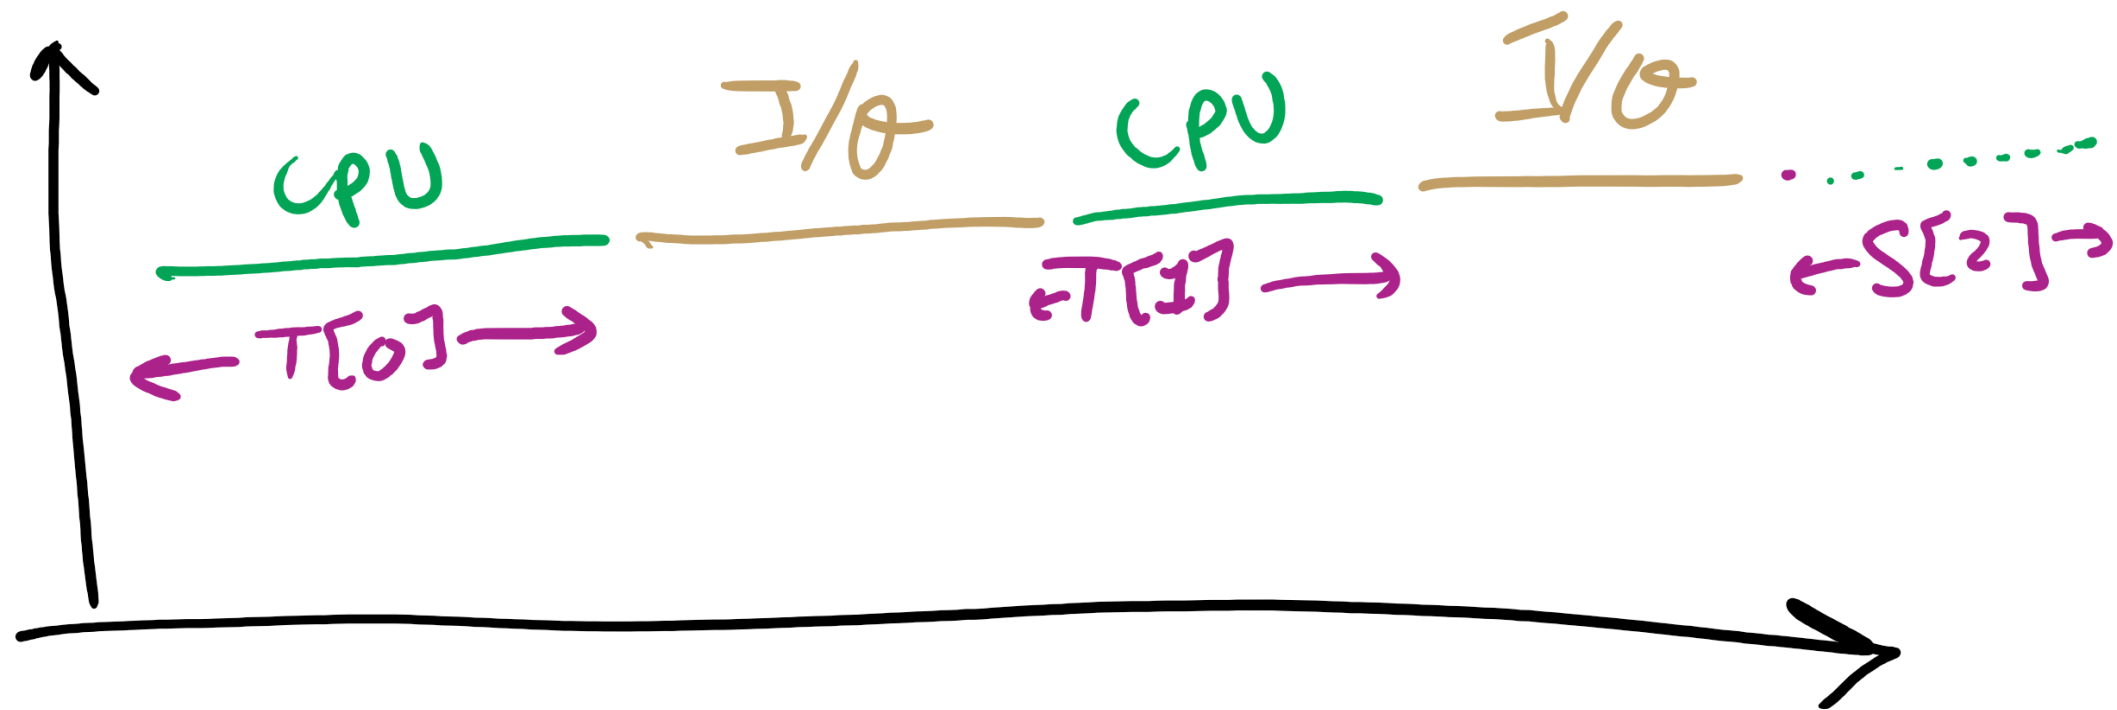
# Algorithm Comparison

- Which one is the best?
- The answer depends on many factors:
  - the system workload (extremely variable)
  - hardware support for the dispatcher
  - relative importance of performance criteria (response time, CPU utilization, throughput...)
  - The evaluation method used (each has its limitations...)

# Back to SJF: CPU Burst Estimation

- Let  $T[i]$  be the execution time for the  $i$ th instance of this process: the actual duration of the  $i$ th CPU burst of this process
- Let  $S[i]$  be the predicted value for the  $i$ th CPU burst of this process. The simplest choice is:
  - $S[n+1] = (1/n)(T[1] + \dots + T[n]) = (1/n) \sum_{\{i=1 \text{ to } n\}} T[i]$
- This can be more efficiently calculated as:
  - $S[n+1] = (1/n) T[n] + ((n-1)/n) S[n]$
- This estimate, however, results in equal weight for each instance

# CPU Burst Estimation






# Estimating the required CPU burst

- Recent instances are more likely to better reflect future behavior
- A common technique to factor the above observation into the estimate is to use **exponential averaging** :
- $S[n+1] = \alpha T[n] + (1 - \alpha) S[n] ; \quad 0 < \alpha < 1$

# CPU burst Estimate (Exponential Average)


- Recent instances have higher weights, whenever  $\alpha > 1/n$
- Expanding the estimated value shows that the weights of past instances decrease exponentially
  - $S[n+1] = \alpha T[n] + (1 - \alpha) \alpha T[n-1] + \dots (1 - \alpha)^{i-1} \alpha T[n-i] +$   
 $\square \dots + (1 - \alpha)^{n-1} S[1]$
  - The predicted value of 1st instance,  $S[1]$ , is usually set to 0 to give priority to new processes

# Exponentially Decreasing Coefficients




# Exponentially Decreasing Coefficients

- $S[1] = 0$  to give high priority to new processes
- Exponential averaging tracks changes in process behavior much faster than simple averaging



# FCFS Problem in HW7




$$\begin{aligned}
 \text{A.R.T.} &= \frac{(7-2) + (1-0) + (19-8) + (10-4) + (14-7)}{5} \\
 &= \frac{5 + 1 + 11 + 6 + 7}{5} = 6\text{ms}
 \end{aligned}$$

$$\text{CPU Util} = \frac{19-4}{19} = \frac{15}{19}$$

# CPU Burst Estimation

$$S[n+1] = \alpha T[n] + (1-\alpha) S[n]$$

$$0 < \alpha < 1$$



# In an ideal world...

- The ideal world has memory that is
  - Very large
  - Very fast
  - Non-volatile (doesn't go away when power is turned off)
- The real world has memory that is:
  - Very large
  - Very fast
  - Affordable!

⇒ Pick any two...
- Memory management goal: make the real world look as much like the ideal world as possible

# Memory hierarchy

- What is the memory hierarchy?
  - Different levels of memory
  - Some are small & fast
  - Others are large & slow
- What levels are usually included?
  - Cache: small amount of fast, expensive memory
    - L1 (level 1) cache: usually on the CPU chip
    - L2 & L3 cache: off-chip, made of SRAM
  - Main memory: medium-speed, medium price memory (DRAM)
  - Disk: many gigabytes of slow, cheap, non-volatile storage
- Memory manager handles the memory hierarchy



# Problem of the Day

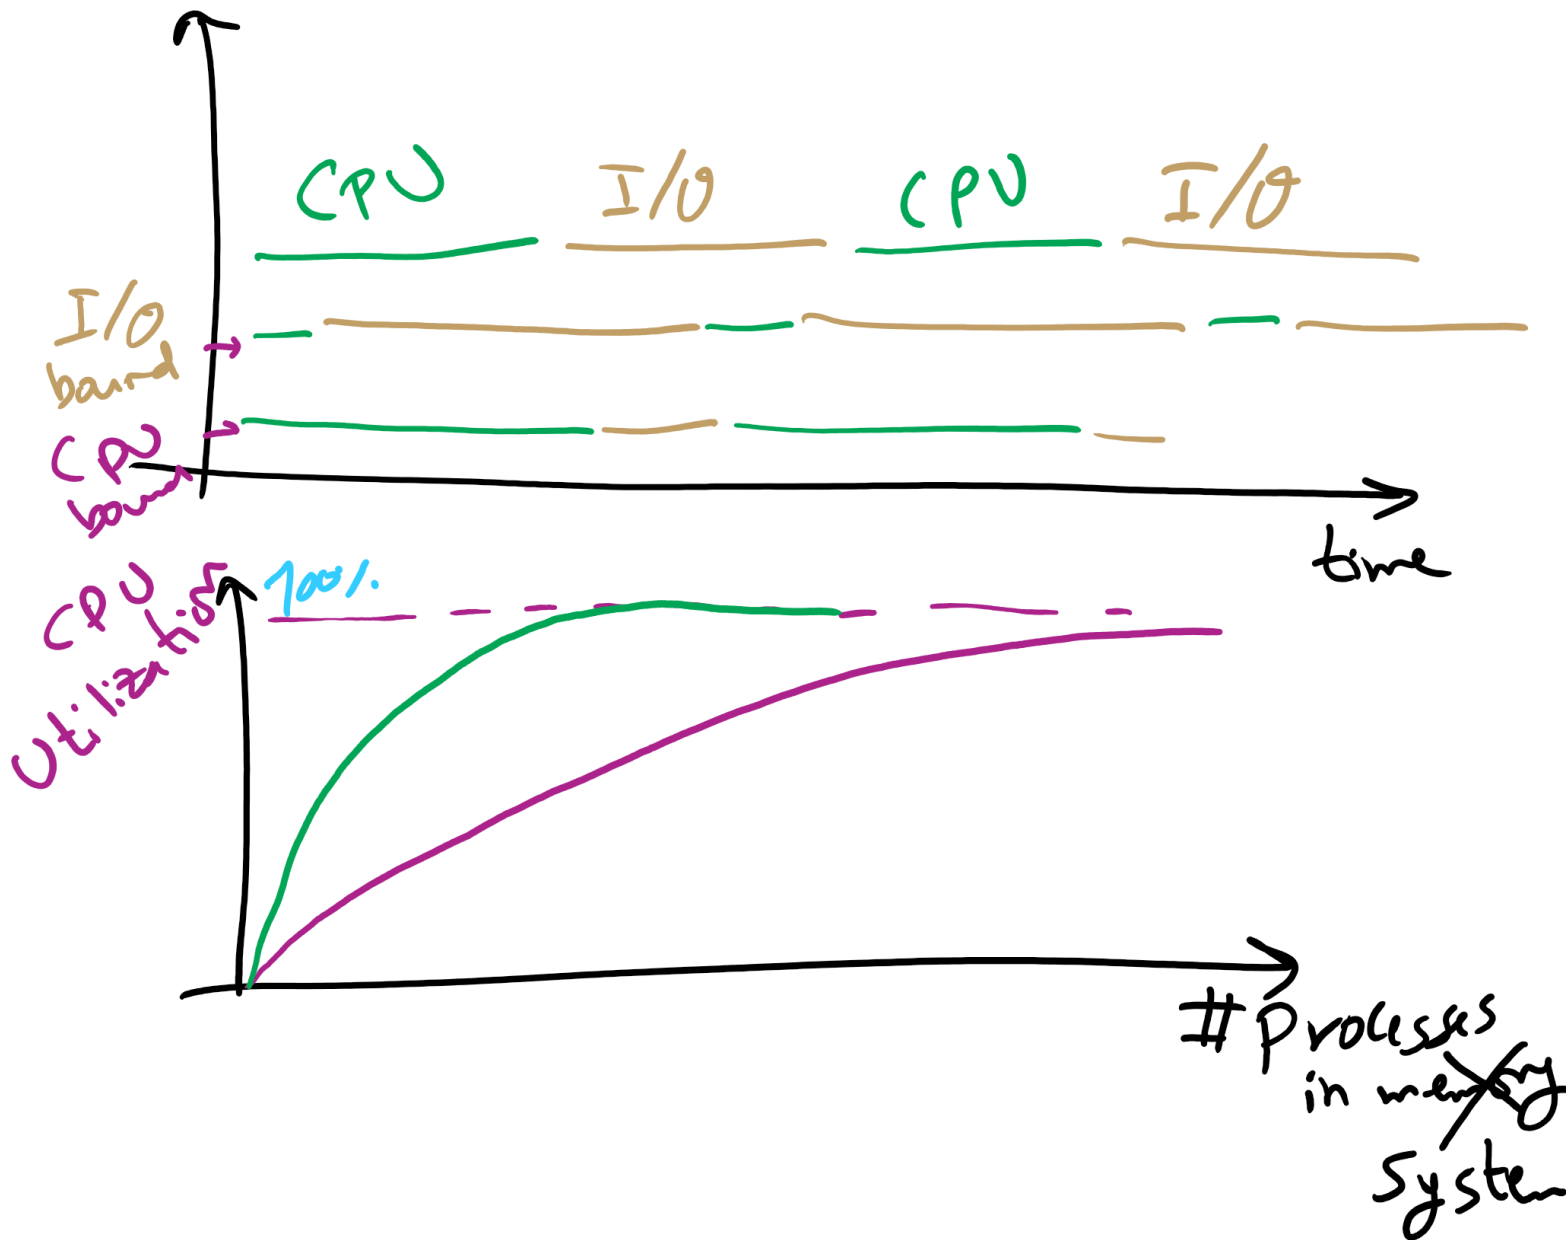
How can we share computer's memory between multiple processes?

How can we protect each process's memory partition from other processes?

# How many programs is enough?


- Several memory partitions (fixed or variable size)
- Lots of processes wanting to use the CPU
- Tradeoff
  - More processes utilize the CPU better
  - Fewer processes use less memory (cheaper!)
- How many processes do we need to keep the CPU fully utilized?
  - This will help determine how much memory we need
  - Is this still relevant with memory costing \$10/GB?

# Why do we need more processes?



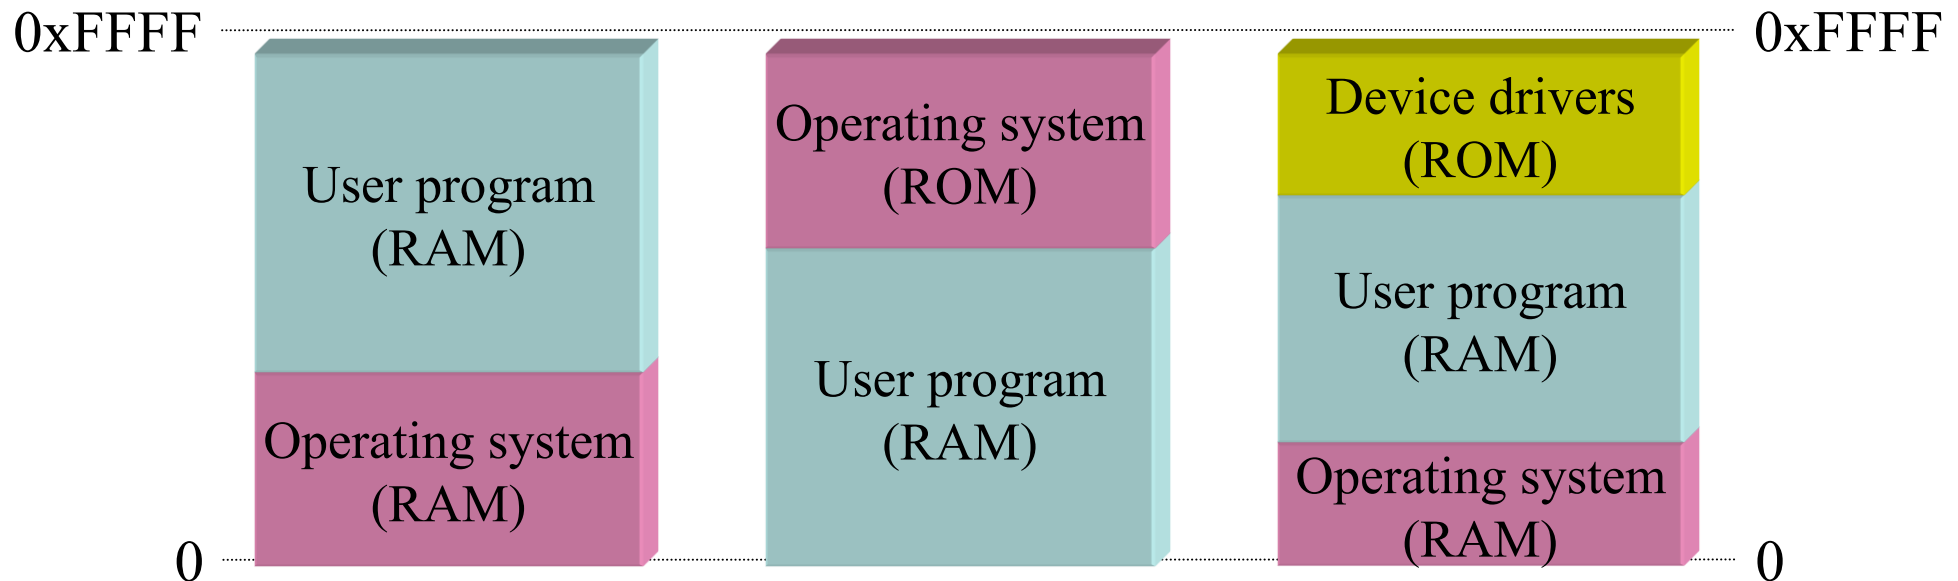
# Modeling multiprogramming

- More I/O wait means less processor utilization
  - At 20% I/O wait, 3–4 processes fully utilize CPU
  - At 80% I/O wait, even 10 processes aren't enough
- This means that the OS should have more processes if they're I/O bound
- More processes => memory management & protection more important!

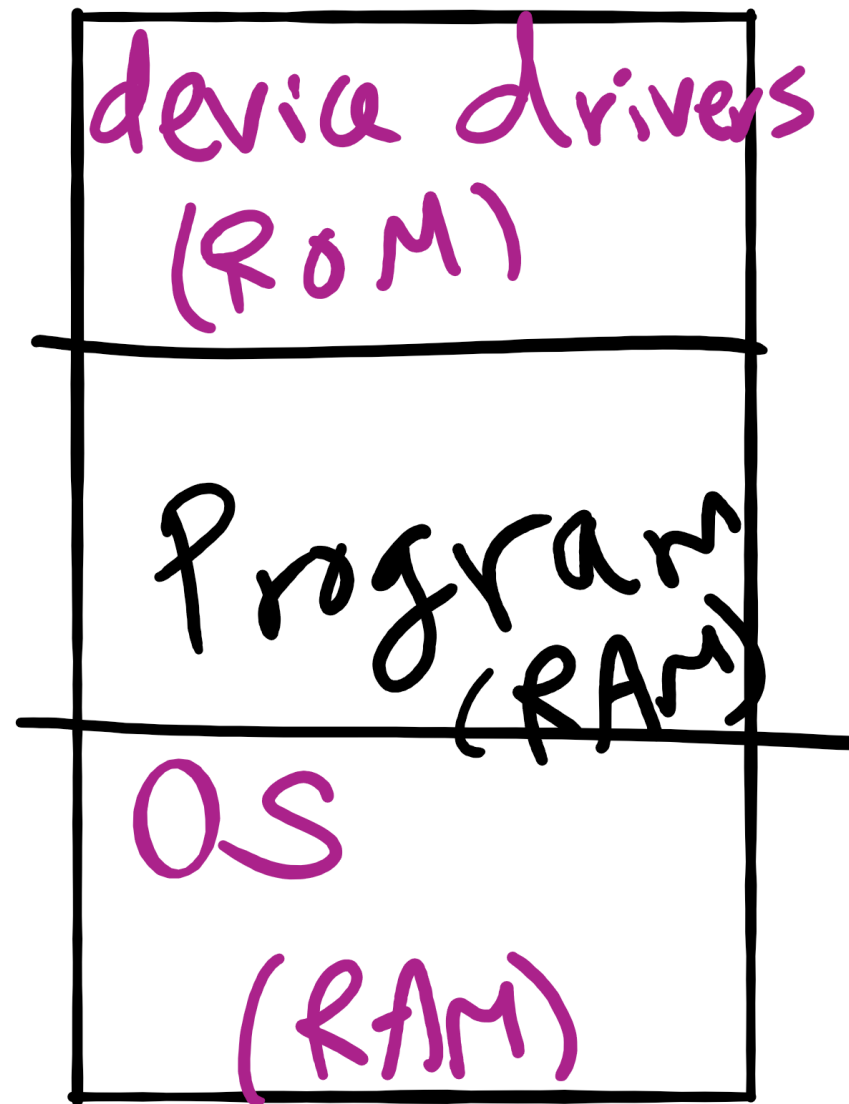


# Basic memory management

- Components include
  - Operating system (perhaps with device drivers)
  - Single process
- Goal: lay these out in memory
  - Memory protection may not be an issue (only one program)
  - Flexibility may still be useful (allow OS changes, etc.)
- No swapping or paging



# Memory Management for Embedded Systems




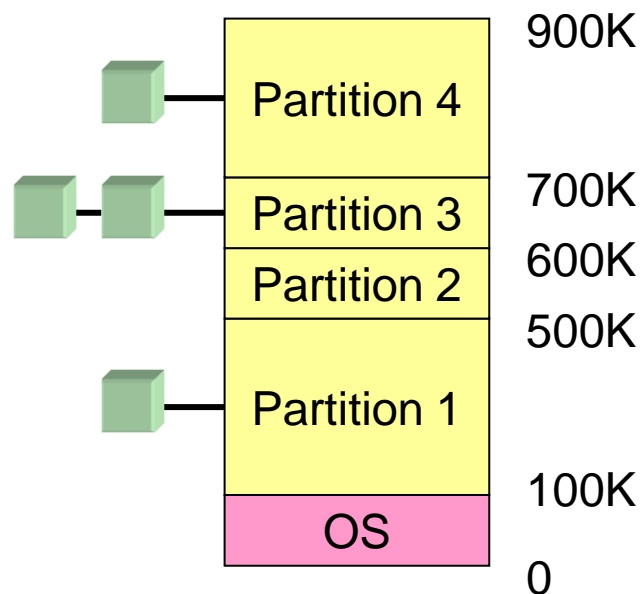
# Fixed partitions: multiple programs

- Fixed memory partitions

- Divide memory into fixed spaces
- Assign a process to a space when it's free

- Mechanisms

- Separate input queues for each partition
- Single input queue: better ability to optimize CPU usage



# Problem of the Day

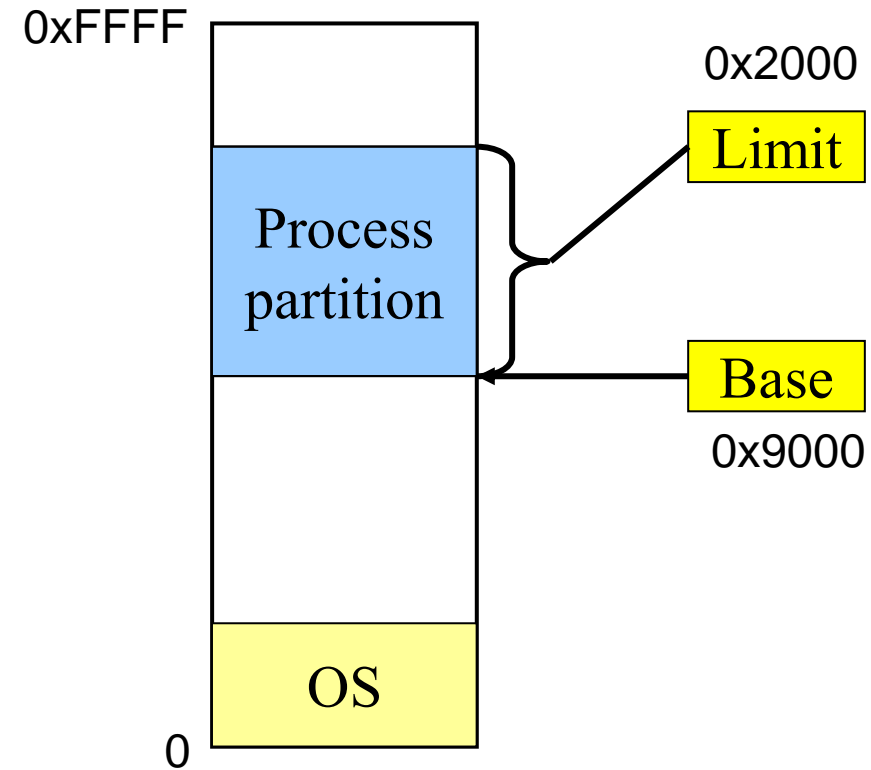
How can we share computer's memory between multiple processes?

How can we protect each process's memory partition from other processes?




# Base and limit registers

- Special CPU registers: base & limit
  - Access to the registers limited to kernel (privileged) mode
  - Registers contain
    - Base: start of the process's memory partition
    - Limit: length of the process's memory partition
- Address generation
  - Physical address: location in actual memory
  - Logical address: location from the process's point of view
  - Physical address = base + logical address
  - Logical address larger than limit => error





Logical address: 0x1204  
Physical address:  
 $0x1204 + 0x9000 = 0xa204$

# Virtual and physical addresses



- Program uses *virtual addresses*
  - Addresses local to the process
  - Hardware translates virtual address to *physical address*
- Translation done by the **Memory Management Unit**
  - Usually on the same chip as the CPU
  - Only physical addresses leave the CPU/MMU chip
- Physical memory indexed by physical addresses

# Address Translation




# Virtual memory

- Basic idea: allow the OS to hand out more memory than exists on the system
- Keep recently used stuff in physical memory
- Move less recently used stuff to disk
- Keep all of this hidden from processes
  - Processes still see an address space from 0 – max address
  - Movement of information to and from disk handled by the OS without process help
- Virtual memory (VM) especially helpful in multiprogrammed system
  - CPU schedules process B while process A waits for its memory to be retrieved from disk

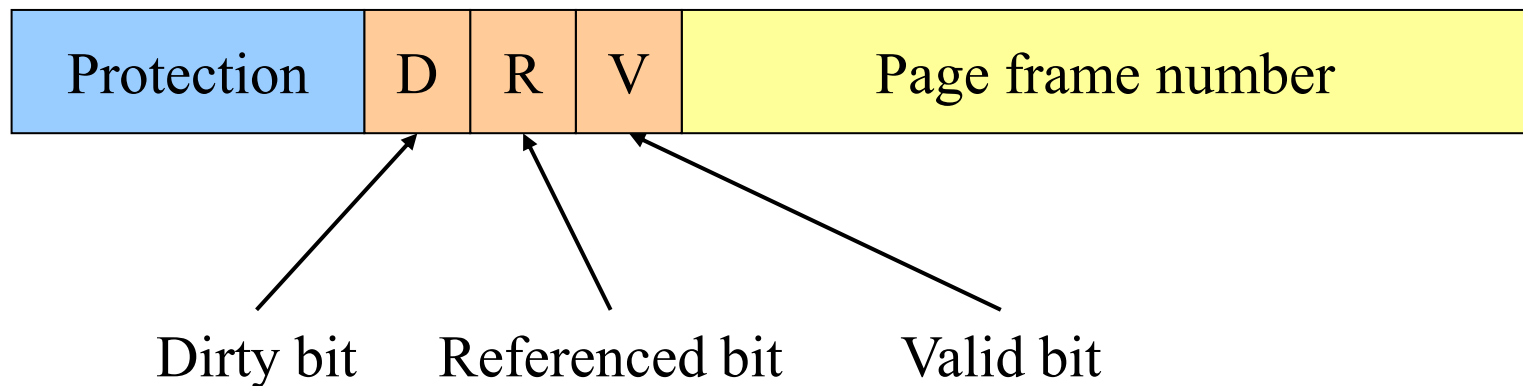
# Paging and page tables

- Virtual addresses mapped to physical addresses
  - Unit of mapping is called a *page*
  - All addresses in the same virtual page are in the same physical page
  - *Page table entry* (PTE) contains translation for a single page
- Table translates virtual page number to physical page number
  - Not all virtual memory has a physical page
  - Not every physical page need be used
- Example:
  - 64 KB virtual memory
  - 32 KB physical memory



# What's in a page table entry?

- Each entry in the page table contains
  - Valid bit: set if this logical page number has a corresponding physical frame in memory
    - If not valid, remainder of PTE is irrelevant
  - Page frame number: page in physical memory
  - Referenced bit: set if data on the page has been accessed
  - Dirty (modified) bit :set if data on the page has been modified
  - Protection information



# Implementing page tables in hardware

- Page table resides in main (physical) memory
- CPU uses special registers for paging
  - Page table base register (PTBR) points to the page table
  - Page table length register (PTLR) contains length of page table: restricts maximum legal logical address
- Translating an address requires two memory accesses
  - First access reads page table entry (PTE)
  - Second access reads the data / instruction from memory
- Reduce number of memory accesses
  - Can't avoid second access (we need the value from memory)
  - Eliminate first access by keeping a hardware cache (called a *translation lookaside buffer* or TLB) of recently used page table entries