



University of  
Pittsburgh

# Introduction to Operating Systems CS 1550



Spring 2023  
Sherif Khattab  
ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013**)

# Announcements

- Upcoming deadlines
  - Homework 8 is due **this Friday**
  - Quiz 1 and Lab 2 due on Tuesday 2/28 at 11:59 pm
  - Project 2 is due Friday 3/17 at 11:59 pm

# In an ideal world...

- The ideal world has memory that is
  - Very large
  - Very fast
  - Non-volatile (doesn't go away when power is turned off)
- The real world has memory that is:
  - Very large
  - Very fast
  - Affordable!

⇒ Pick any two...
- Memory management goal: make the real world look as much like the ideal world as possible

# Memory hierarchy

- What is the memory hierarchy?
  - Different levels of memory
  - Some are small & fast
  - Others are large & slow
- What levels are usually included?
  - Cache: small amount of fast, expensive memory
    - L1 (level 1) cache: usually on the CPU chip
    - L2 & L3 cache: off-chip, made of SRAM
  - Main memory: medium-speed, medium price memory (DRAM)
  - Disk: many gigabytes of slow, cheap, non-volatile storage
- Memory manager handles the memory hierarchy

# Problem of the Day


How can we share computer's memory between multiple processes?

How can we protect each process's memory partition from other processes?

# How many programs is enough?


- Several memory partitions (fixed or variable size)
- Lots of processes wanting to use the CPU
- Tradeoff
  - More processes utilize the CPU better
  - Fewer processes use less memory (cheaper!)
- How many processes do we need to keep the CPU fully utilized?
  - This will help determine how much memory we need
  - Is this still relevant with memory costing \$10/GB?

# Why do we need more processes?



# Modeling multiprogramming

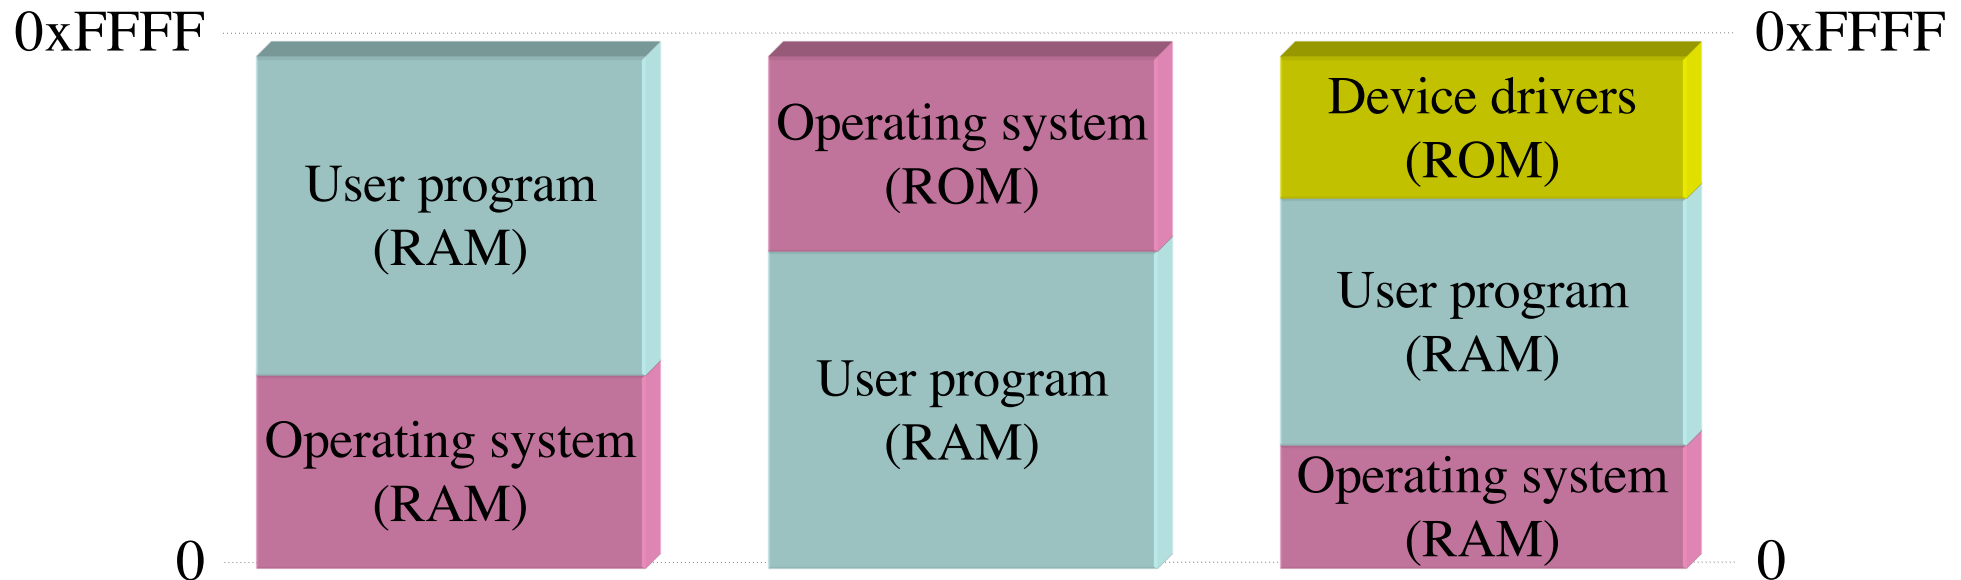
- More I/O wait means less processor utilization
  - At 20% I/O wait, 3–4 processes fully utilize CPU
  - At 80% I/O wait, even 10 processes aren't enough
- This means that the OS should have more processes if they're I/O bound
- More processes => memory management & protection more important!



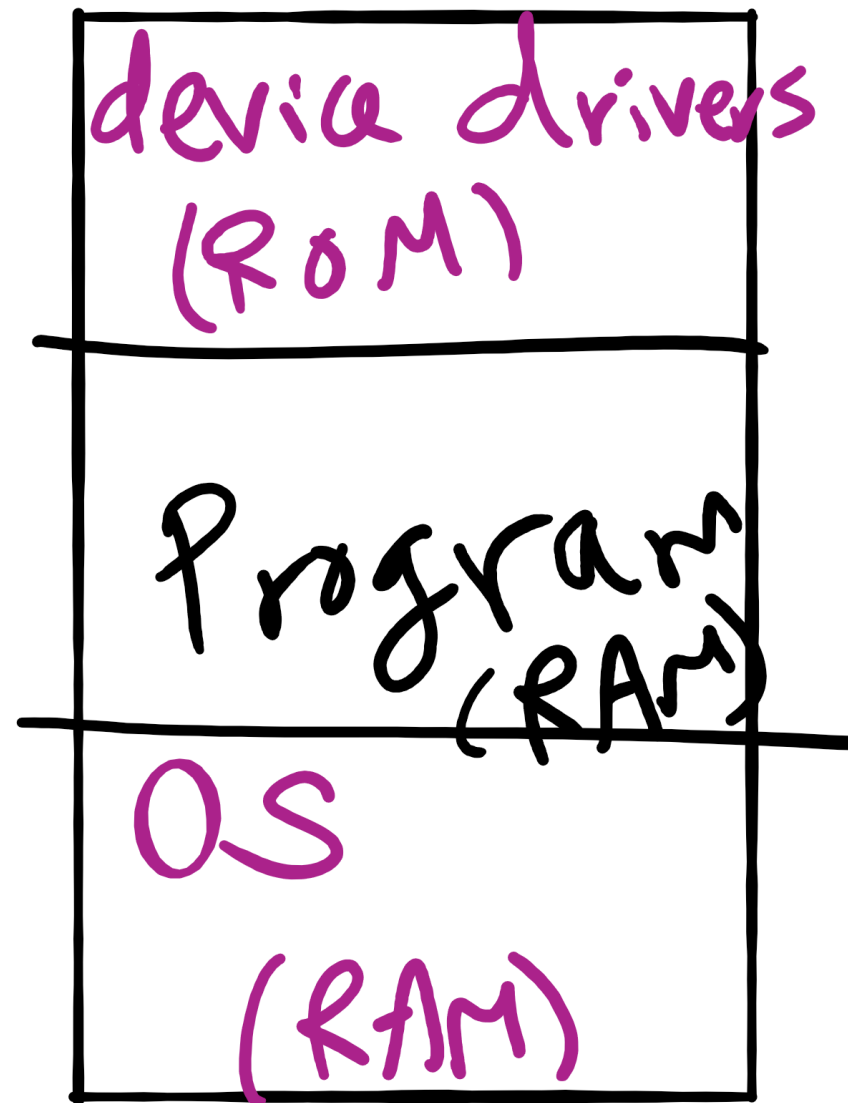


# Basic memory management

- Components include
  - Operating system (perhaps with device drivers)
  - Single process
- Goal: lay these out in memory
  - Memory protection may not be an issue (only one program)
  - Flexibility may still be useful (allow OS changes, etc.)
- No swapping or paging



# Memory Management for Embedded Systems




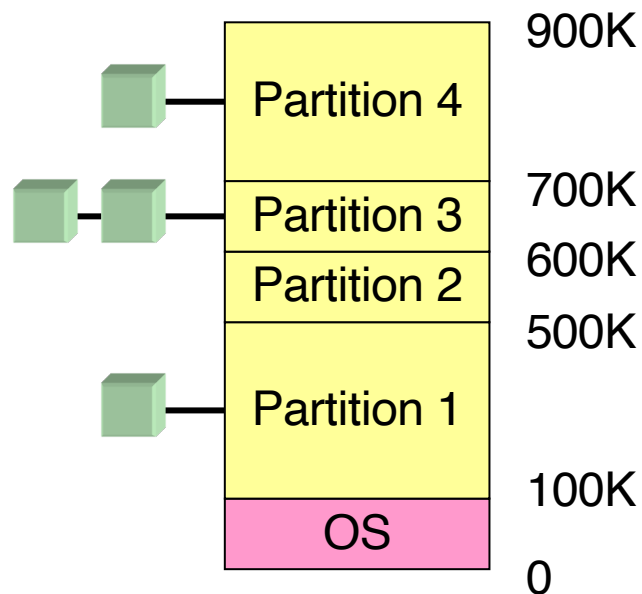
# Fixed partitions: multiple programs

- Fixed memory partitions

- Divide memory into fixed spaces
- Assign a process to a space when it's free

- Mechanisms

- Separate input queues for each partition
- Single input queue: better ability to optimize CPU usage



# Problem of the Day

How can we share computer's memory between multiple processes?

How can we protect each process's memory partition from other processes?

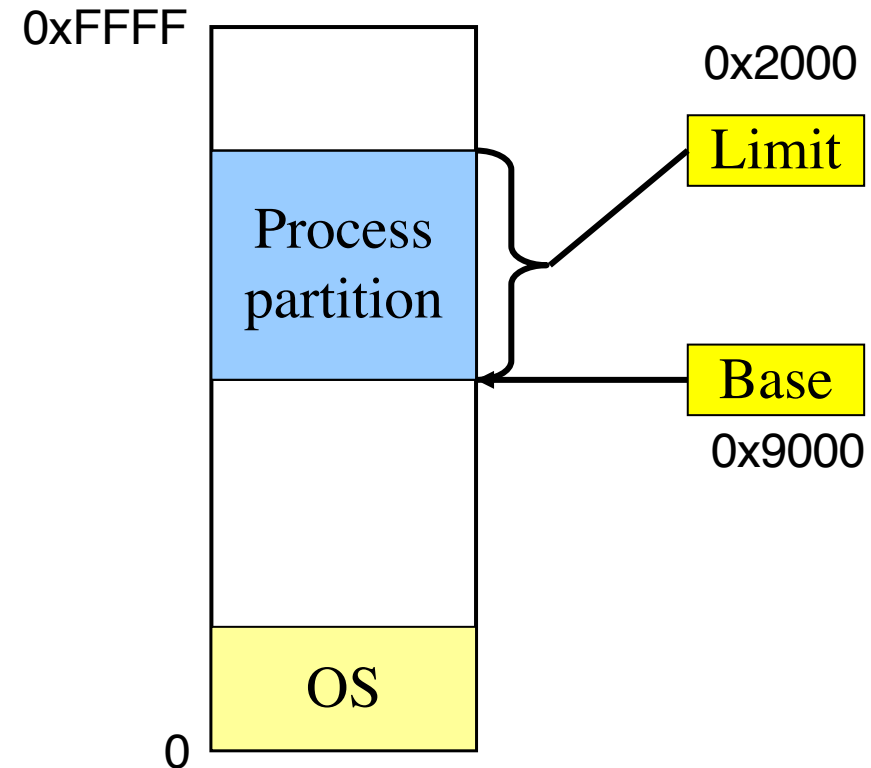
# Base and limit registers

- Special CPU registers: base & limit

- Access to the registers limited to kernel (privileged) mode
- Registers contain
  - Base: start of the process's memory partition
  - Limit: length of the process's memory partition

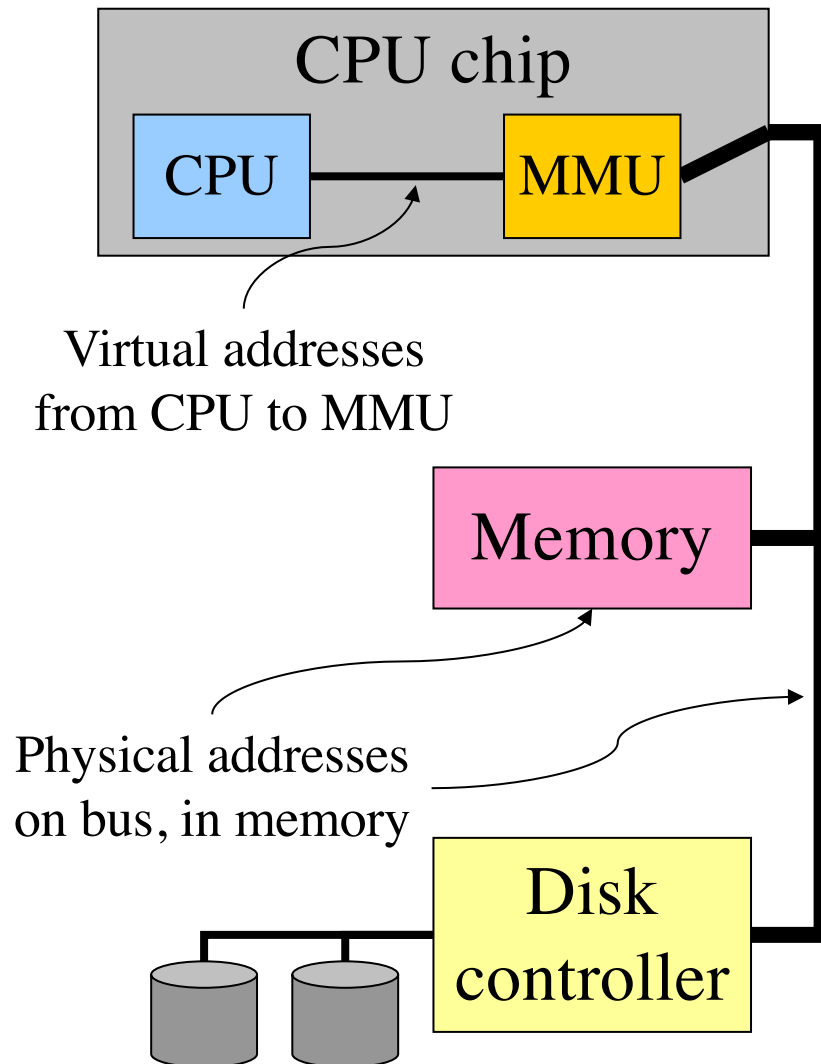
- Address generation

- Physical address: location in actual memory
- Logical address: location from the process's point of view
- Physical address = base + logical address
- Logical address larger than limit => error



Logical address: 0x1204  
Physical address:  
 $0x1204 + 0x9000 = 0xa204$

# Virtual and physical addresses



- Program uses *virtual addresses*
  - Addresses local to the process
  - Hardware translates virtual address to *physical address*
- Translation done by the **Memory Management Unit**
  - Usually on the same chip as the CPU
  - Only physical addresses leave the CPU/MMU chip
- Physical memory indexed by physical addresses


# Address Translation

```
int x = 5;  
int *p = &x;  
printf("/p\n", p);
```

2000

Physical Memory

x ?



# Virtual memory

- Basic idea: allow the OS to hand out more memory than exists on the system
- Keep recently used stuff in physical memory
- Move less recently used stuff to disk
- Keep all of this hidden from processes
  - Processes still see an address space from 0 – max address
  - Movement of information to and from disk handled by the OS without process help
- Virtual memory (VM) especially helpful in multiprogrammed system
  - CPU schedules process B while process A waits for its memory to be retrieved from disk



# Paging and page tables

- Virtual addresses mapped to physical addresses
  - Unit of mapping is called a *page*
  - All addresses in the same virtual page are in the same physical page
  - *Page table entry* (PTE) contains translation for a single page
- Table translates virtual page number to physical page number
  - Not all virtual memory has a physical page
  - Not every physical page need be used
- Example:
  - 64 KB virtual memory
  - 32 KB physical memory

