University of Pittsburgh

Introduction to Operating Systems
CS 1550

Spring 2023

Sherif Khattab

ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013)**

Announcements

- Upcoming deadlines
  - Homework 9 due this Friday
  - Project 3 is due Friday 4/7 at 11:59 pm
  - Lab 4 is due on Tuesday 4/11 at 11:59 pm

Previous Lecture

- How to simulate page replacement algorithms
  - Clock, LRU, OPT

This Lecture …

- Problem of Too Large Page Tables
  - multi-level page tables + TLB
  - inverted page tables
- Miscellaneous Memory Management Issues

Page Table Size

- How big can a page table be?
  - 64-bit machine
  - 4 KB page size
  - How many pages?
  - How many PTEs?
  - How big is a PTE?
  - How big is the page table of one process?

Virtual Address

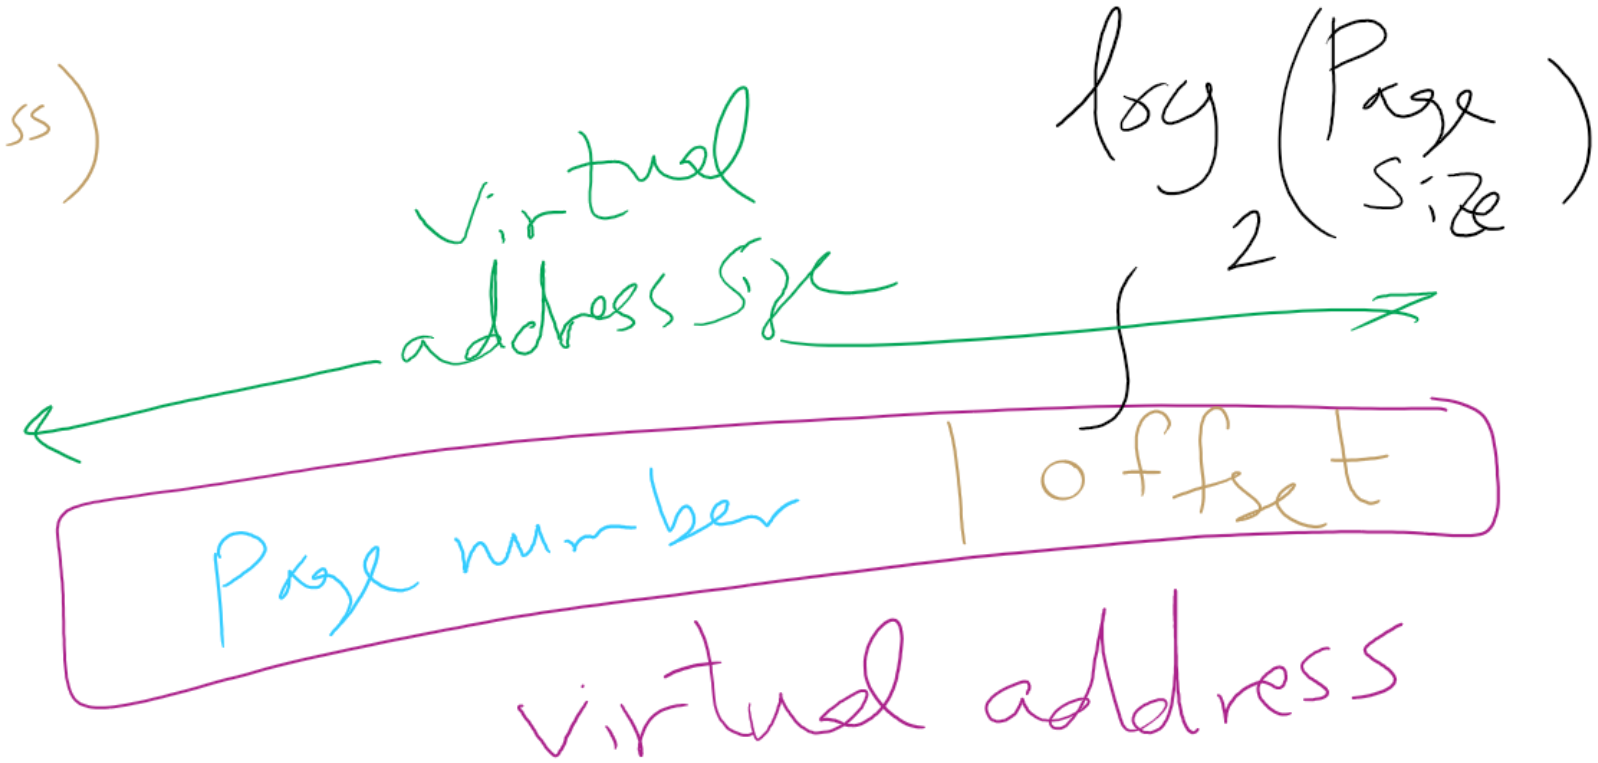

Page number vs. Offset

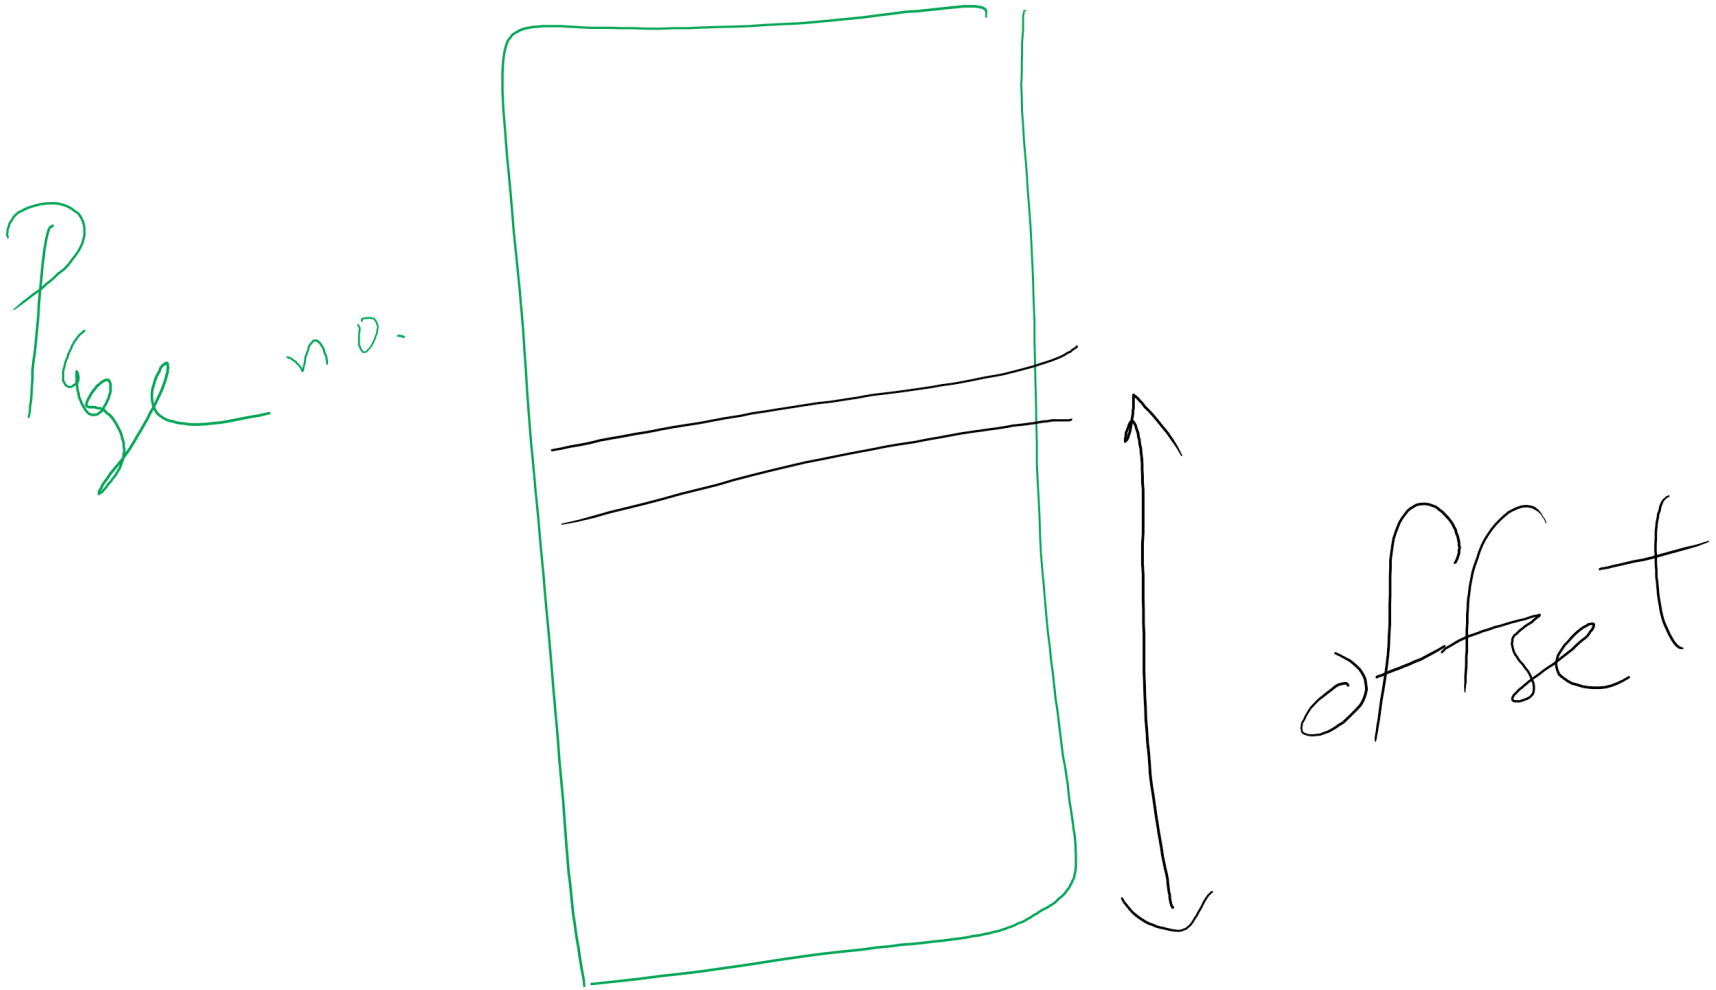

Memory Sizes

KiB = 2^{10}

MiB = 2^{20}

GiB = 2^{30}

TiB = 2^{40}

PiB = 2^{50}

Example on splitting a virtual address

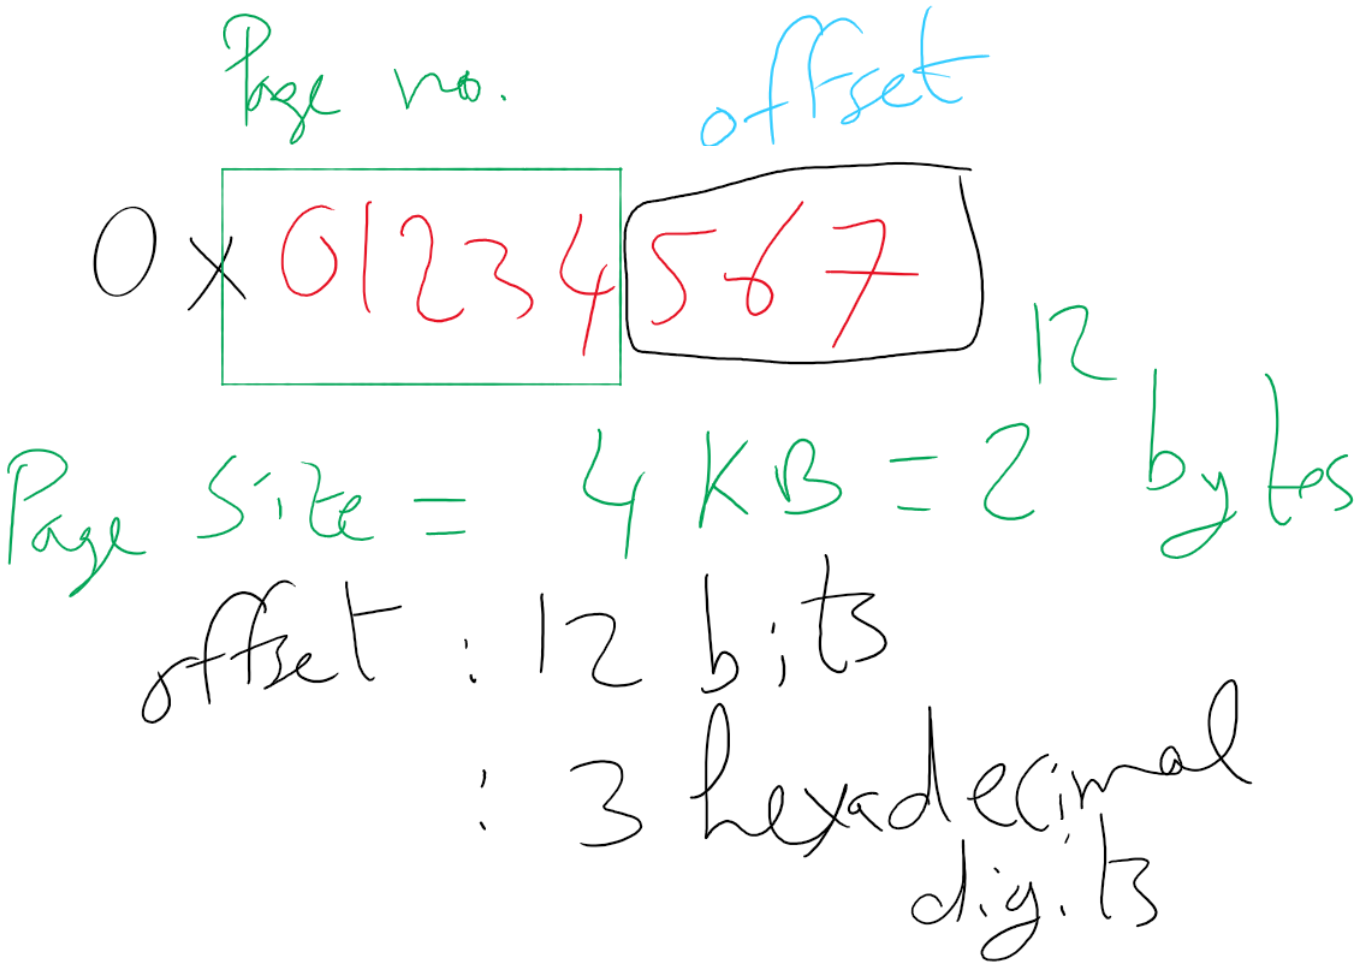

PTEs

$$\#PTEs = \frac{\text{Virtual address space size}}{\text{Page size}} = \frac{2^{(\text{logical address size})}}{\text{Page size}}$$

$$= 2^{(\text{Page no. size})}$$

$\lceil$

PTE Size

$$\text{frame number size} = \left\lceil \log_2 \#\text{frames} \right\rceil$$

$$\#\text{frames} = \frac{\text{RAM size}}{\text{frame size}}$$

Page Table Size

Page Table Size = # PTEs * PTE size

Page Table Size

- How big can a page table be?
  - 64-bit machine → virtual address size = 64 bits
  - 4 KB page size →page offset is $\log_2 2^{12} = 12$ bits
  - How many pages?
  - How many PTEs?
  - #PTEs = 2^{64} / page size = 2^{64} / 2^{12} = 2^{52}
  - How big is a PTE?
  - Assume 4 GB RAM (Physical Memory) = 2^{32} bytes
  - # Frames = 2^{32} / page size = 2^{32} / 2^{12} = 2^{20}
  - Frame no. size = $\log_2 2^{20}$ = 20 bits
  - PTE size = 20 + (let's say) 12 bits for flags = 32 bits = 4 bytes
  - How big is the page table of **one process**?
  - #PTEs * PTE size = 2^{52} * 4 = 2^{54} bytes = **16 PB**

Problem of the Day

- How big is the page table of **one process**?
- #PTEs * PTE size = 2^{52} * 4 = 2^{54} bytes = **16 PB**
- Such page table is too big to put in memory!
- But why does it have to be in memory?
- Let's see how address translation happens …

Address Translation Structures

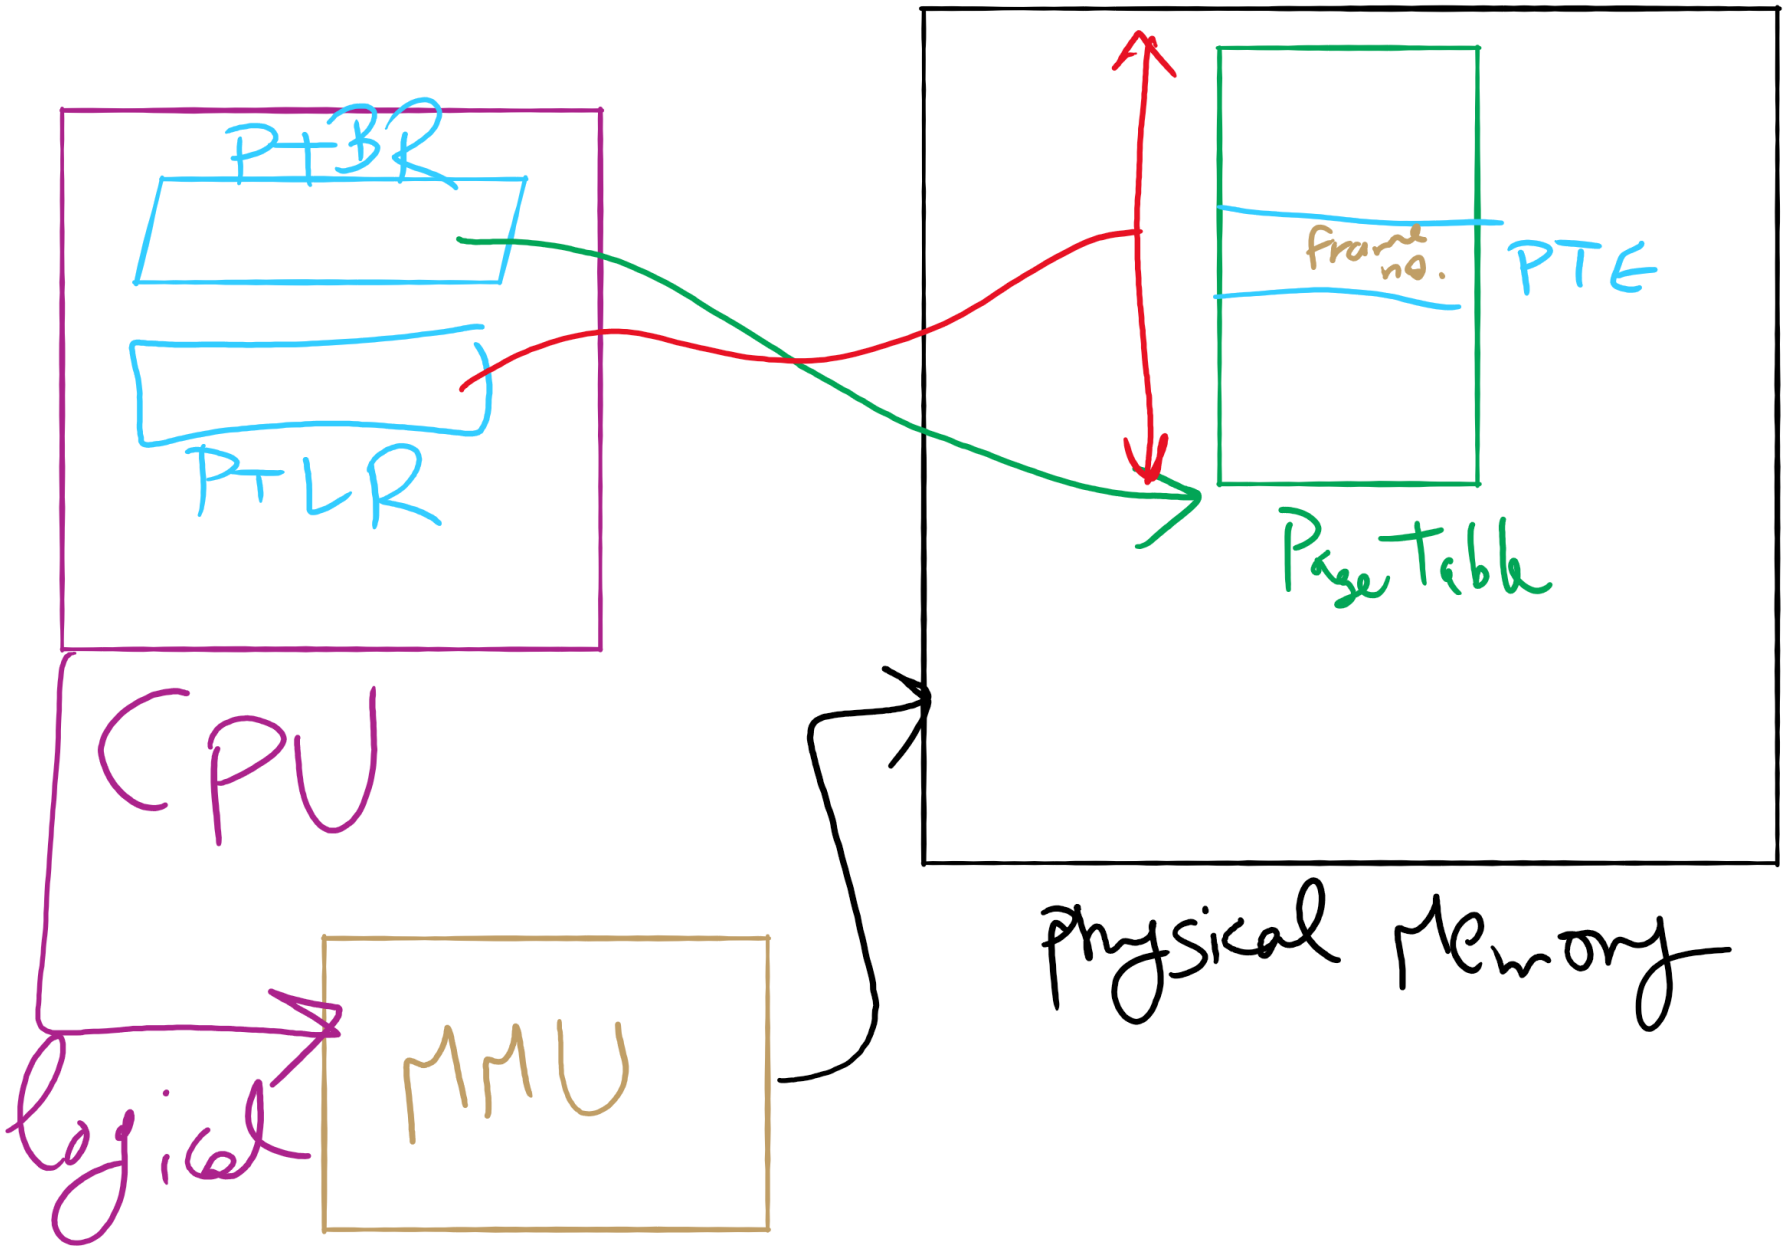

Memory & paging structures

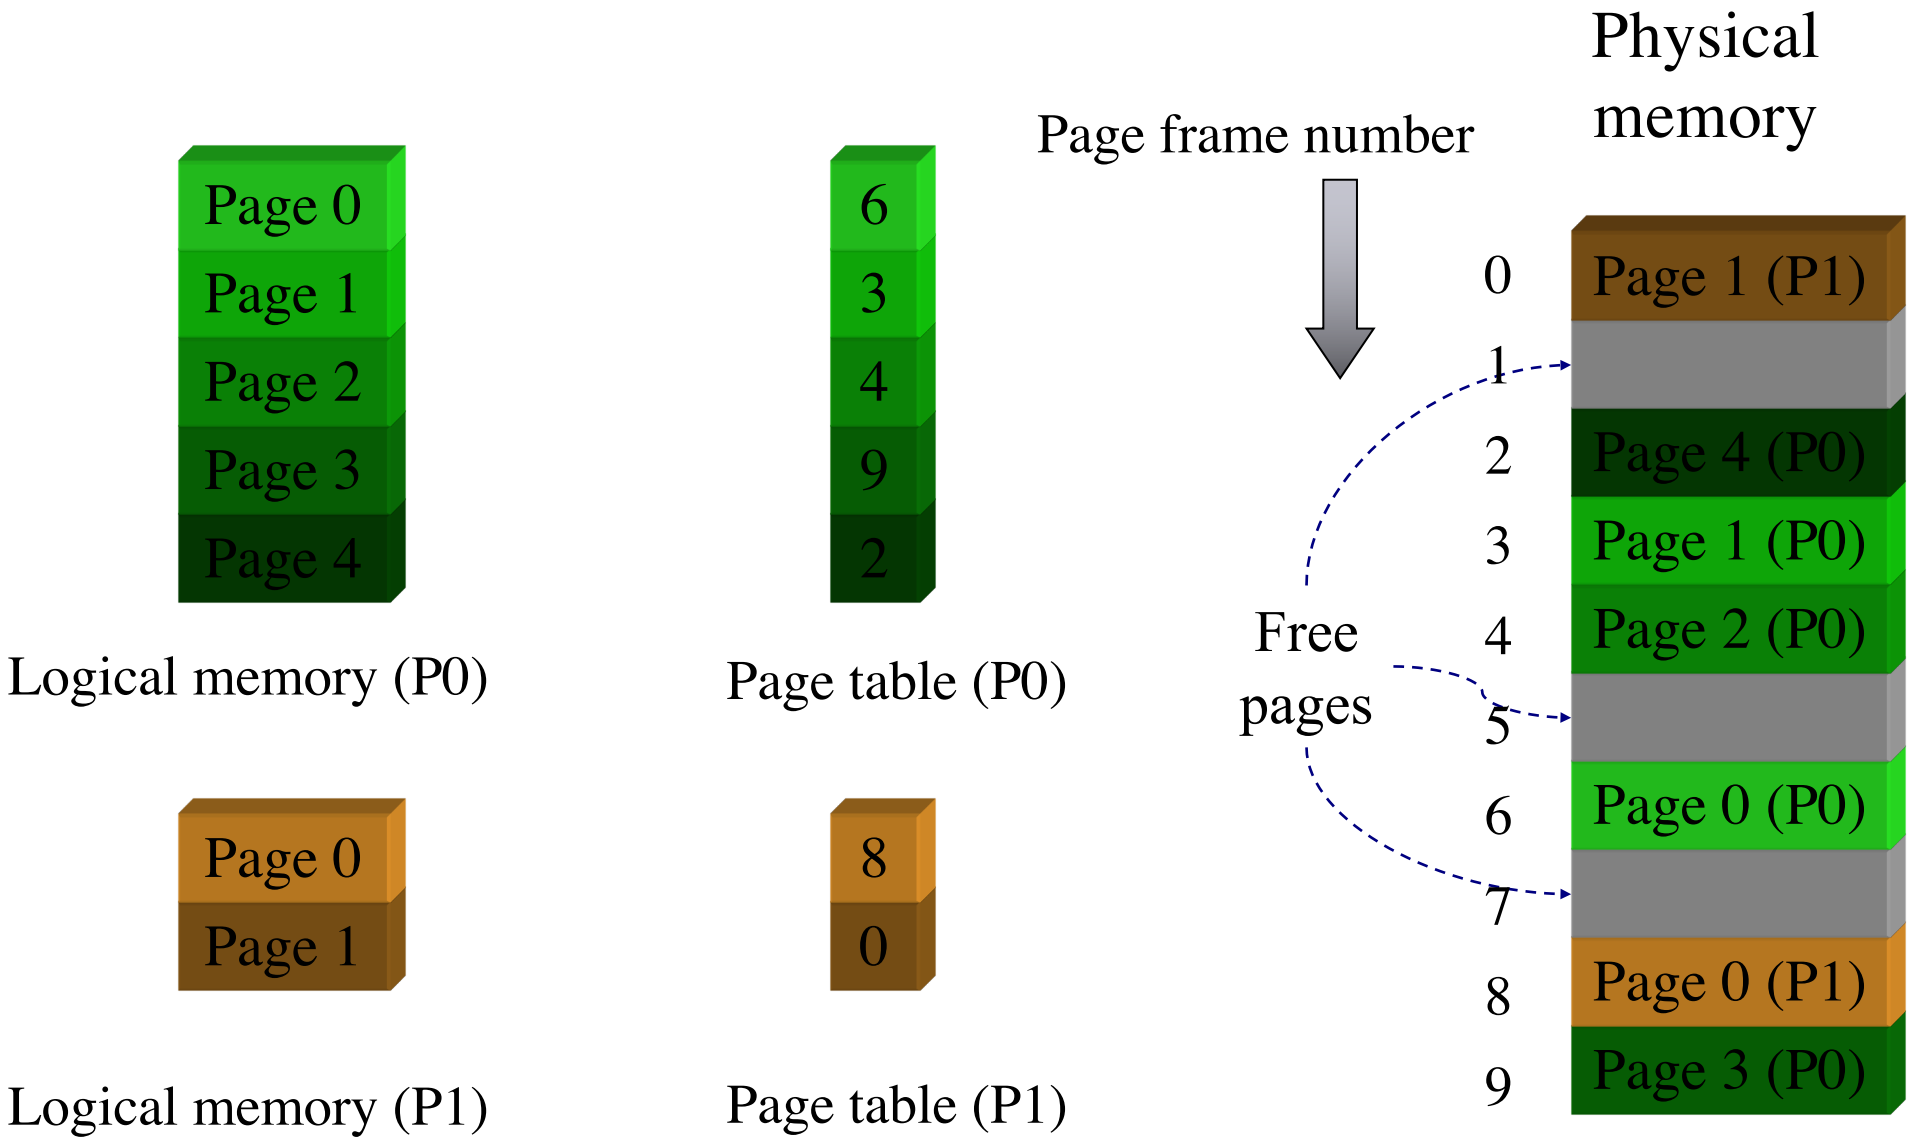

Mapping logical => physical address

- Split address from CPU into two pieces
  - Page number (*p*)
  - Page offset (*d*)
- Page number
  - Index into page table
  - Page table contains base address of page in physical memory
- Page offset
  - Added to base address to get actual physical memory address
- Page size = 2^d bytes

Example:

- 4 KB (=4096 byte) pages
- 32 bit logical addresses

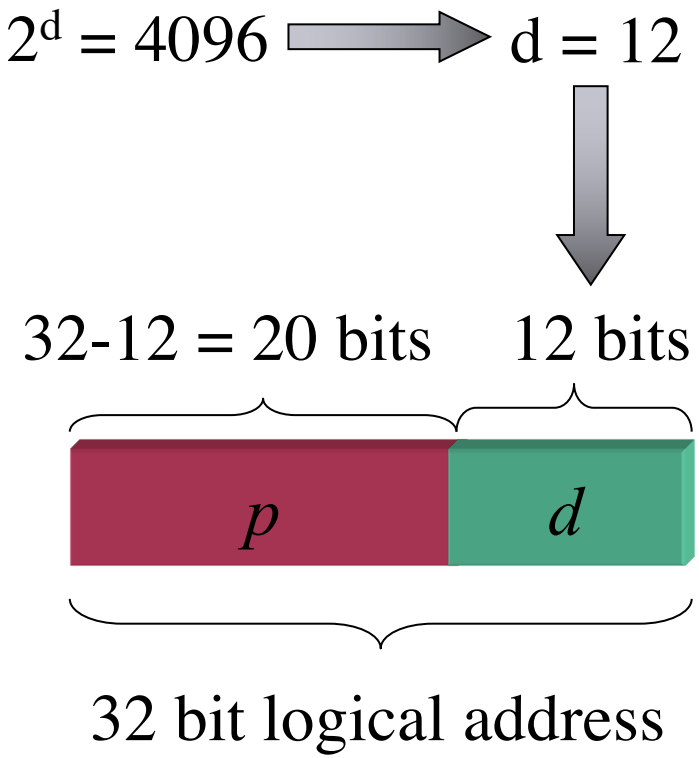

Address translation architecture

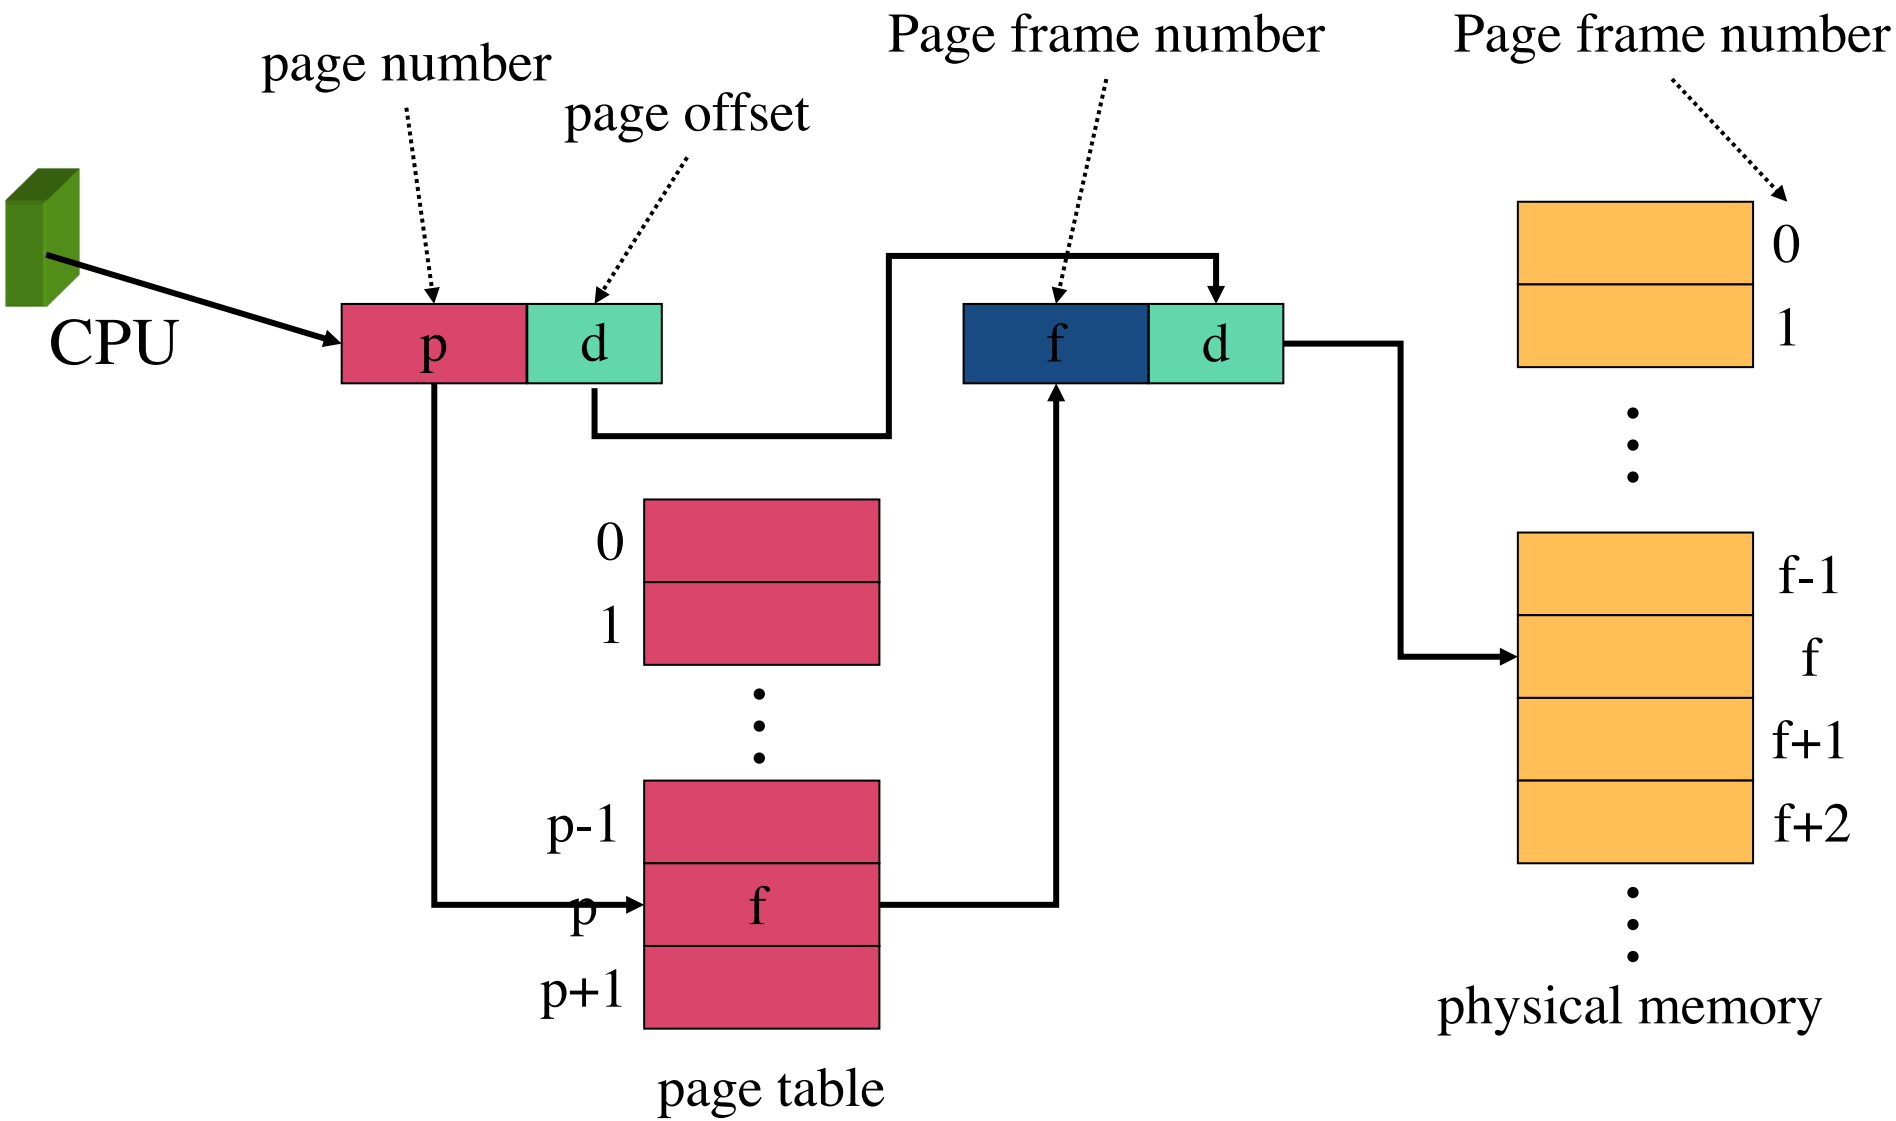

Solution 1: Two-level page tables

- 1st level page table called **page directory**
  - PDEs have pointers to 2nd level page tables
- PTEs in 2nd level page table have actual physical page numbers
- All addresses are physical
- Protection bits kept in 2nd level

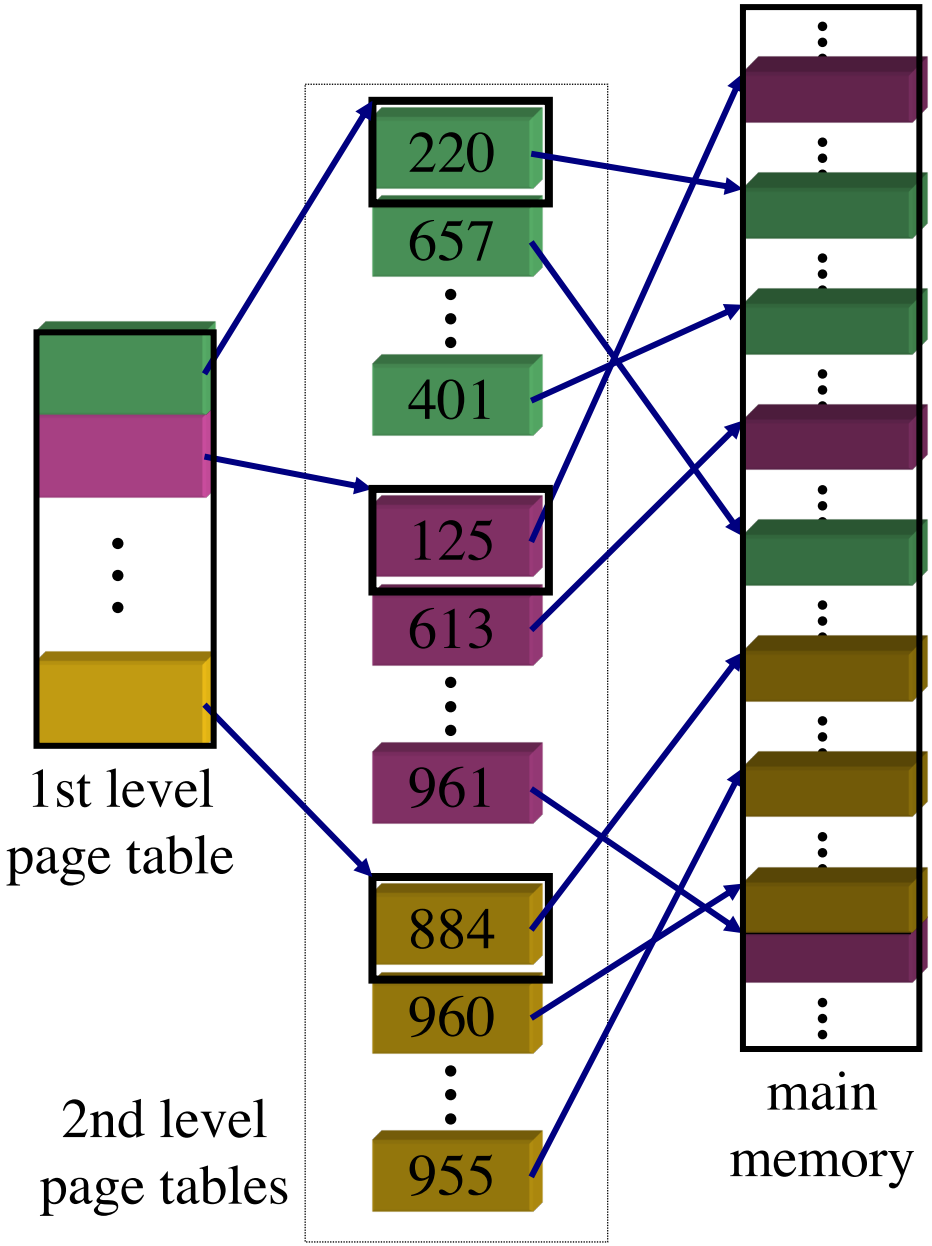

How does that solve the problem?

- 2nd level page tables don't have to be in memory all at once
  - can be loaded **on demand**
  - PDEs marked invalid in first page table

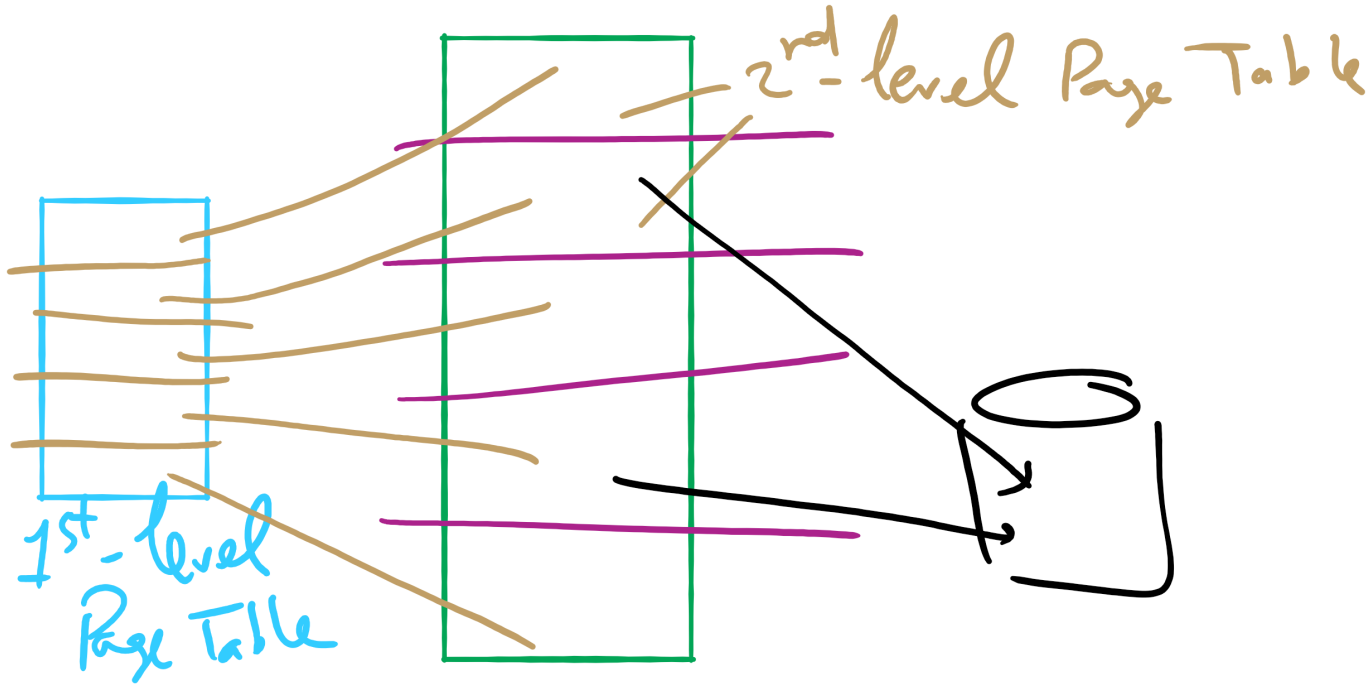

More on two-level page tables

- Total number of bits indexing 1st and 2nd level table is constant for a given page size and logical address length. **Why?**
- More bits in 1st level: fine granularity at 2nd level
- Fewer bits in 1st level: maybe less wasted space?

Two-level paging: address translation

- 8 KB pages
- 32-bit logical address
- p_1 is an index into the 1st level page table
- p_2 is an index into the 2nd level page table pointed to by p_1

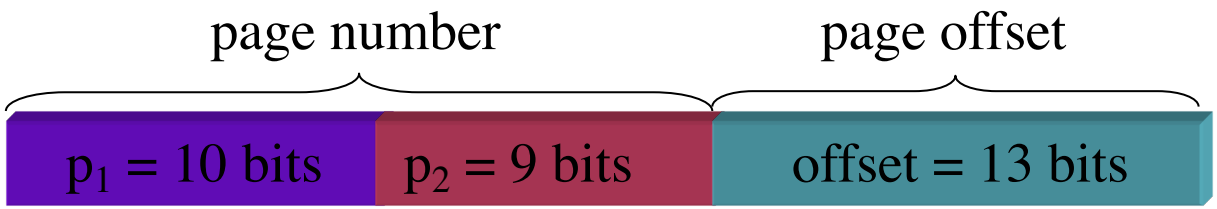

2-level address translation example

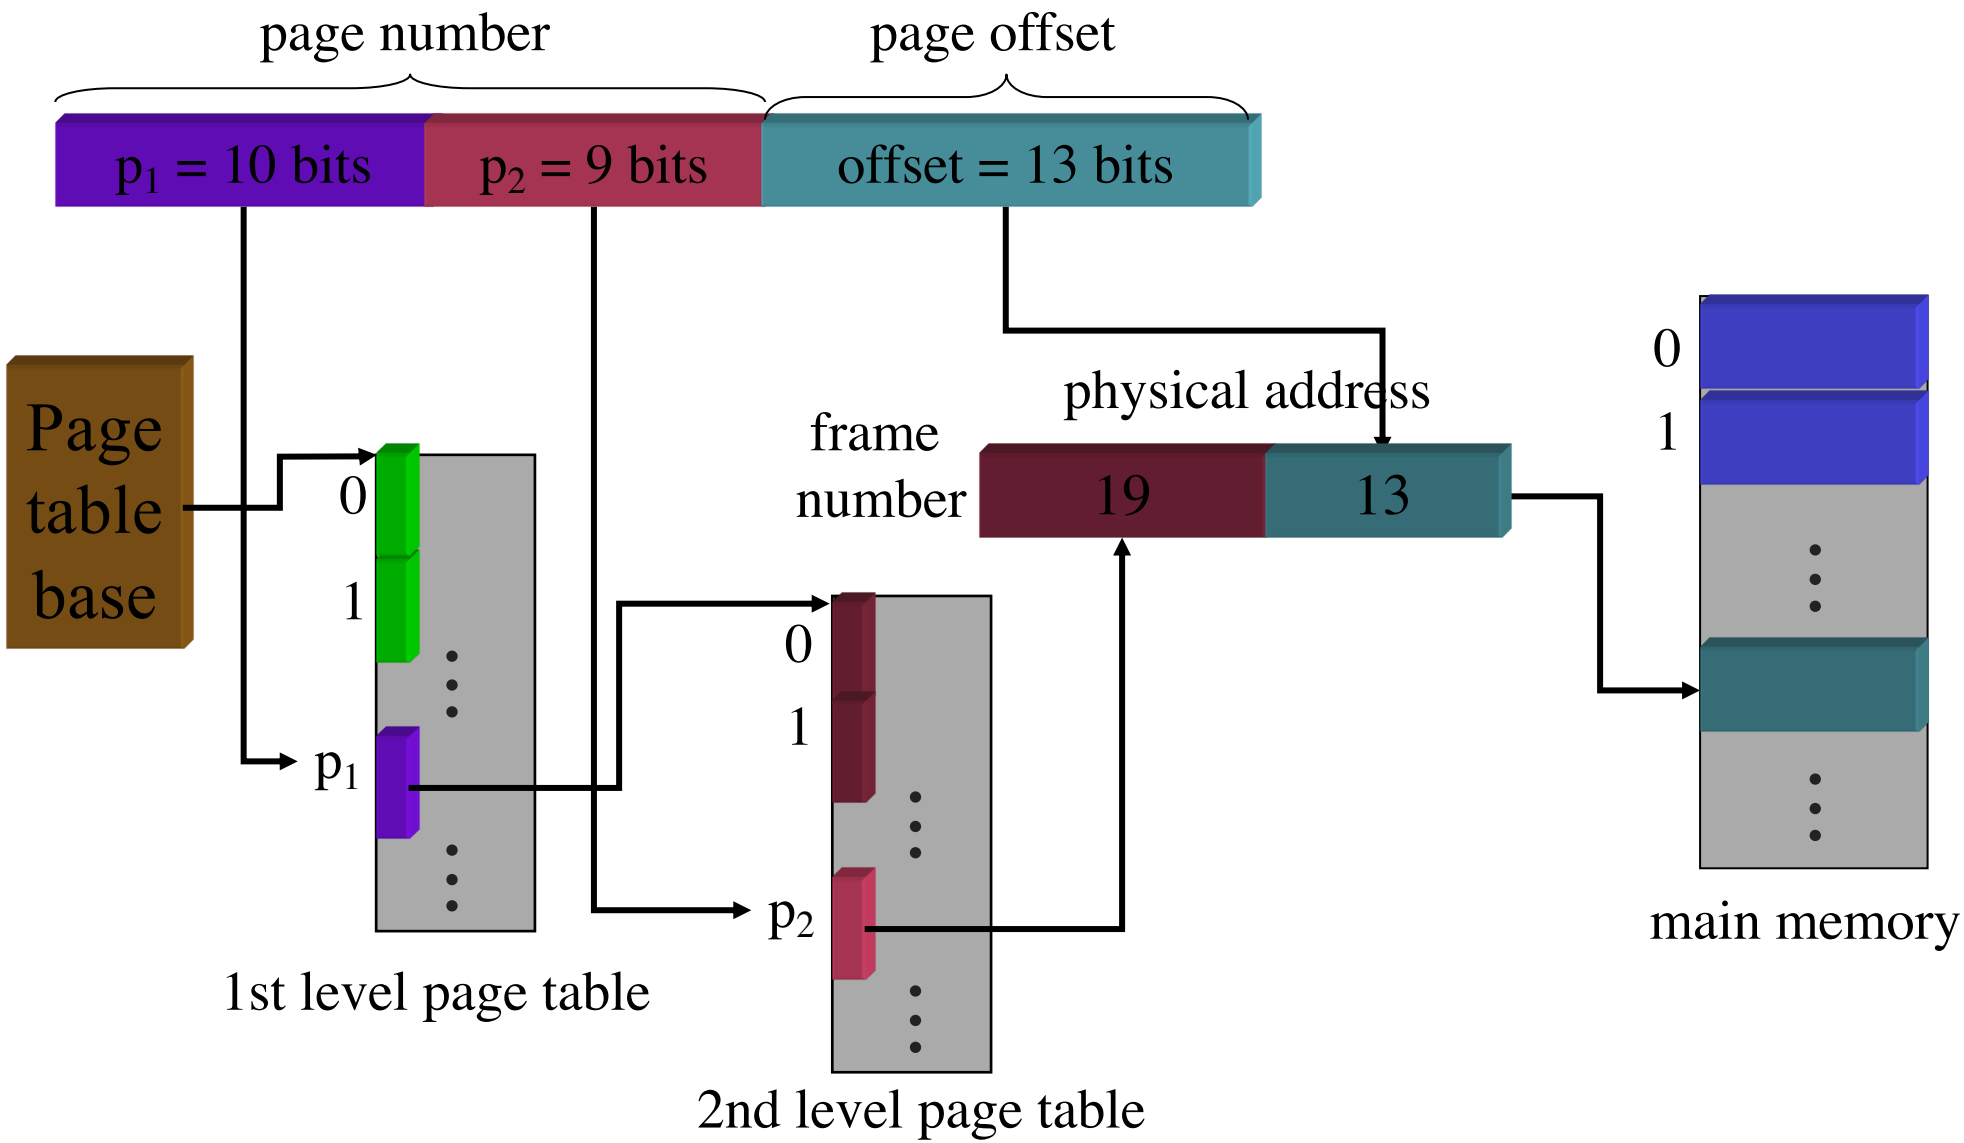

Address Translation: 2-level Page Table

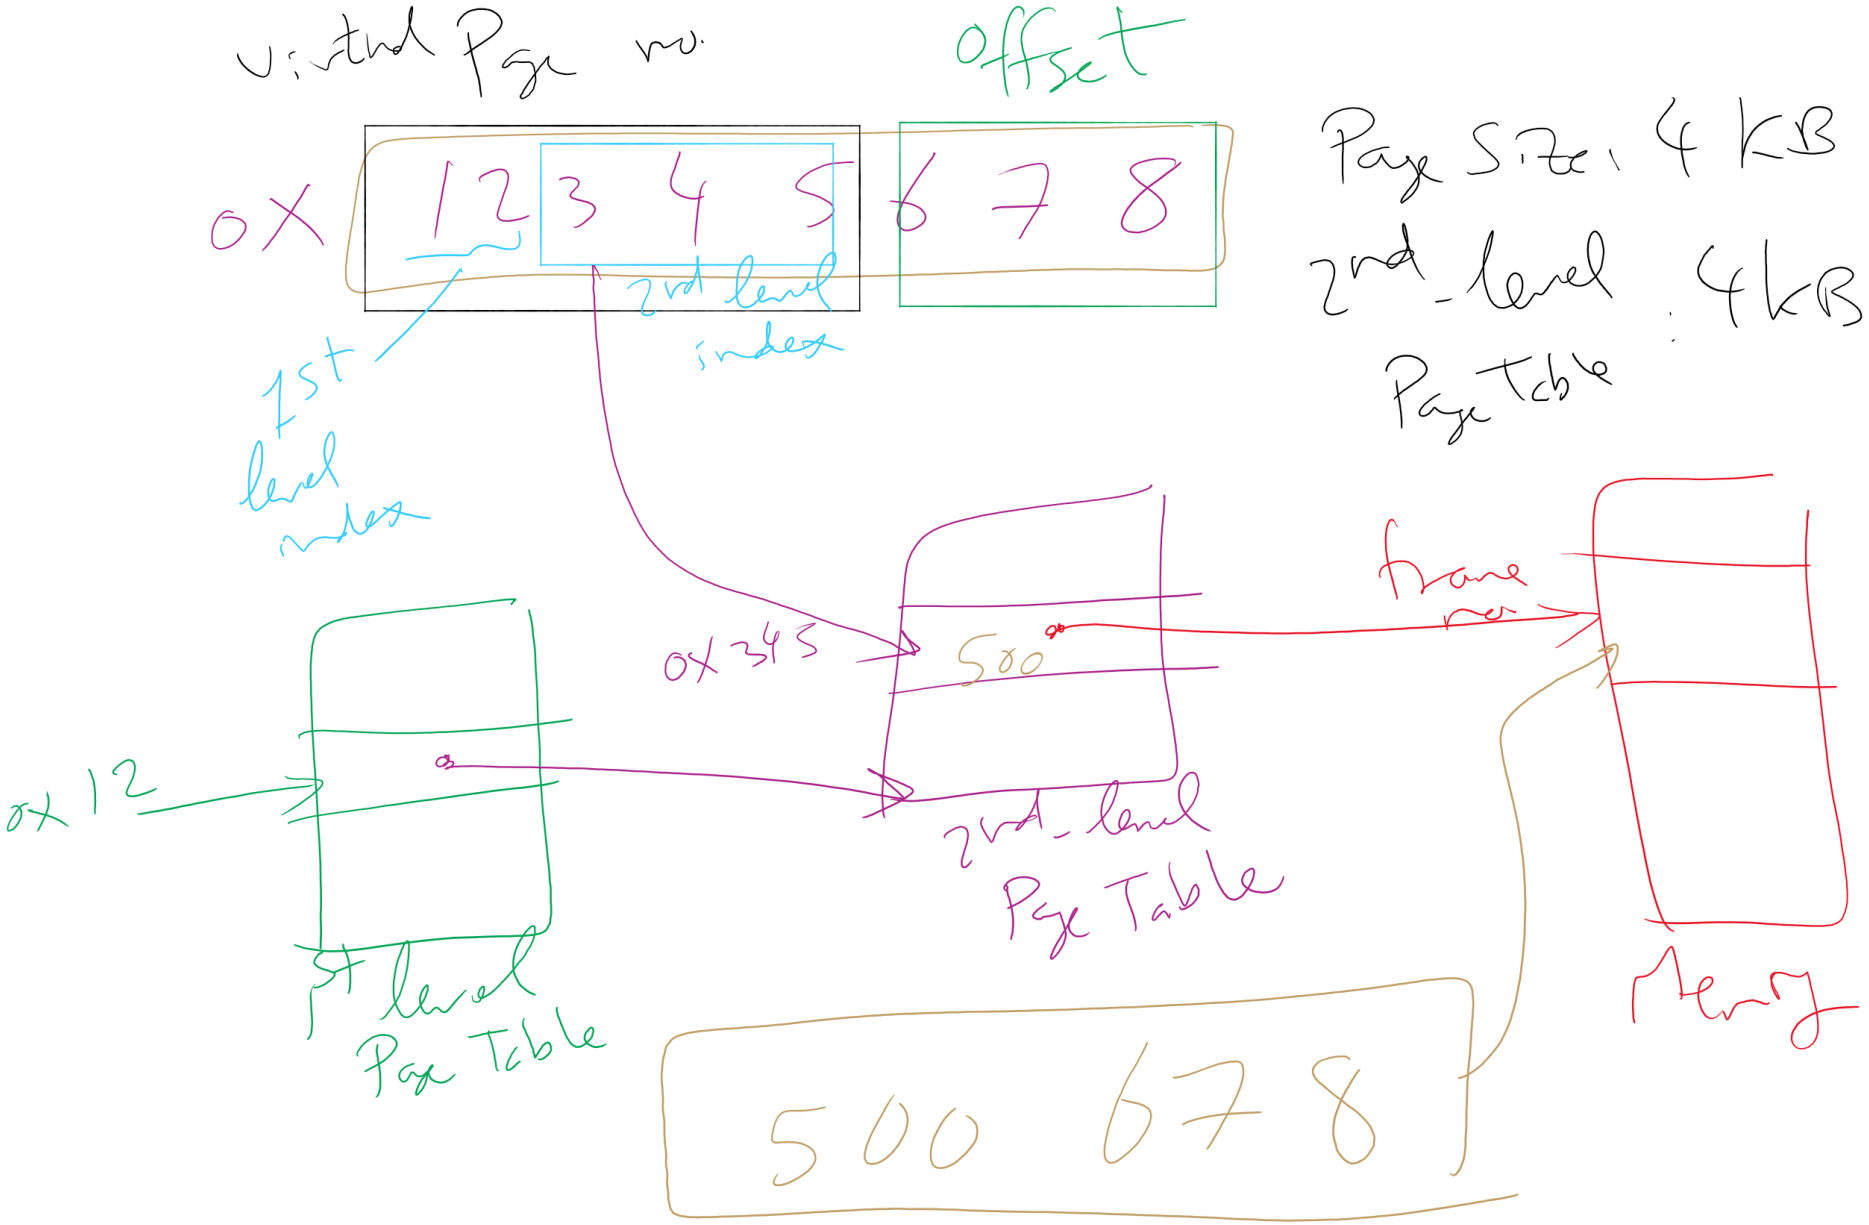

Translation Lookaside Buffer (TLB)

- Search the TLB for the desired logical page number
  - Search entries in parallel
  - Use standard cache techniques
- If desired logical page number is found, get frame number from TLB
- If desired logical page number isn't found
  - Get frame number from page table in memory
  - Replace an entry in the TLB with the logical & physical page numbers from this reference

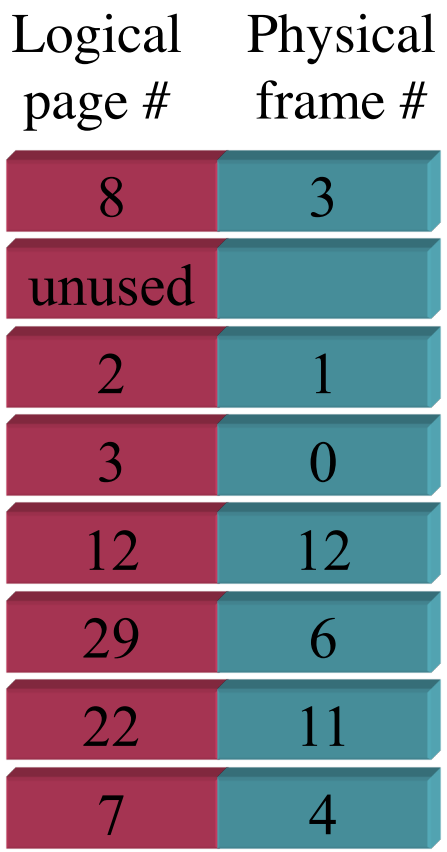

Example TLB

Handling TLB misses

- If PTE isn't found in TLB, OS needs to do the lookup in the page table
- Lookup can be done in hardware or software
- Hardware TLB replacement
  - CPU hardware does page table lookup
  - Can be faster than software
  - Less flexible than software, and more complex hardware
- Software TLB replacement
  - OS gets TLB exception
  - Exception handler does page table lookup & places the result into the TLB
  - Program continues after return from exception
  - Larger TLB (lower miss rate) can make this feasible

TLB

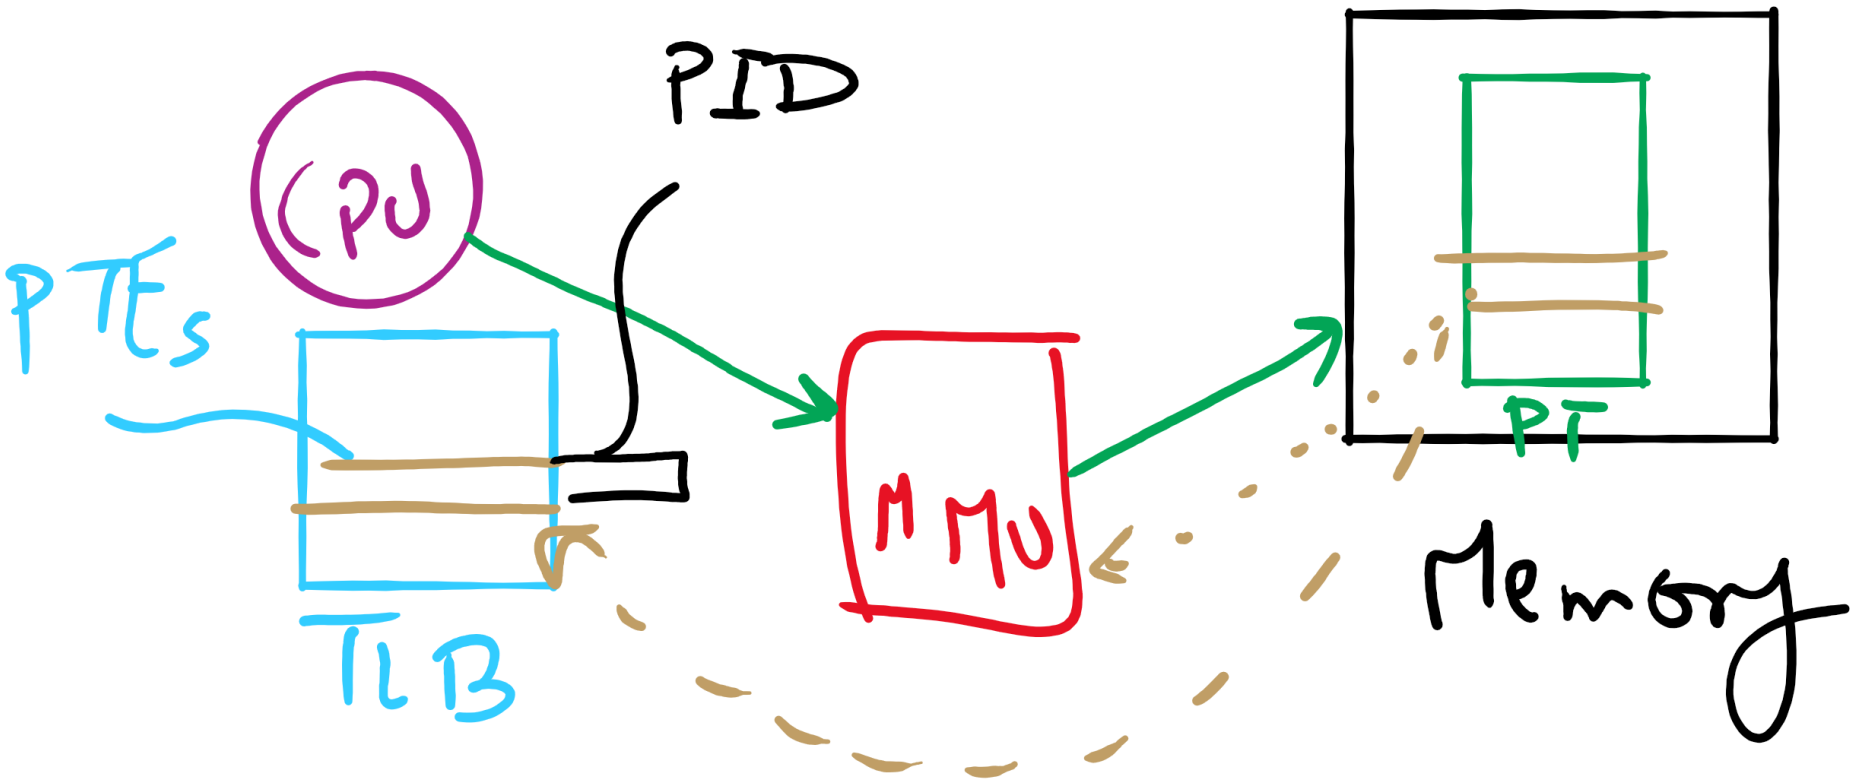

How long do memory accesses take?

- Assume the following times:
  - TLB lookup time = a (often zero - overlapped in CPU)
  - Memory access time = m
- Hit ratio (h) is percentage of time that a logical page number is found in the TLB
  - Larger TLB usually means higher h
  - TLB structure can affect h as well
- Effective access time (an average) is calculated as:
  - EAT = (m + a)h + (m + m + a)(1-h)
  - EAT =a + (2-h)m
- Interpretation
  - Reference always requires TLB lookup, 1 memory access
  - TLB misses also require an additional memory reference

Effective Access Time

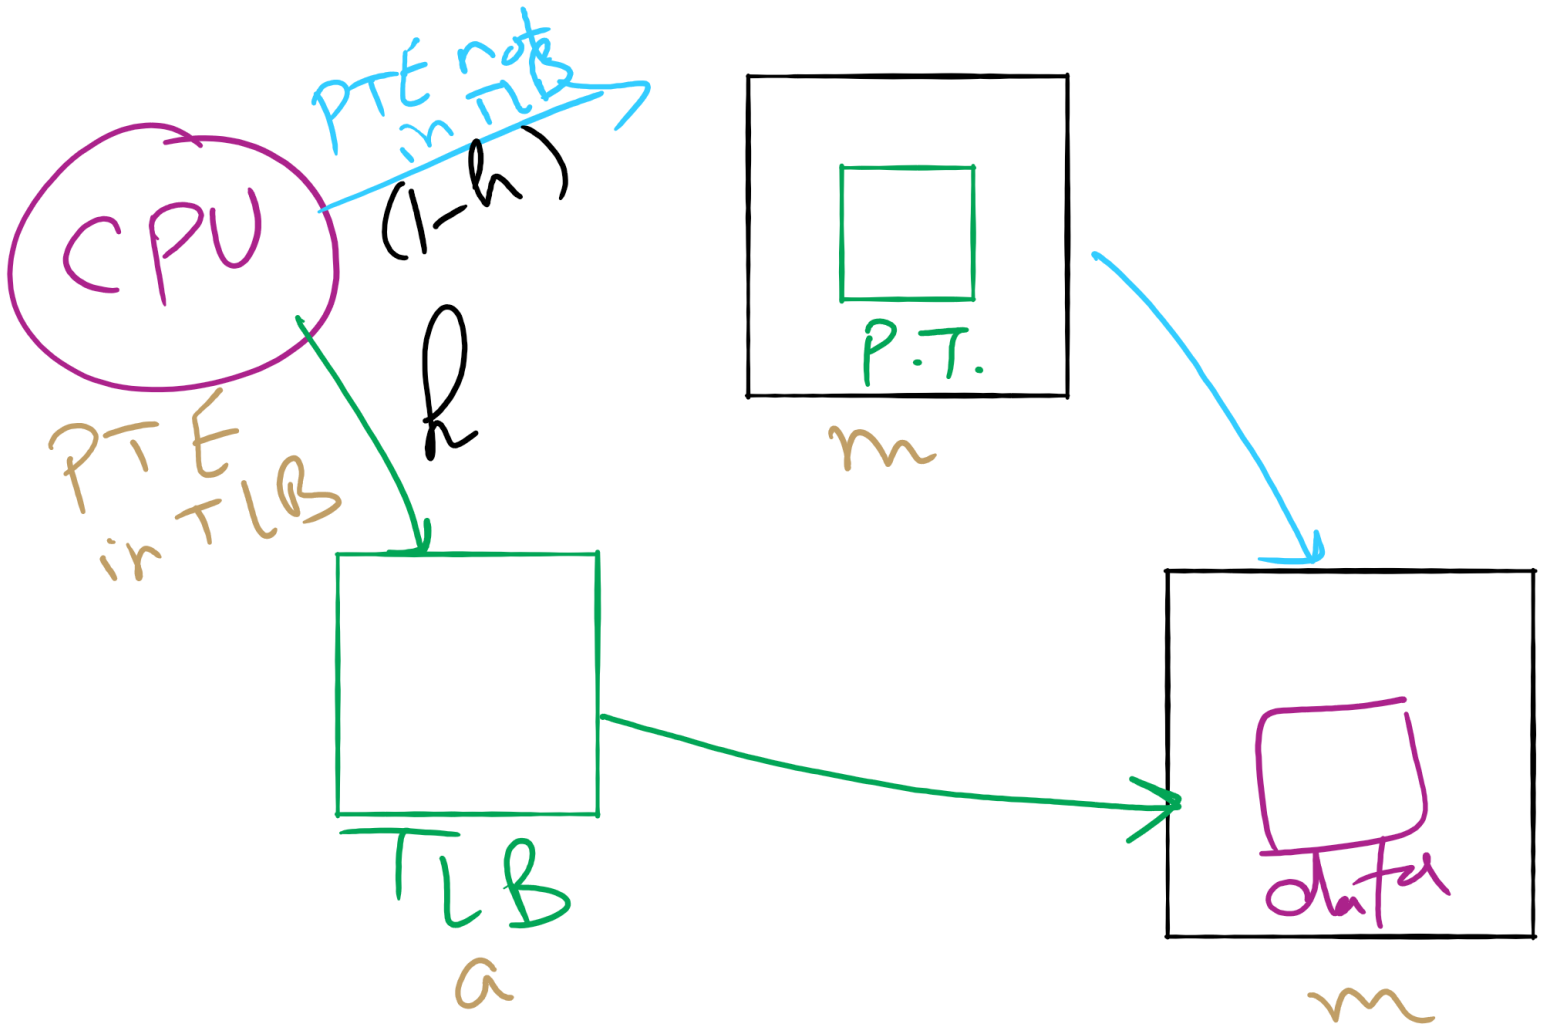

$$EAT = h(a+m) + (1-h)(a+m+m)$$

EAT Calculation

$$EAT = h(a+m) + (1-h)(a+m+m)$$

2-level Page Table → $2m$

Solution 2: Inverted page table

- One entry for each **frame** in memory
- One table for the entire system
- PTE contains
  - Virtual address pointing to this frame
  - Information about the process that owns this page
- Search page table by
  - Hashing the virtual page number and process ID
  - Starting at the entry corresponding to the hash result
  - Search until either the entry is found or a limit is reached
- Page frame number is **index** of PTE in the table
- Improve performance by using more advanced hashing algorithms

Inverted page table architecture

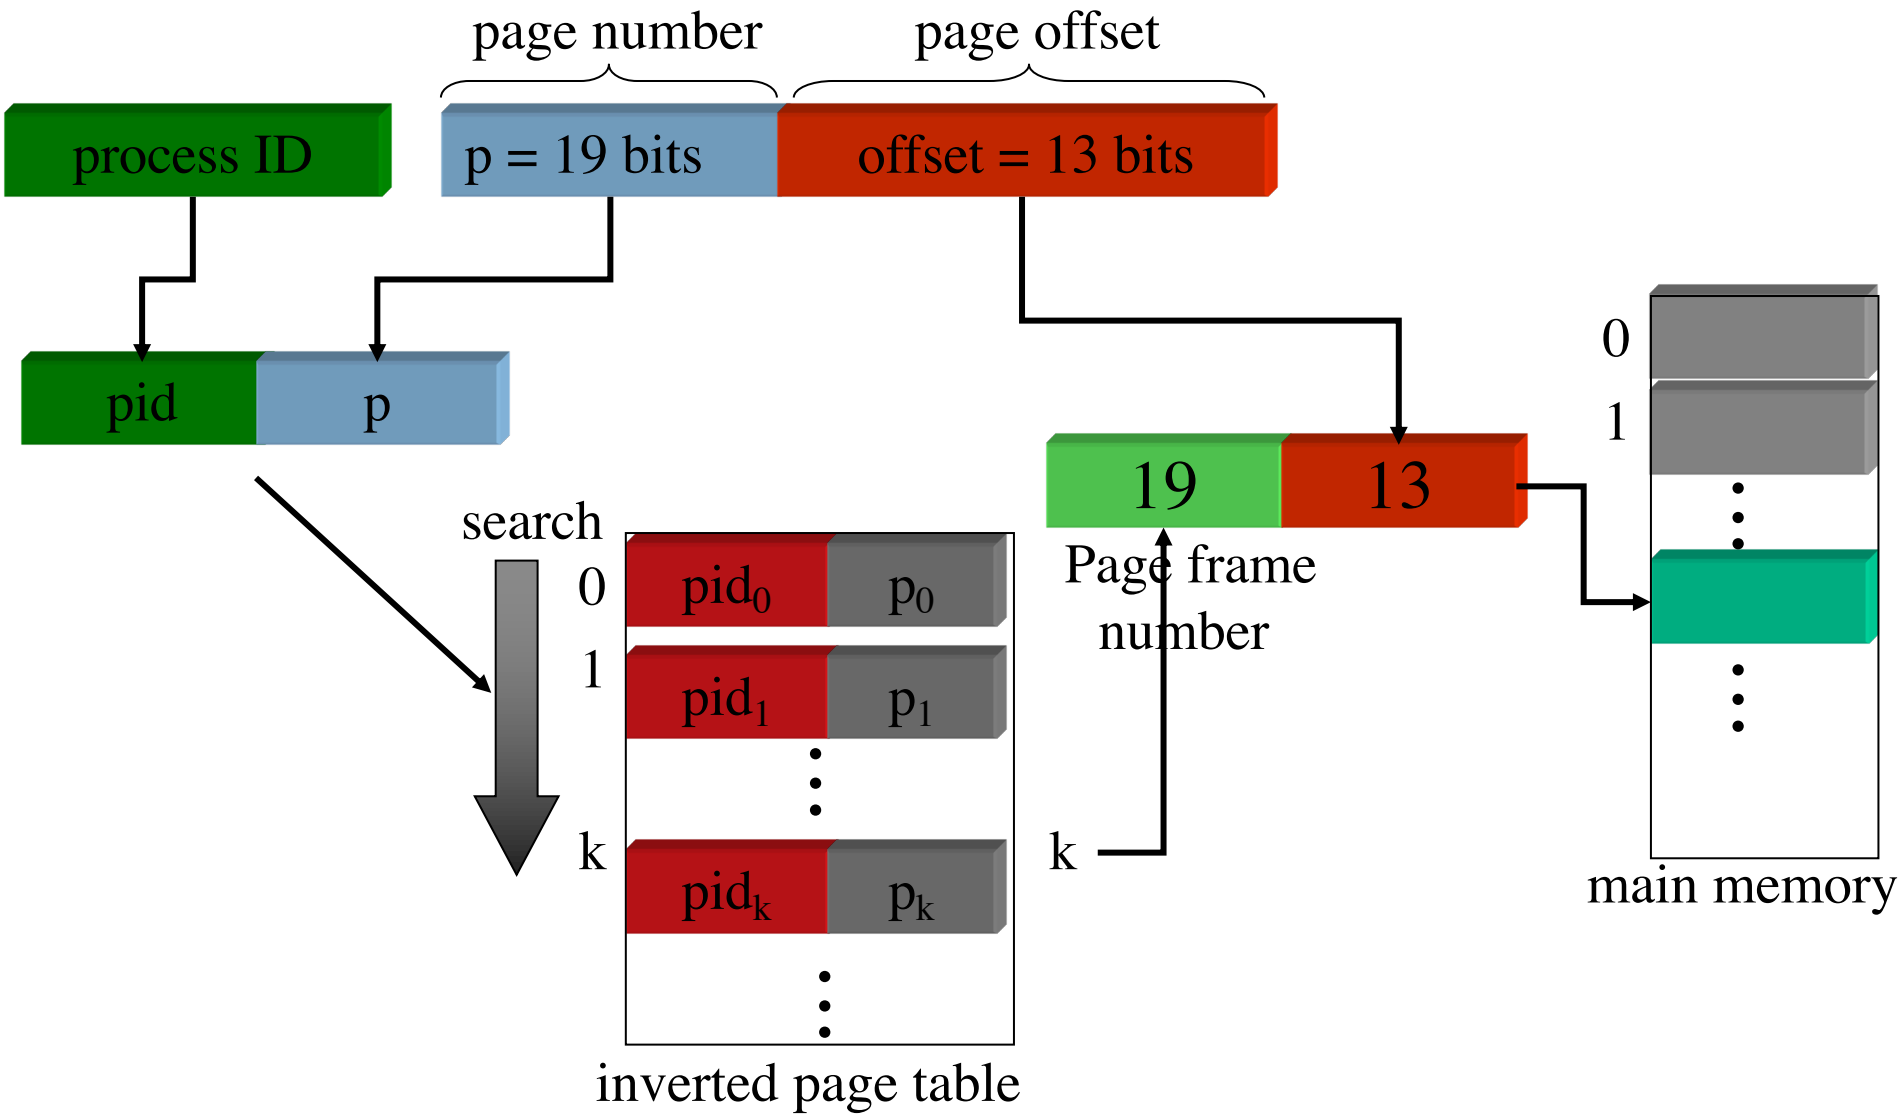

Inverted Page Table Example

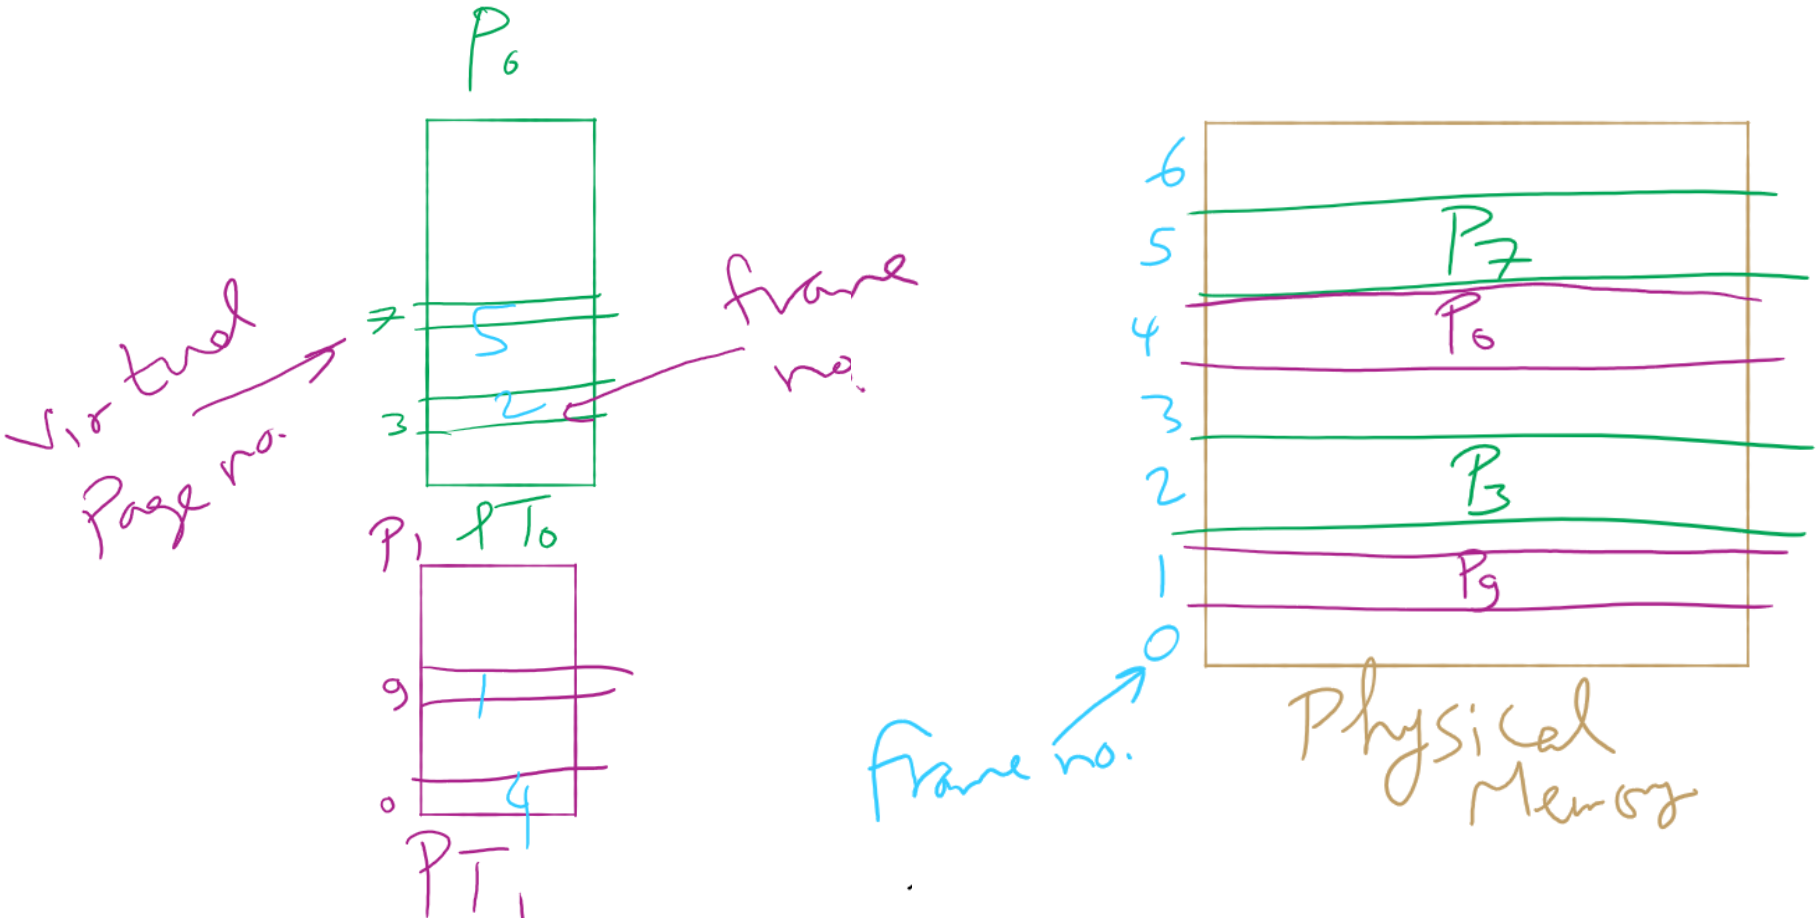

Inverted Page Table Example

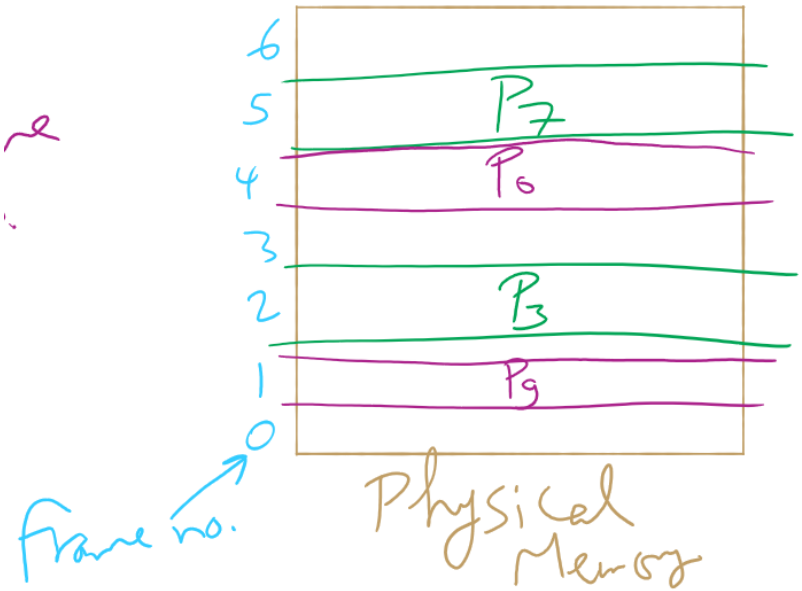

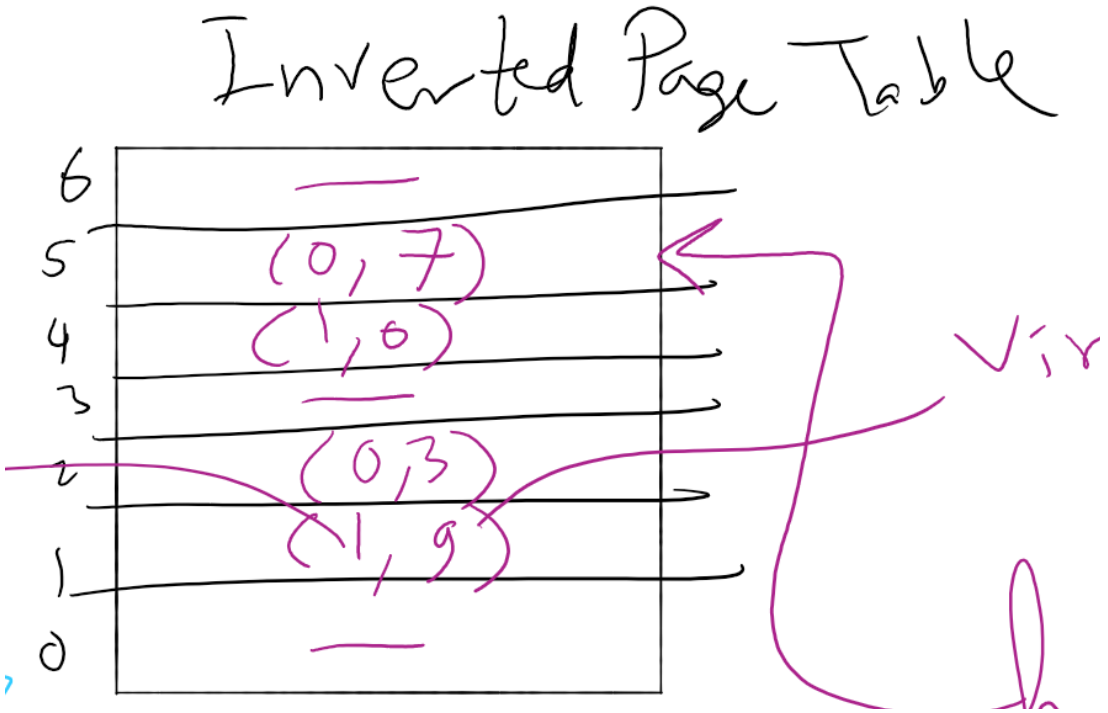

Inverted Page Table Example