



University of  
Pittsburgh

# Introduction to Operating Systems CS 1550



Spring 2023  
Sherif Khattab  
ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013**)

# Announcements

- Upcoming deadlines
  - **All deadlines moved to Monday May 1<sup>st</sup> at 11:59 pm**
    - But please don't wait to last minute!
  - Homework 11, 12, Bonus Homework
  - Lab 4 and Lab 5
  - Quiz 3 and Quiz 4
  - Project 4 (**no late deadline**)

# Final Exam

- **Wednesday 4/26 8:00-9:50**
  - same classroom
  - coffee will be served!
- Same format as midterm
- **Non-cumulative**
- Study guide and practice test on Canvas
- **Review Session** during Finals' Week
  - Date and time TBD
  - recorded

# Bonus Opportunities

- **Bonus Homework**
  - worth up to 1%
  - lowest two homework assignments still dropped
- **Post-Course Quiz on Canvas**
  - worth 1%
- bonus point for class when
  - OMETs response rate  $\geq 80\%$** 
    - Currently at 28%
    - Deadline is Sunday 4/23


# Previous Lecture ...

- How does a file system hide disk access delays?
- How do device drivers program I/O devices?
- Disk arm scheduling
  - FCFS, SSTF, SCAN, C-SCAN, LOOK, C-LOOK

# Today ...

- Miscellaneous issues in File Systems
- Protection in operating systems

# Hard Linking vs. Copying




## hard linking

- doesn't create a new i-node
- increments the **link count** inside original i-node
- e.g., **ln** command in Linux **without -s**

## soft (symbolic linking)

creates a new i-node that contains **path** of original file

e.g., **ln -s** command in Linux




## copying

- creates a new i-node
- e.g., **cp** command in Linux

# File-related kernel structures: open-file tables


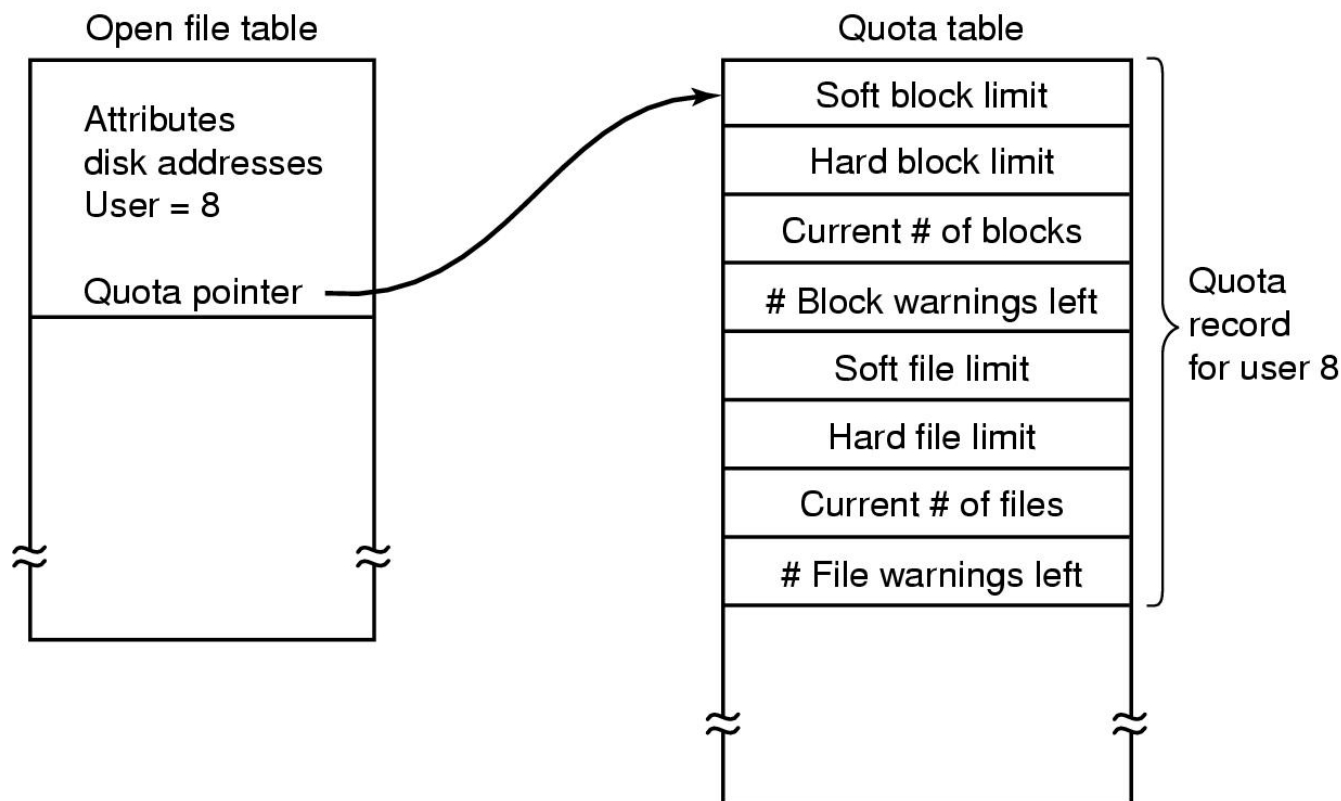
Process



per-process and global open-file tables



# File-related kernel structures: quota table




# Journaling File System

- **Problem:** regular file system: changes to files and directories result in **multiple separate writes** to disk
  - e.g., deleting a file → three writes
  - **power failures** result in file system **inconsistency**
- **Solution:** Write the changes **twice**
  - first to an on-disk ***journal***
    - for efficiency, journal can be put on **SSD or NVRAM**
  - then **commit** changes to main part
  - atomic operations
- modified **data** may or may not be written to the journal
  - implications?
- Examples
  - Windows NTFS
  - Linux ext3, ext4


# Journaling File System

- Interaction with disk arm scheduling?
  - may cause a different write order than wanted
  - **solution:** flush write cache at certain points (barriers)
    - ext3 and ext4 file systems
- Journaling vs. Log-structured file system
  - main file system itself is a journal
  - journaling is not needed in LFS


# Buffering device input




Unbuffered  
input



Buffering in  
user space




Buffer in kernel  
Copy to user space



Double buffer  
in kernel

# Double Buffering



# Max Partition Size

Block size	FAT-12	FAT-16	FAT-32
0.5 KB	2 MB		
1 KB	4 MB		
2 KB	8 MB	128 MB	
4 KB	16 MB	256 MB	1 TB
8 KB		512 MB	2 TB
16 KB		1024 MB	2 TB
32 KB		2048 MB	2 TB

# Max Partition Size

Block size	FAT-12	FAT-16	FAT-32
0.5 KB	2 MB		
1 KB	4 MB		
2 KB	8 MB	128 MB	
4 KB	16 MB	256 MB	1 TB
8 KB		512 MB	2 TB
16 KB		1024 MB	2 TB
32 KB		2048 MB	2 TB

32-bit block no.  
(block no. size)

$$\text{Max. Partition size} = 2 \times \text{block size}$$

FAT-12 :  $2^{12} \times \frac{1}{2} \text{ KB}$   
 $= 2^{12} \times 2^9 = 2^{21} = 2 \text{ MB}$

FAT-16 :  $2^{16} \times 2 \text{ KB}$   
 $= 2^{16} \times 2^{11} = 2^{27} = 128 \text{ MB}$

FAT-32 :  $2^{28} \times 4 \text{ KB}$   
 $= 2^{28} \times 2^{12} = 2^{40} = 1 \text{ TB}$

sector no. size

$$2 \times 512 = 2^{32} \times 512$$

$$= 2^{41} = 2 \text{ TB}$$

# Max Partition Size

Block size	FAT-12	FAT-16	FAT-32
0.5 KB	2 MB		
1 KB	4 MB		
2 KB	8 MB	128 MB	
4 KB	16 MB	256 MB	1 TB
8 KB		512 MB	2 TB
16 KB		1024 MB	2 TB
32 KB		2048 MB	2 TB

$$\begin{array}{l} \text{Max.} \\ \text{Partition} \\ \text{Size} \end{array} = 2^{32\text{-bit block no.}} \times \text{block size}$$

(block no. size)



# Max Partition Size

Block size	FAT-12	FAT-16	FAT-32
0.5 KB	2 MB		
1 KB	4 MB		
2 KB	8 MB	128 MB	
4 KB	16 MB	256 MB	1 TB
8 KB		512 MB	2 TB
16 KB		1024 MB	2 TB
32 KB		2048 MB	2 TB

32-bit block no.  
(block no. size)

Max. Partition size =  $2^{32} \times \text{block size}$

FAT-12 :  $2^{12} \times \frac{1}{2} \text{ KB}$   
 $= 2^{12} \times 2^9 = 2^{21} = 2 \text{ MB}$

# Max Partition Size

Block size	FAT-12	FAT-16	FAT-32
0.5 KB	2 MB		
1 KB	4 MB		
2 KB	8 MB	128 MB	
4 KB	16 MB	256 MB	1 TB
8 KB		512 MB	2 TB
16 KB		1024 MB	2 TB
32 KB		2048 MB	2 TB

FAT-16 :  $2^{16} * 2 \text{ KB}$   
 $= 2^{16} * 2^{11} = 2^{27} = 128 \text{ MB}$

# Max Partition Size


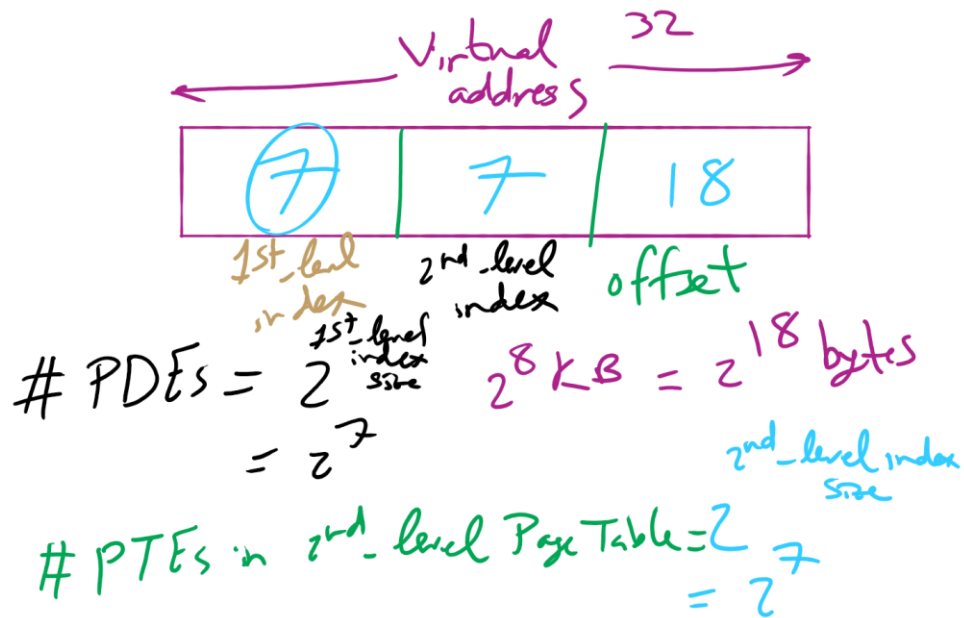
Block size	FAT-12	FAT-16	FAT-32
0.5 KB	2 MB		
1 KB	4 MB		
2 KB	8 MB	128 MB	
4 KB	16 MB	256 MB	1 TB
8 KB		512 MB	2 TB
16 KB		1024 MB	2 TB
32 KB		2048 MB	2 TB

FAT-32 :  $2^{28} * 4 \text{ KB}$   
 $= 2^{28} * 2^{12} = 2^{40} = 1 \text{ TB}$

sector no.  
size

$2 * 512 = 2^{32} * 512$   
 $= 2^{41} = 2 \text{ TB}$

# HW 10: Q 2-4



# HW 10: Q 10-13

single-level Page Table

$$EAT = h(a+m) + (1-h)(a+r)_m$$

2-level Page Table

$$EAT = h(a+r) + (1-h)(a+m+r)_{+m}$$

(PID, VPN)  
(30, 1)

0x1200  
offset 12  
Page No 4 KB = 2  
⇒ offset : 12 bits

Frame no. : 3

31200  
3200

0x4500  
offset  
frame no. : 4  
PID  
(10, 7)  
frame no. : 4  
0x7500


# Effective Disk Access Time

$$EDAT = h * \text{hit time} + (1-h) * \text{miss time}$$

Diagram illustrating the Effective Disk Access Time (EDAT) formula with handwritten annotations:

- $h$  (hit probability) is annotated with  $0.93$ .
- $\text{hit time}$  is annotated with  $0.01$ .
- $(1-h)$  (miss probability) is annotated with  $0.07$ .
- $\text{miss time}$  is annotated with  $0.01 + 3 * 10$ .


# Minimum and Maximum Seek Time



$$\text{Max Seek time} = \text{min seek time} * (\# \text{ tracks} - 1)$$

↑  
# cylinders

# Average Rotational Delay



$$\begin{aligned}\text{Average rotational delay} &= \frac{\text{one full rotation time}}{2} \\ &= \frac{1 / \text{rotation speed}}{2} \\ &= \frac{(1/4800)^2 \times 60 \times 1000}{2} \text{ ms}\end{aligned}$$




# Problem of the Day: Protection

- Protection is about controlling access of programs, processes, or users to the system resources (e.g., memory pages, files, devices, CPUs)
  - How to decide who can access what?
  - Specifications must be
    - Correct
    - Efficient
    - Easy to use (or nobody will use them!)

# Protection domains

- A process operates within a protection domain
  - resources accessible by the process
  - each domain lists objects with permitted operations
- Domains can share objects & permissions
  - Objects can have different permissions in different domains
  - There need be no overlap between object permissions in different domains
- How can this arrangement be specified more formally?



# Protection matrix

	File1	File2	File3	File4	File5	Printer1	Camera
Domain1	Read	Read Write					
Domain2			Read	Read Write Execute	Read Write	Write	
Domain3			Write			Write	Read

- Each domain has a row in the matrix
- Each object (resource) has a column in the matrix
- Entry for <object, column> has the permissions
- Who's allowed to modify the protection matrix?
  - What changes can they make?
- How is this implemented efficiently?

# Domains as objects in the protection matrix

Domain	File1	File2	File3	File4	File5	Printer1	Camera	Dom1	Dom2	Dom3
1	Read	Read Write						Modify		
2			Read	Read Write Execute	Read Write	Write		Modify		
3			Write			Write	Read		Enter	

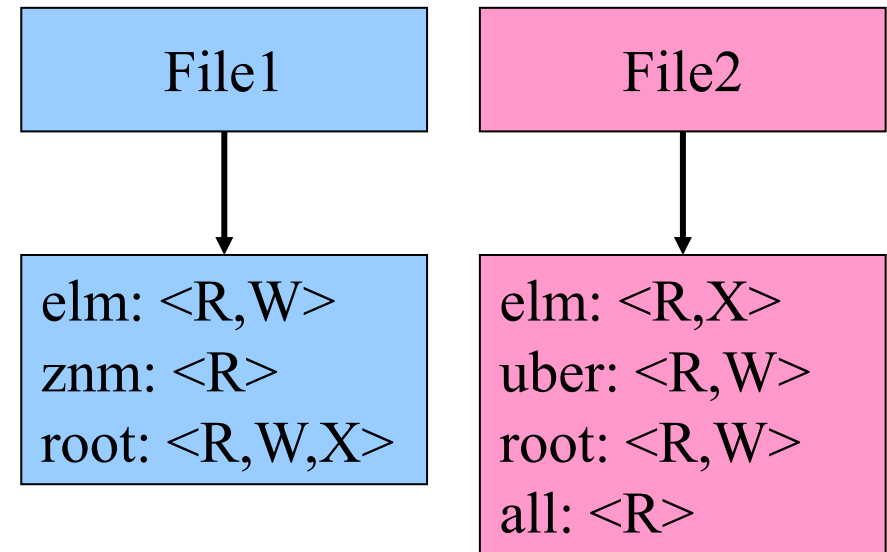
- Specify permitted operations on domains in the matrix
  - Domains may (or may not) be able to modify themselves
  - Domains can modify other domains
  - Some domain transfers (switching) permitted, others not
- Doing this allows flexibility in specifying domain permissions
  - Retains ability to restrict modification of domain policies

# Representing the protection matrix

- Need to find an efficient representation of the protection matrix (also called the *access matrix*)
- Most entries in the matrix are empty!
- Compress the matrix by:
  - Associating permissions with each object: *access control list*
  - Associating permissions with each domain: *capabilities*
- How is this done, and what are the tradeoffs?

# Access control lists (ACLs)

- Each object has a list attached to it
- List has
  - Protection domain (User name, Group of users, Other)
  - Access rights (Read, Write, Execute, Others)
- No entry for domain => no rights for that domain
- Operating system checks permissions when access is needed



# Access control lists in the real world

- Unix file system
  - Access list for each file has exactly three domains on it
    - User (owner)
    - Group
    - Others
  - Rights include read, write, execute: interpreted differently for directories and files
- AFS
  - Access lists only apply to directories: files inherit rights from the directory they're in
  - Access list may have many entries on it with possible rights:
    - read, write, lock (for files in the directory)
    - lookup, insert, delete (for the directories themselves),
    - administer (ability to add or remove rights from the ACL)

# ACL in UNIX

Special flags			owner (user)			group of owner			others		
set user id	set group id	sticky bit	r	w	x	r	w	x	r	w	x
1	1	1	1	1	0	0	0	1	0	0	1
							1			1	
				6							
			1	1	1	1	0	1	1	0	1
							5			5	
			1	1	0	1	0	0	1	0	0
				6			4			4	




# Domains as objects in the protection matrix

Domain	File1	File2	File3	File4	File5	Printer1	Camera	Dom1	Dom2	Dom3
1	Read	Read Write						Modify		
2			Read	Read Write Execute	Read Write	Write		Modify		
3			Write			Write	Read		Enter	

- Specify permitted operations on domains in the matrix
  - Domains may (or may not) be able to modify themselves
  - Domains can modify other domains
  - Some domain transfers (switching) permitted, others not
- Doing this allows flexibility in specifying domain permissions
  - Retains ability to restrict modification of domain policies

# Capabilities

- Each process has a capability list
- List has one entry per object the process can access
  - Object name
  - Object permissions
- Objects not listed are not accessible
- How are these secured?
  - Kept in kernel
  - Cryptographically secured



# Cryptographically protected capability

Server	Object	Rights	$H(\textit{Object}, \textit{Rights}, \textit{Check})$
--------	--------	--------	---

- Rights include generic rights (read, write, execute) and
  - Copy capability
  - Copy object
  - Remove capability
  - Destroy object
- Server has a secret (*Check*) and uses it to verify capabilities presented to it
  - Alternatively, use public-key signature techniques

# Protecting the access matrix: summary

- OS must ensure that the access matrix isn't modified (or even accessed) in an unauthorized way
- Access control lists
  - Reading or modifying the ACL is a system call
  - OS makes sure the desired operation is allowed
- Capability lists
  - Can be handled the same way as ACLs: reading and modification done by OS
  - Can be handed to processes and verified cryptographically later on
    - May be better for widely distributed systems where capabilities can't be centrally checked